

# V-Series Avalon-MM DMA Interface for PCIe Solutions User Guide



Subscribe



Send Feedback

Last updated for Quartus Prime Design Suite: 16.1

**UG-01154**  
2018.07.31

101 Innovation Drive  
San Jose, CA 95134  
[www.altera.com](http://www.altera.com)



# Contents

<b>Datasheet.....</b>	<b>1-1</b>
V-Series Avalon-MM DMA Interface for PCIe* Datasheet.....	1-1
Features .....	1-2
Comparison of Avalon-ST, Avalon-MM and Avalon-MM with DMA Interfaces for V-Series	
Devices.....	1-4
Release Information .....	1-7
V-Series Device Family Support .....	1-7
Design Examples.....	1-8
Debug Features .....	1-8
IP Core Verification .....	1-8
Compatibility Testing Environment .....	1-9
Resource Utilization .....	1-9
V-Series Recommended Speed Grades .....	1-9
Creating a Design for PCI Express.....	1-13
 <b>Getting Started with the Avalon-MM DMA .....</b>	 <b>2-1</b>
Understanding the Avalon-MM DMA Ports.....	2-3
Generating the Testbench .....	2-5
Understanding the Simulation Generated Files .....	2-7
Understanding Simulation Log File Generation.....	2-7
Simulating the Example Design in ModelSim.....	2-7
Running a Gate-Level Simulation.....	2-8
Generating Synthesis Files.....	2-8
Compiling the Design .....	2-8
Creating a Quartus Prime Project .....	2-8
 <b>Parameter Settings.....</b>	 <b>3-1</b>
System Settings.....	3-1
Base Address Register (BAR) Settings .....	3-4
Device Identification Registers .....	3-5
PCI Express and PCI Capabilities Parameters .....	3-6
Device Capabilities .....	3-6
Error Reporting .....	3-7
Link Capabilities .....	3-7
MSI and MSI-X Capabilities .....	3-8
Power Management .....	3-9
PCIe Address Space Settings.....	3-10
 <b>IP Core Interfaces .....</b>	 <b>4-1</b>
V-Series DMA Avalon-MM DMA Interface to the Application Layer.....	4-1

Avalon-MM DMA Interfaces when Descriptor Controller Is Internally Instantiated.....	4-3
Read Data Mover .....	4-6
Write DMA Avalon-MM Master Port .....	4-8
RX Master Module .....	4-9
Non-Bursing Slave Module .....	4-11
32-Bit Control Register Access (CRA) Slave Signals .....	4-12
Avalon-ST Descriptor Control Interface when Instantiated Separately .....	4-13
Descriptor Controller Interfaces when Instantiated Internally .....	4-16
Clock Signals .....	4-19
Reset, Status, and Link Training Signals.....	4-19
MSI Interrupts for Endpoints .....	4-25
Physical Layer Interface Signals .....	4-26
Transceiver Reconfiguration .....	4-26
Serial Data Signals .....	4-27
PIPE Interface Signals .....	4-38
Test Signals .....	4-42
<b>Registers.....</b>	<b>5-1</b>
Correspondence between Configuration Space Registers and the PCIe Specification .....	5-1
Type 0 Configuration Space Registers .....	5-6
Type 1 Configuration Space Registers .....	5-7
PCI Express Capability Structures.....	5-7
Intel-Defined VSEC Registers.....	5-8
CvP Registers.....	5-10
Advanced Error Reporting Capability.....	5-12
Uncorrectable Internal Error Mask Register .....	5-12
Uncorrectable Internal Error Status Register .....	5-13
Correctable Internal Error Mask Register .....	5-14
Correctable Internal Error Status Register .....	5-15
DMA Descriptor Controller Registers .....	5-15
Read DMA Descriptor Controller Registers.....	5-16
Write DMA Descriptor Controller Registers.....	5-19
Read DMA and Write DMA Descriptor Format .....	5-21
Read DMA Example .....	5-24
Software Program for Simultaneous Read and Write DMA .....	5-29
Control Register Access (CRA) Avalon-MM Slave Port .....	5-30
<b>Reset and Clocks.....</b>	<b>6-1</b>
Reset Sequence for Hard IP for PCI Express IP Core and Application Layer .....	6-3
Clocks .....	6-5
Clock Domains .....	6-5
Clock Summary .....	6-7
<b>Error Handling .....</b>	<b>7-1</b>
Physical Layer Errors .....	7-1
Data Link Layer Errors .....	7-2

Transaction Layer Errors .....	7-3
Error Reporting and Data Poisoning .....	7-5
Uncorrectable and Correctable Error Status Bits .....	7-6
<b>PCI Express Protocol Stack.....</b>	<b>8 9-1</b>
Top-Level Interfaces .....	8 9-3
Avalon-MM DMA Interface.....	8 9-3
Clocks and Reset .....	8 9-3
Transceiver Reconfiguration .....	8 9-3
Interrupts .....	8 9-3
PIPE .....	8 9-4
Data Link Layer .....	8 9-4
Physical Layer .....	8 9-6
<b>V-Series Avalon-MM DMA for PCI Express .....</b>	<b>-1</b>
Understanding the Internal DMA Descriptor Controller.....	-1
Understanding the External DMA Descriptor Controller.....	-3
<b>Transceiver PHY IP Reconfiguration .....</b>	<b>10-1</b>
Connecting the Transceiver Reconfiguration Controller IP Core .....	10-1
Transceiver Reconfiguration Controller Connectivity for Designs Using CvP .....	10-4
<b>Frequently Asked Questions for V-Series Avalon-MM DMA Interface for PCIe.....</b>	<b>A-1</b>
<b>V-Series Interface for PCIe Solutions User Guide Archive .....</b>	<b>B-1</b>
<b>Document Revision History.....</b>	<b>C-1</b>
Document Revision History for the V-Series Avalon-MM DMA Interface for PCIe Solutions User Guide.....	C-1

2018.07.31

UG-01154



Subscribe



Send Feedback

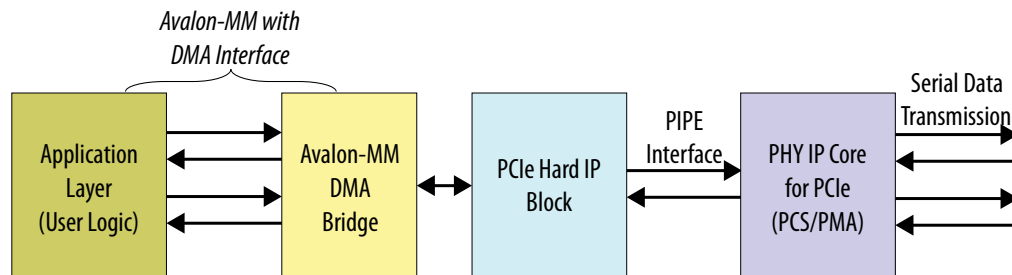
## V-Series Avalon-MM DMA Interface for PCIe\* Datasheet

Intel® V-Series FPGAs include a configurable, hardened protocol stack for PCI Express\* that is compliant with *PCI Express Base Specification 2.1 or 3.0*.

The V-Series Avalon® Memory-Mapped (Avalon-MM) DMA for PCI Express removes some of the complexities associated with the PCIe\* protocol. For example, the IP core handles TLP encoding and decoding. In addition, it includes Read DMA and Write DMA engines. If you have already architected your own DMA system with the Avalon-MM interface, you may want to continue to use it. However, you may want to take advantage of the simplicity of having the DMA engines already implemented. Intel recommends this variant for new users. Depending of the device you select, this variant is available in Qsys for 128- and 256-bit interfaces to the Application Layer. The Avalon-MM interface and DMA engines are implemented in FPGA soft logic.

**Figure 1-1: V-Series PCIe Variant with Avalon-MM DMA Interface**

The following figure shows the high-level modules and connecting interfaces for this variant.



**Note:** This variant was renamed in the Quartus® II 14.0 release. The name in the Quartus II 13.1 release was Avalon-MM 256-bit Hard IP for PCI Express IP Core.

**Table 1-1: PCI Express Data Throughput**

The following table shows the aggregate bandwidth of a PCI Express link for Gen1, Gen2, and Gen3 for 2, 4, and 8 lanes. The protocol specifies 2.5 giga-transfers per second for Gen1, 5.0 giga-transfers per second for Gen2, and 8.0 giga-transfers per second for Gen3. This table provides bandwidths for a single transmit (TX) or receive (RX) channel. The numbers double for duplex operation. Gen1 and Gen2 use 8B/10B

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

**ALTERA**  
now part of Intel

encoding which introduces a 20% overhead. In contrast, Gen3 uses 128b/130b encoding which introduces only a 1.5% overhead.

Units are Gigabits per second (Gbps).

	Link Width		
	×2	×4	×8
PCI Express Gen1 (2.5 Gbps)	N/A	N/A	16 Gbps
PCI Express Gen2 (5.0 Gbps)	8 Gbps	16 Gbps	32 Gbps
PCI Express Gen3 (8.0 Gbps)	15.75 Gbps	31.51 Gbps	63Gbps

#### Related Information

- [V-Series Interface for PCIe Solutions User Guide Archive](#) on page 12-1
- [Introduction to Altera IP Cores](#)  
Provides general information about all Intel FPGA IP cores, including parameterizing, generating, upgrading, and simulating IP cores.
- [Creating Version-Independent IP and Qsys Simulation Scripts](#)  
Create simulation scripts that do not require manual updates for software or IP version upgrades.
- [Project Management Best Practices](#)  
Guidelines for efficient management and portability of your project and IP files.
- [PCI Express Base Specification 2.1 or 3.0](#)

## Features

New features in the Quartus® Prime 16.1 software release:

- Increased maximum DMA transfer size for the 128- and 256-bit interfaces to 1 megabyte (MB).

The V-Series Avalon-MM DMA for PCI Express supports the following features:

- Complete protocol stack including the Transaction, Data Link, and Physical Layers implemented as hard IP.
- Native support for Gen1 x8, Gen2 x4, Gen2 x8, Gen3 x2, Gen3 x4, Gen3 x8 for Endpoints. The variant downtrains when plugged into a lesser link width or changes to a different maximum link rate.
- Dedicated 16 kilobyte (KB) receive buffer.
- Optional hard reset controller for Gen2.
- Support for 128- or 256-bit Avalon-MM interface to Application Layer with embedded DMA up to Gen3 ×8 data rate.
- Support for 32- or 64-bit addressing for the Avalon-MM interface to the Application Layer.
- Qsys design example demonstrating parameterization, design modules, and connectivity.
- Extended credit allocation settings to better optimize the RX buffer space based on application type.

- Optional end-to-end cyclic redundancy code (ECRC) generation and checking and advanced error reporting (AER) for high reliability applications.
- Support for Configuration Space Bypass Mode (in Arria® V GZ and Stratix® V), allowing you to design a custom Configuration Space and support multiple functions.
- Support for Gen3 PIPE simulation.
- Easy to use:
  - Flexible configuration.
  - No license requirement.
  - Design examples to get started.

**Table 1-2: Feature Comparison for 128- and 256-Bit Avalon-MM with DMA Interface to the Application Layer**

Feature	128-Bit Interface	256-Bit Interface
Gen1	x8	Not supported
Gen2	x4, x8	x8
Gen3	x4	x4, x8
Root Port	Not supported	Supported
Tags supported	16	16 or 256
Maximum descriptor size	1 MB	1 MB
Maximum payload size	128 or 256	128 or 256
Immediate write <sup>(1)</sup>	Not supported	Supported

<sup>(1)</sup> The Immediate Write provides a fast mechanism to send a Write TLP upstream. The descriptor stores the 32-bit payload, replacing the Source Low Address field of the descriptor.

## Comparison of Avalon-ST, Avalon-MM and Avalon-MM with DMA Interfaces for V-Series Devices

**Table 1-3: Feature Comparison for all Hard IP for PCI Express IP Cores**

The table compares the features of the three mainstream Hard IP for PCI Express IP Cores. Refer to the *Stratix V Avalon-ST Interface with SR-IOV PCIe Solutions User Guide* for the features of that variant.

Feature	Avalon-ST Interface	Avalon-MM Interface	Avalon-MM DMA
IP Core License	Free	Free	Free
Native Endpoint	Supported	Supported	Supported
Legacy Endpoint <sup>(2)</sup>	Supported	Not Supported	Not Supported
Root port	Supported	Supported	Not Supported
Gen1	×1, ×2, ×4, ×8	×1, ×2, ×4, ×8	×8
Gen2	×1, ×2, ×4, ×8	×1, ×2, ×4, ×8	×4, ×8
Gen3	×1, ×2, ×4, ×8	×1, ×2, ×4	×4, ×8
64-bit Application Layer interface	Supported	Supported	Not supported
128-bit Application Layer interface	Supported	Supported	Supported
256-bit Application Layer interface	Supported	Not Supported	Supported
Maximum payload size	128, 256, 512, 1024, 2048 bytes	128, 256 bytes	128, 256 bytes
Number of tags supported for non-posted requests	256	8 for 64-bit interface 16 for 128-bit interface	16 or 256
62.5 MHz clock	Supported	Supported	Not Supported

<sup>(2)</sup> Not recommended for new designs.





Feature	Avalon-ST Interface	Avalon-MM Interface	Avalon-MM DMA
Automatically handle out-of-order completions (transparent to the Application Layer)	Not supported	Supported	Not Supported
Automatically handle requests that cross 4 KB address boundary (transparent to the Application Layer)	Not supported	Supported	Supported
Polarity Inversion of PIPE interface signals	Supported	Supported	Supported
ECRC forwarding on RX and TX	Supported	Not supported	Not supported
Number of MSI requests	1, 2, 4, 8, 16, or 32	1, 2, 4, 8, 16, or 32	1, 2, 4, 8, 16, or 32
MSI-X	Supported	Supported	Supported
Legacy interrupts	Supported	Supported	Supported
Expansion ROM	Supported	Not supported	Not supported

**Table 1-4: TLP Support Comparison for all Hard IP for PCI Express IP Cores**

The table compares the TLP types that the Hard IP for PCI Express IP Cores variants can transmit. Each entry indicates whether this TLP type is supported (for transmit) by Endpoints (EP), Root Ports (RP), or both (EP/RP). For the Avalon-MM DMA interface, a software application programs a descriptor controller to specify DMA transfers between host and IP memory. The Read DMA Avalon-MM Master port and Write DMA Avalon-MM Master port send read and write TLPs, respectively. The optional TX Slave module supports single, non-bursting Memory Write TLPs to send status updates to the host.

TLP (Transmit Support)	Avalon-ST Interface	Avalon-MM Interface	Avalon-MM DMA
Memory Read Request ( $M_{rd}$ )	EP/RP	EP/RP	EP/RP (Read DMA Avalon-MM Master)
Memory Read Lock Request ( $M_{rdLk}$ )	EP/RP		Not supported

TLP (Transmit Support)	Avalon-ST Interface	Avalon-MM Interface	Avalon-MM DMA
Memory Write Request ( <i>MW<sub>r</sub></i> )	EP/RP	EP/RP	EP/RP (Write DMA Avalon-MM Master) (TX Slave - optional)
I/O Read Request ( <i>IOR<sub>d</sub></i> )	EP/RP	EP/RP	Not supported
I/O Write Request ( <i>IOW<sub>r</sub></i> )	EP/RP	EP/RP	Not supported
Config Type 0 Read Request ( <i>CfgRd0</i> )	RP	RP	Not supported
Config Type 0 Write Request ( <i>CfgWr0</i> )	RP	RP	Not supported
Config Type 1 Read Request ( <i>CfgRd1</i> )	RP	RP	Not supported
Config Type 1 Write Request ( <i>CfgWr1</i> )	RP	RP	Not supported
Message Request ( <i>Msg</i> )	EP/RP	Not supported	Not supported
Message Request with Data ( <i>MsgD</i> )	EP/RP	Not supported	Not supported
Completion ( <i>Cp1</i> )	EP/RP	EP/RP	EP/RP (Read & Write DMA Avalon-MM Masters)
Completion with Data ( <i>Cp1D</i> )	EP/RP	Not supported	EP/RP (Read & Write DMA Avalon-MM Masters)
Completion-Locked ( <i>Cp1Lk</i> )	EP/RP	Not supported	Not supported
Completion Lock with Data ( <i>Cp1DLk</i> )	EP/RP	Not supported	Not supported
Fetch and Add AtomicOp Request ( <i>FetchAdd</i> )	EP	Not supported	Not supported

The *V-Series Avalon-MM DMA Interface for PCIe Solutions User Guide* explains how to use this IP core and not the PCI Express protocol. Although there is inevitable overlap between these two purposes, use this document only in conjunction with an understanding of the *PCI Express Base Specification*.

## Release Information

Table 1-5: Hard IP for PCI Express Release Information

Item	Description
Version	15.1
Release Date	May 2018
Ordering Codes	No ordering code is required
Product IDs	The Product ID and Vendor ID are not required because this IP core does not require a license.
Vendor ID	

Intel verifies that the current version of the Intel® Quartus Prime software compiles the previous version of each IP core, if this IP core was included in the previous release. Intel reports any exceptions to this verification in the *Intel IP Release Notes* or clarifies them in the Intel Quartus Prime IP Update tool. Intel does not verify compilation with IP core versions older than the previous release.

## V-Series Device Family Support

Table 1-6: Device Family Support

Device Family	Support
Arria V, Arria V GZ, Cyclone V, Stratix V	Final. The IP core is verified with final timing models. The IP core meets all functional and timing requirements for the device family and can be used in production designs.
Other device families	Refer to the <i>Related Information</i> below for other device families:

### Related Information

- [Arria V Avalon-MM Interface for PCIe Solutions User Guide](#)
- [Arria V Avalon-ST Interface for PCIe Solutions User Guide](#)
- [Arria V GZ Avalon-MM Interface for PCIe Solutions User Guide](#)
- [Arria V GZ Avalon-ST Interface for PCIe Solutions User Guide](#)
- [Arria 10 Avalon-MM Interface for PCIe Solutions User Guide](#)
- [Arria 10 Avalon-MM DMA Interface for PCIe Solutions User Guide](#)
- [Arria 10 Avalon-ST Interface for PCIe Solutions User Guide](#)

- [Cyclone V Avalon-MM Interface for PCIe Solutions User Guide](#)
- [Cyclone V Avalon-ST Interface for PCIe Solutions User Guide](#)
- [IP Compiler for PCI Express User Guide](#)
- [Stratix V Avalon-MM Interface for PCIe Solutions User Guide](#)
- [Stratix V Avalon-ST Interface for PCIe Solutions User Guide](#)
- [Stratix V Avalon-ST Interface with SR-IOV for PCIe Solutions User Guide](#)

## Design Examples

Platform Designer example designs are available for the V-Series Avalon-MM DMA for PCI Express IP Core. You can download them from the `<install_dir>/ip/altera/altera_pcie/altera_pcie_hip_256_avmm/example_design/<dev>` directory.

Starting from the 18.0 release of the Intel Quartus Prime software, you can generate both Endpoint example designs and Root Port example designs.

### Related Information

[Getting Started with the Avalon-MM DMA](#) on page 2-1

## Debug Features

Debug features allow observation and control of the Hard IP for faster debugging of system-level problems.

## IP Core Verification

To ensure compliance with the PCI Express specification, Intel performs extensive verification. The simulation environment uses multiple testbenches that consist of industry-standard bus functional models (BFMs) driving the PCI Express link interface. Intel performs the following tests in the simulation environment:

- Directed and pseudorandom stimuli test the Application Layer interface, Configuration Space, and all types and sizes of TLPs
- Error injection tests inject errors in the link, TLPs, and Data Link Layer Packets (DLLPs), and check for the proper responses
- PCI-SIG<sup>®</sup> Compliance Checklist tests that specifically test the items in the checklist
- Random tests that test a wide range of traffic patterns

Intel provides the following two example designs that you can leverage to test your PCBs and complete compliance base board testing (CBB testing) at PCI-SIG.

### Related Information

- [PCI SIG Gen3 x8 Merged Design - Stratix V](#)
- [PCI SIG Gen2 x8 Merged Design - Stratix V](#)

## Compatibility Testing Environment

Intel has performed significant hardware testing to ensure a reliable solution. In addition, Intel internally tests every release with motherboards and PCI Express switches from a variety of manufacturers. All PCI-SIG compliance tests are run with each IP core release.

## Resource Utilization

Because the PCIe protocol stack is implemented in hardened logic, it uses no core device resources (no ALMs and no embedded memory).

The Avalon-MM with DMA V-Series variants include an Avalon-MM DMA bridge implemented in soft logic that operates as a front end to the hardened protocol stack. The following table shows the typical expected device resource utilization for selected configurations using the current version of the Quartus Prime software targeting an V-Series device. With the exception of M20K memory blocks, the numbers of ALMs and logic registers are rounded up to the nearest 50.

**Table 1-7: Resource Utilization V-Series Avalon-MM DMA for PCI Express**

Data Rate, Number of Lanes, and Interface Width	ALMs	M20K Memory Blocks	Logic Registers
Gen2 x4 128	4300	29	5800
Gen2 x8 128	12700	19	22300
Gen3 x8 256	18000	47	31450

**Note:** Soft calibration of the transceiver module requires additional logic. The amount of logic required depends upon the configuration.

## V-Series Recommended Speed Grades

Altera recommends setting the Quartus Prime Analysis & Synthesis Settings **Optimization Technique** to **Speed** when the Application Layer clock frequency is 250 MHz. For information about optimizing synthesis, refer to *Setting Up and Running Analysis and Synthesis* in Quartus Prime Help. For more information about how to effect the **Optimization Technique** settings, refer to *Area and Timing Optimization* in volume 2 of the Quartus Prime Handbook.

**Table 1-8: Arria V Recommended Speed Grades for All Link Widths, Link Widths, and Application Layer Clock Frequencies**

Link Rate	Link Width	Interface Width	Application Clock Frequency (MHz)	Recommended Speed Grades
Gen1	×1	64 bits	62.5 <sup>(3)</sup> ,125	-4,-5,-6
	×2	64 bits	125	-4,-5,-6
	×4	64 bits	125	-4,-5,-6
	×8	128 bits	125	-4,-5,-6
Gen2	×1	64 bits	125	-4,-5
	×2	64 bits	125	-4,-5
	×4	128 bits	125	-4,-5

**Table 1-9: Arria V GZ Recommended Speed Grades for All Widths, Link Widths, and Application Layer Clock Frequencies**

Link Rate	Link Width	Interface Width	Application Clock Frequency (MHz)	Recommended Speed Grades
Gen1	x1	64 bits	62.5 <sup>(4)</sup> ,125	-1, -2, -3, -4
	x2	64 bits	125	-1, -2, -3, -4
	x4	64 bits	125	-1, -2, -3, -4
	x8	64 bits	250	-1, -2, -3
	x8	128 Bits	125	-1, -2, -3, -4

<sup>(3)</sup> This is a power-saving mode of operation<sup>(4)</sup> This is a power-saving mode of operation<sup>(5)</sup> The -4 speed grade is also possible for this configuration; however, it requires significant effort by the end user to close timing.

Link Rate	Link Width	Interface Width	Application Clock Frequency (MHz)	Recommended Speed Grades
Gen2	x1	64 bits	125	-1, -2, -3, -4
	x2	64 bits	125	-1, -2, -3, -4
	x4	64 bits	250	-1, -2, -3 <sup>(5)</sup>
	x4	128 bits	125	-1, -2, -3, -4
	x8	128 bits	250	-1, -2, -3 <sup>(5)</sup>
	x8	256 bits	125	-1, -2, -3, -4
Gen3	x1	64 bits	125	-1, -2, -3, -4
	x2	64 bits	250	-1, -2, -3, -4
	x2	128 bits	125	-1, -2, -3, -4
	x4	128 bits	250	-1, -2, -3 <sup>(5)</sup>
	x4	256 bits	125	-1, -2, -3, -4
	x8	256 bits	250	-1, -2, -3 <sup>(5)</sup>

**Table 1-10: Cyclone V Recommended Speed Grades for All Link Widths, Link Widths, and Application Layer Clock Frequencies**

The Gen2 data rate requires Cyclone V GT devices.

Link Rate	Link Width	Interface Width	Application Clock Frequency (MHz)	Recommended Speed Grades
Gen1	x1	64 bits	62.5 <sup>(6)</sup> , 125	-6, -7, -8
	x2	64 bits	125	-6, -7, -8
	x4	64 bits	125	-6, -7, -8

<sup>(6)</sup> This is a power-saving mode of operation

Link Rate	Link Width	Interface Width	Application Clock Frequency (MHz)	Recommended Speed Grades
Gen2	x1	64 bits	125	-7
	x2	64 bits	125	-7
	x4	128 bits	125	-7

**Table 1-11: Stratix V Recommended Speed Grades for All Widths, Link Widths, and Application Layer Clock Frequencies**

Link Rate	Link Width	Interface Width	Application Clock Frequency (MHz)	Recommended Speed Grades
Gen1	x1	64 bits	62.5 <sup>(7)</sup> , 125	-1, -2, -3, -4
	x2	64 bits	125	-1, -2, -3, -4
	x4	64 bits	125	-1, -2, -3, -4
	x8	64 bits	250	-1, -2, -3 <sup>(8)</sup>
	x8	128 Bits	125	-1, -2, -3, -4
Gen2	x1	64 bits	125	-1, -2, -3, -4
	x2	64 bits	125	-1, -2, -3, -4
	x4	64 bits	250	-1, -2, -3
	x4	128 bits	125	-1, -2, -3, -4
	x8	128 bits	250	-1, -2, -3
	x8	256 bits	125	-1, -2, -3, -4

<sup>(7)</sup> This is a power-saving mode of operation

<sup>(8)</sup> The -4 speed grade is also possible for this configuration; however, it requires significant effort by the end user to close timing.



Link Rate	Link Width	Interface Width	Application Clock Frequency (MHz)	Recommended Speed Grades
Gen3	x1	64 bits	125	-1, -2, -3, -4
	x2	64 bits	250	-1, -2, -3, -4
	x2	128 bits	125	-1, -2, -3, -4
	x4	128 bits	250	-1, -2, -3
	x4	256 bits	125	-1, -2, -3, -4
	x8	256 bits	250	-1, -2, -3

**Related Information**

- [Area and Timing Optimization](#)
- [Altera Software Installation and Licensing Manual](#)
- [Setting up and Running Analysis and Synthesis](#)

## Creating a Design for PCI Express

### Before you begin

Select the PCIe variant that best meets your design requirements.

- Is your design an Endpoint or Root Port?
- What Generation do you intend to implement?
- What link width do you intend to implement?
- What bandwidth does your application require?
- Does your design require Configuration via Protocol (CvP)?

**Note:** The following steps only provide a high-level overview of the design generation and simulation process. For more details, refer to the *Quick Start Guide* chapter.

1. Select parameters for that variant.
2. For Intel Arria 10 devices, you can use the new Example Design tab of the component GUI to generate a design that you specify. Then, you can simulate this example and also download it to an Intel Arria 10 FPGA Development Kit. Refer to the Intel Arria 10/Intel Cyclone® 10 GX PCI Express\* IP Core Quick Start Guide for details.
3. For all devices, you can simulate using an Intel-provided example design. All static PCI Express example designs are available under `<install_dir>/ip/altera/altera_pcie/altera_pcie_<dev>_ed/example_design/<dev>`. Alternatively, create a simulation model and use your own custom or third-party BFM. The Platform Designer Generate menu generates simulation models. Intel supports ModelSim\* - Intel FPGA Edition for all IP. The PCIe cores support the Aldec RivieraPro\*, Cadence NCSim\*, Mentor Graphics ModelSim, and Synopsys VCS\* and VCS-MX\* simulators.

The Intel testbench and Root Port or Endpoint BFM provide a simple method to do basic testing of the Application Layer logic that interfaces to the variation. However, the testbench and Root Port BFM are not intended to be a substitute for a full verification environment. To thoroughly test your application, Intel suggests that you obtain commercially available PCI Express verification IP and tools, or do your own extensive hardware testing, or both.

4. Compile your design using the Quartus Prime software. If the versions of your design and the Quartus Prime software you are running do not match, regenerate your PCIe design.
5. Download your design to an Intel development board or your own PCB. Click on the *All Development Kits* link below for a list of Intel's development boards.
6. Test the hardware. You can use Intel's Signal Tap Logic Analyzer or a third-party protocol analyzer to observe behavior.
7. Substitute your Application Layer logic for the Application Layer logic in Intel's testbench. Then repeat Steps 3–6. In Intel's testbenches, the PCIe core is typically called the DUT (device under test). The Application Layer logic is typically called APPS.

#### Related Information

- [Parameter Settings](#) on page 3-1
- [Getting Started with the Avalon-MM DMA](#) on page 2-1
- [All Development Kits](#)
- [Intel Wiki PCI Express](#)

For complete design examples and help creating new projects and specific functions, such as MSI or MSI-X related to PCI Express. Intel Applications engineers regularly update content and add new design examples. These examples help designers like you get more out of the Intel PCI Express IP core and may decrease your time-to-market. The design examples of the Intel Wiki page provide useful guidance for developing your own design. However, the content of the Intel Wiki is not guaranteed by Intel.



2018.07.31

UG-01154



Subscribe



Send Feedback

You can download the Platform Designer design example, `pcie_de_ep_dma_g3x8_integrated.qsys`, from the `<install_dir>/ip/altera/altera_pcie/altera_pcie_hip_256_avmm/example_design/<dev>` directory.

**Note:** This design example provides instructions for generating simulation and synthesis files, but does not generate all the files necessary to download the design to hardware. Refer to *AN 690: PCI Express Avalon-MM DMA Reference Design* and *AN 708: PCI Express DMA Reference Design Using External DDR3 Memory for Stratix V and Arria V GZ Devices* for reference designs that include all files necessary to download your design to an FPGA Development Kit.

The design example includes the following components:

## Avalon-MM DMA for PCI Express

This IP core includes highly efficient DMA Read and DMA Write modules. The DMA Read and Write modules effectively move large blocks of data between the PCI Express address domain and the Avalon-MM address domain using burst data transfers. Depending on the configuration you select, the DMA Read and DMA Write modules use either a 128- or 256-bit Avalon-MM datapath.

In addition to high performance data transfer, the DMA Read and DMA Write modules ensure that the requests on the PCI link adhere to the *PCI Express Base Specification, 3.0*. The DMA Read and Write engines also perform the following functions:

- Divide the original request into multiple requests to avoid crossing 4KByte boundaries.
- Divide the original request into multiple requests to ensure that the maximum payload size is equal to or smaller than the maximum payload size for write requests and maximum read request size for read requests.
- Supports out-of-order completions when the original request is divided into multiple requests to adhere to the read request size. The Read Completions can come back in any order. The Read DMA Avalon-MM master port supports out-of-order Completions by writing the Read Completions to the correct locations. The Read DMA Avalon-MM master port does not have an internal reordering buffer.

Using the DMA Read and DMA Write modules, you can specify descriptor entry table entries with large payloads.

## On-Chip Memory IP core

This IP core stores the DMA data. This memory has a 256-bit data width.

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

**ALTERA**  
now part of Intel

## Descriptor Controller

The Descriptor Controller manages the Read DMA and Write DMA modules. Host software programs the Descriptor Controller internal registers with the location of the descriptor table. The Descriptor Controller instructs the Read DMA module to copy the entire table to its internal FIFO. It then pushes the table entries to DMA Read or DMA Write modules to transfer data. The Descriptor Controller also sends DMA status upstream via an Avalon-MM TX slave port.

In this example design the Descriptor Controller parameter, **Instantiate internal descriptor controller**, is on. Consequently, the Descriptor Controller is integrated into the Avalon-MM DMA bridge as shown in the figure below. Embedding the Descriptor Controller in the Avalon-MM DMA bridge simplifies the design. If you plan to replace the Descriptor Controller IP core with your own implementation, do not turn on the **Instantiate internal descriptor controller** in the parameter editor when parameterizing the IP core.

The Descriptor Controller supports the following features:

- A single duplex channel.
- Minimum transfer size of one dword (4 bytes).
- Maximum transfer size of 1 M (1024 \* 1024) - 4 bytes.

**Note:** Although the Descriptor Controller supports a maximum transfer size of (1 M (1024 \* 1024) - 4 bytes), the on-chip memory in this design example is smaller. Consequently, this design example cannot handle the maximum transfer size.

- Endpoints, only.
- Provides status to host software by generating an MSI interrupt when the DMA transfer completes.

## Transceiver Reconfiguration Controller IP Core

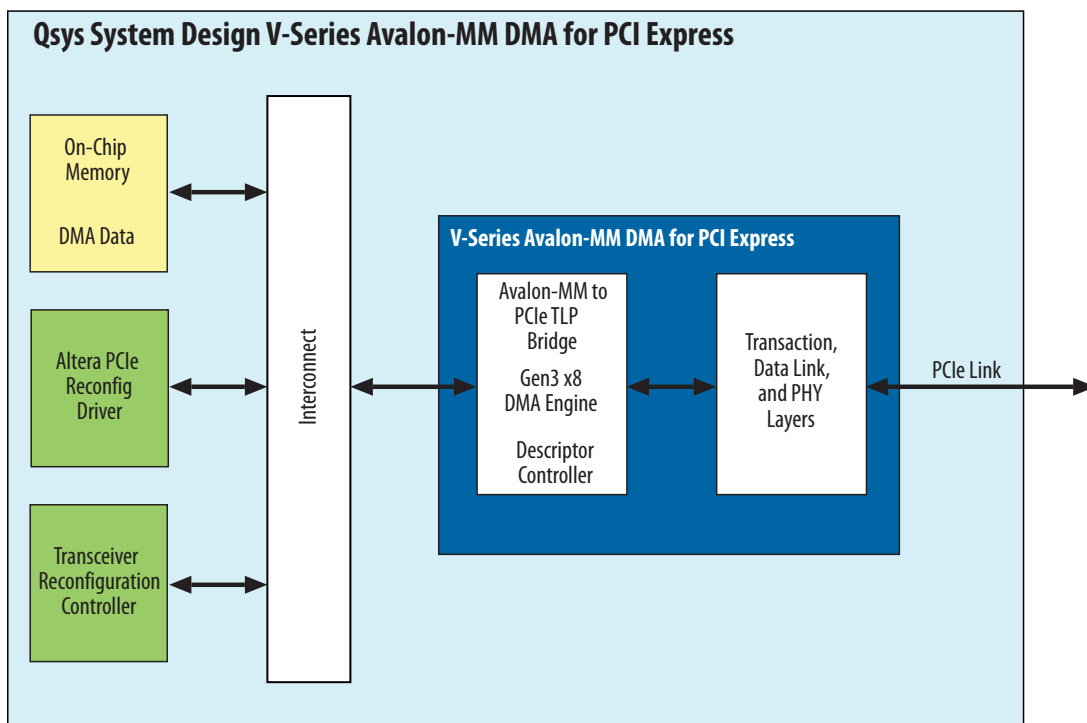
The Transceiver Reconfiguration Controller performs offset cancellation to compensate for variations due to process, voltage, and temperature (PVT).

The following provides a high-level block diagram of the V-Series Avalon-MM DMA for PCI Express Design Example.

## Intel PCIe Reconfig Driver IP Core

The PCIe Reconfig Driver drives the Transceiver Reconfiguration Controller. This driver is a plain text Verilog HDL file that you can modify if necessary to meet your system requirements.

## Block Diagram of the Avalon-MM DMA for PCI Express Example Design



### Design Example Limitations

This design example is intended to show basic DMA functionality. It is not a substitute for a robust verification testbench. If you modify this testbench, be sure to verify that the modifications result in the correct behavior.

#### Related Information

- [V-Series Avalon-MM DMA for PCI Express](#) on page 9-1
- [DMA Descriptor Controller Registers](#) on page 5-15
- [AN 690: PCI Express Avalon-MM DMA Reference Design](#)  
This reference design includes an Avalon-MM with DMA interface to the Application Layer. It illustrates chaining DMA performance using internal memory.
- [AN 708: PCI Express DMA Reference Design Using External DDR3 Memory for Stratix V and Arria V GZ Devices](#)  
This reference design includes an Avalon-MM with DMA interface to the Application Layer. It illustrates chaining DMA performance using external DDR3 memory.

## Understanding the Avalon-MM DMA Ports

The Avalon-MM DMA bridge includes ports to implement the DMA functionality. The following figure and table below illustrate and describe these ports.

Table 2-1: Avalon-MM DMA Platform Designer System Port Descriptions

Function	Port	Description
TXS	Txs	<p>This is an Avalon-MM slave port. In a typical application, an Avalon-MM master uses this port to send memory reads or writes to the PCIe domain.</p> <p>The Descriptor Controller uses it to write DMA status back to descriptor space in the PCIe domain when the DMA completes its operation. The Descriptor Controller also uses this port to send MSI interrupts upstream.</p>
Read Data Mover	dma_rd_master	<p>This is an Avalon-MM master port.</p> <p>The Read Data Mover moves data from the PCIe domain to the on-chip memory during normal read DMA operation. The Read Data Mover also fetches the descriptors from the PCIe domain and writes them to the FIFO in the Descriptor Controller. There are two separate descriptor tables for the read and write DMA descriptors. The <code>dma_rd_master</code> connects to <code>wr_dts_slave</code> port to load the write DMA descriptor FIFO and <code>rd_dts_slave</code> port to load the read DMA descriptor FIFO.</p>
Write Data Mover	dma_wr_master	<p>This is an Avalon-MM master port.</p> <p>The Write Data Mover reads data from the on-chip memory and then writes data to the PCIe domain.</p>
Descriptor Controller FIFOs	wr_dts_slave and rd_dts_slave	<p>This is an Avalon-MM slave port for the Descriptor Controller FIFOs. When the Read Data Mover fetches the descriptors from system memory, it writes the descriptors to the FIFO using this port. Because there are separate descriptor tables for read and write, there are two ports.</p> <p>The address range for the write DMA FIFO is 0x100_0000—0x100_1FFF.</p> <p>The address range for the read DMA FIFO is 0x100_2000—0x100_3FFF.</p>
Control in the Descriptor Controller	wr_dcm_master and rd_dcm_master	<p>The control block in the Descriptor Controller has one transmit and one receive port, one for read DMA and another one for write DMA. The receive port connects to the <code>RXM_BAR0</code> and the transmit port connects to the <code>Txs</code>.</p> <p>The receive path from the <code>RXM_BAR0</code> connects internally. It is not shown in the connections panel. For the transmit path, both read and write DMA ports connect to the <code>Txs</code> externally as shown in the connections panel.</p>

Function	Port	Description
RXM_BAR0	not shown in connections panel	This is an Avalon-MM master port. It passes the memory access from PCIe host to PCIe BAR0. The host uses this port to program the Descriptor Controller. Because this Platform Designer system uses an internal descriptor controller, the port connection is not shown in Platform Designer. The connection is inside the <code>a10_pcie_hip_0</code> module.
RXM_BAR4	Rxm_BAR4	<p>This is an Avalon-MM master port. It passes the memory access from PCIe host to PCIe BAR4. In the Platform Designer system, it connects to the on-chip memory. The PCIe host accesses the memory through PCIe BAR4</p> <p>In a typical application, system software controls this port to initialize random data in the on-chip memory. Software also reads the data back to verify correct operation.</p>

## Generating the Testbench

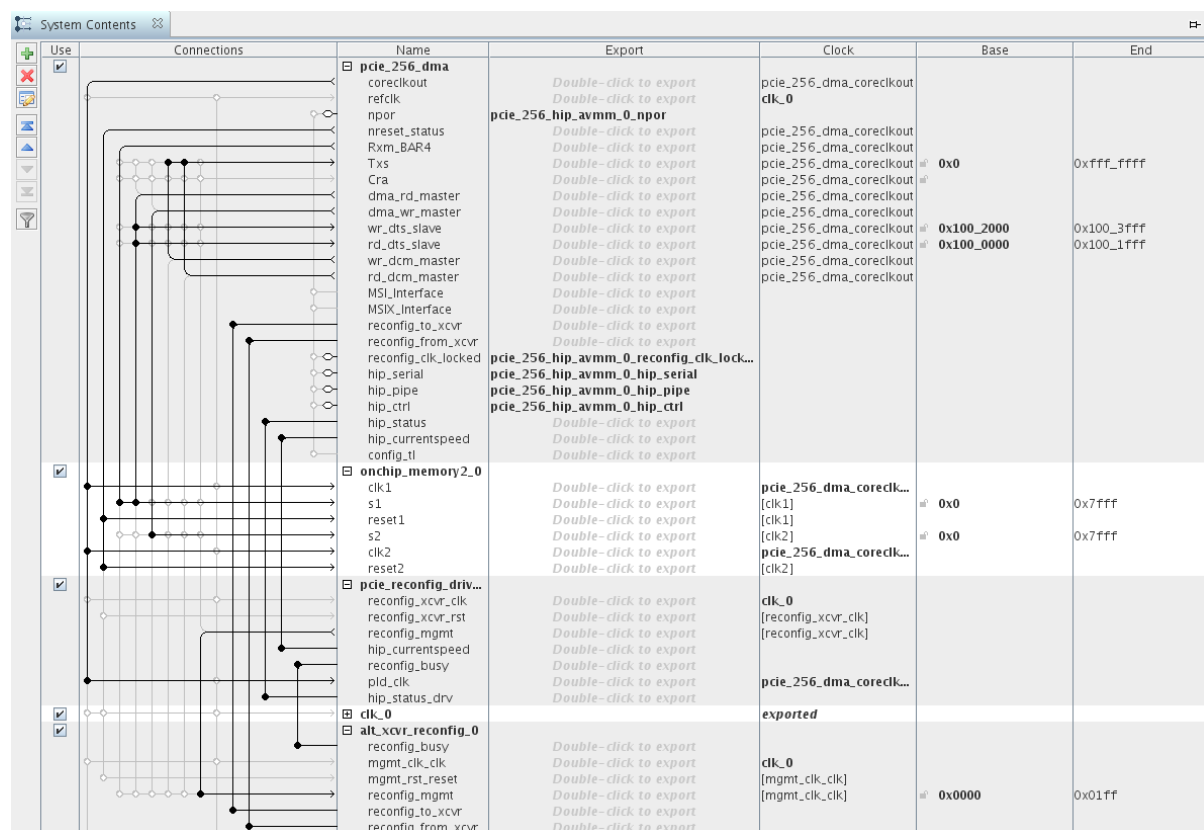
1. Copy the example design, **pcie\_de\_ep\_dma\_g3x8\_integrated.qsys**, from the installation directory:  
`<install_dir>/ip/altera/altera_pcie/altera_pcie_hip_256_avmm/example_design/` to your working directory.
2. Start Platform Designer, by typing the following command:

```
qsys-edit
```

3. Open `pcie_de_ep_dma_g3x8_integrated.qsys`.

The following figure shows the Platform Designer system.

Figure 2-1: V-Series Avalon-MM DMA for PCI Express Platform Designer System Design



4. Click **Generate > Generate Testbench System**.

5. Specify the following parameters:

Table 2-2: Parameters to Specify in the Generation Dialog Box

Parameter	Value
<b>Testbench System</b>	
Create testbench Platform Designer system	Standard, BFM's for standard Platform Designer interfaces
Create testbench simulation model	Verilog
Allow mixed-language simulation	You can leave this option off.
<b>Output Directory</b>	
Path	<working_dir>//pcie_de_ep_dma_g3x8_integrated

6. Click **Generate**.



## Understanding the Simulation Generated Files

**Table 2-3: Platform Designer Generation Output Files**

Directory	Description
<code>&lt;testbench_dir&gt;/&lt;variant_name&gt;/testbench</code>	Includes testbench subdirectories for the Aldec, Cadence, Mentor, and Synopsys simulation tools with the required libraries and simulation scripts.
<code>&lt;testbench_dir&gt;/&lt;variant_name&gt;/testbench/&lt;cad_vendor&gt;</code>	Includes the HDL source files and scripts for the simulation testbench.
<code>&lt;testbench_dir&gt;/&lt;variant_name&gt;/testbench/&lt;variant_namer&gt;_tb</code>	Includes HDL design files

## Understanding Simulation Log File Generation

Starting with the Quartus II 14.0 software release, simulation automatically creates a log file, `altpcie_monitor_<dev>_dlhip_tlp_file_log.log` in your simulation directory.

**Table 2-4: Sample Simulation Log File Entries**

Time	TLP Type	Payload (Bytes)	TLP Header
17989 RX	CfgRd0	0004	04000001_0000000F_01080008
17989 RX	MRd	0000	00000000_00000000_01080000
18021 RX	CfgRd0	0004	04000001_0000010F_0108002C
18053 RX	CfgRd0	0004	04000001_0000030F_0108003C
18085 RX	MRd	0000	00000000_00000000_0108000C

## Simulating the Example Design in ModelSim

1. In a terminal window, change directory to `<workingdir>/pcie_de_ep_dma_g3x8_integrated/testbench/mentor/`.
2. Start the ModelSim simulator.
3. To run the simulation, type the following commands in a terminal window:
  - a. `do msim_setup.tcl`
  - b. `ld_debug`  
The `ld_debug` command compiles all design files and elaborates the top-level design without any optimization.
  - c. `run -all`

The simulation performs the following operations:

- Various configuration accesses after the link is initialized
- Setup of the DMA controller to read data from the BFM's shared memory
- Setup of the DMA controller to write the same data back to the BFM's shared memory
- Data comparison and report of any mismatch

## Running a Gate-Level Simulation

The PCI Express testbenches run simulations at the register transfer level (RTL). However, it is possible to create your own gate-level simulations. Contact your Intel Sales Representative for instructions and an example that illustrates how to create a gate-level simulation from the RTL testbench.

## Generating Synthesis Files

1. On the **Generate** menu, select **Generate HDL**.
2. For **Create HDL design files for synthesis**, select **Verilog**.  
You can leave the default settings for all other items.
3. Click **Generate** to generate files for synthesis.
4. Click **Finish** when the generation completes.

## Compiling the Design

1. On the Quartus Prime Processing menu, click **Start Compilation**.
2. After compilation, expand the **TimeQuest Timing Analyzer** folder in the Compilation Report. Note whether the timing constraints are achieved in the Compilation Report.

If your design does not initially meet the timing constraints, you can find the optimal Fitter settings for your design by using the Design Space Explorer. To use the Design Space Explorer, click **Launch Design Space Explorer** on the Tools menu.

## Creating a Quartus Prime Project

You can create a new Quartus Prime project with the New Project Wizard, which helps you specify the working directory for the project, assign the project name, and designate the name of the top-level design entity.

1. On the Quartus Prime File menu, click then **New Project Wizard**, then **Next**.
2. Click **Next** in the **New Project Wizard: Introduction** (The introduction does not appear if you previously turned it off.)
3. On the **Directory, Name, Top-Level Entity** page, enter the following information:
  - a. For **What is the working directory for this project**, browse to `<project_dir>/pcie_de_ep_dma_g3x8_integrated/`.
  - b. For **What is the name of this project?** browse to the `<project_dir>/pcie_de_ep_dma_g3x8_integrated/synthesis` directory and select `pcie_de_ep_dma_g3x8_integrated.v`.
  - c. Click **Next**.
4. For **Project Type** select **Empty project**.

5. Click **Next**.
6. On the **Add Files** page, add `<project_dir>/pcie_de_ep_dma_g3x8_integrated/synthesis/ep_g3x8_avmm256_integrated.qip` to your Quartus Prime project.
7. Click **Next** to display the **Family & Device Settings** page.
8. On the **Device** page, choose the following target device family and options:
  - Note:** Currently, you cannot target an Intel Cyclone 10 GX Development Kit when generating an example design for Intel Cyclone 10 GX.
  - a. In the **Family** list, select **Stratix V (GS/GT/GX/E)**.
  - b. In the **Devices** list, select **Stratix V GX PCIe**.
  - c. In the **Available devices** list, select **5SGXEA7K2F40C2**.
9. Click **Next** to close this page and display the **EDA Tool Settings** page.
10. From the **Simulation** list, select **ModelSim**. From the **Format** list, select the HDL language you intend to use for simulation.
11. Click **Next** to display the **Summary** page.
12. Check the **Summary** page to ensure that you have entered all the information correctly.
13. Click **Finish**.
14. Save your project.

2018.07.31

UG-01154



Subscribe



Send Feedback

## System Settings

**Table 3-1: System Settings for PCI Express**

Parameter	Value	Description																								
Number of Lanes	x1, x2, x4, x8	Specifies the maximum number of lanes supported. Avalon-MM Interface with DMA does not support x1 configurations.																								
Lane Rate	Gen1 (2.5 Gbps) Gen2 (2.5/5.0 Gbps) Gen3 (2.5/5.0/8.0 Gbps)	Specifies the maximum data rate at which the link can operate.																								
Application interface width	128-bit 256-bit	<p>Specifies the width of the interface to the Application Layer. The following table indicates the possible combinations.</p> <table> <tr> <th>Application Interface Width</th><th>Configuration</th><th>Application Layer Clock Frequency</th></tr> <tr> <td>128 bits</td><td>Gen1 x8</td><td>125 MHz</td></tr> <tr> <td>128 bits</td><td>Gen2 x4</td><td>125 MHz</td></tr> <tr> <td>128 bits</td><td>Gen2 x8</td><td>250 MHz</td></tr> <tr> <td>256 bits</td><td>Gen2 x8</td><td>125 MHz</td></tr> <tr> <td>128 bits</td><td>Gen3 x4</td><td>250 MHz</td></tr> <tr> <td>256 bits</td><td>Gen3 x4</td><td>125 MHz</td></tr> <tr> <td>256 bits</td><td>Gen3 x8</td><td>250 MHz</td></tr> </table>	Application Interface Width	Configuration	Application Layer Clock Frequency	128 bits	Gen1 x8	125 MHz	128 bits	Gen2 x4	125 MHz	128 bits	Gen2 x8	250 MHz	256 bits	Gen2 x8	125 MHz	128 bits	Gen3 x4	250 MHz	256 bits	Gen3 x4	125 MHz	256 bits	Gen3 x8	250 MHz
Application Interface Width	Configuration	Application Layer Clock Frequency																								
128 bits	Gen1 x8	125 MHz																								
128 bits	Gen2 x4	125 MHz																								
128 bits	Gen2 x8	250 MHz																								
256 bits	Gen2 x8	125 MHz																								
128 bits	Gen3 x4	250 MHz																								
256 bits	Gen3 x4	125 MHz																								
256 bits	Gen3 x8	250 MHz																								

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

Parameter	Value	Description
<b>RX Buffer credit allocation - performance for received requests</b>	<b>Minimum</b> <b>Low</b> <b>Balanced</b>	<p>Determines the allocation of posted header credits, posted data credits, non-posted header credits, completion header credits, and completion data credits in the 16 KByte RX buffer. The settings allow you to adjust the credit allocation to optimize your system. The credit allocation for the selected setting displays in the message pane.</p> <p>Refer to the <i>Throughput Optimization</i> chapter in the <i>Stratix V Avalon-ST Interface for PCIe Solutions User Guide</i> for more information about optimizing performance.</p> <p>The <b>Message</b> window dynamically updates the number of credits for Posted, Non-Posted Headers and Data, and Completion Headers and Data as you change this selection.</p> <ul style="list-style-type: none"> <li>• <b>Minimum RX Buffer credit allocation -performance for received requests</b> )—configures the minimum PCIe specification allowed for non-posted and posted request credits, leaving most of the RX Buffer space for received completion header and data. Select this option for variations where application logic generates many read requests and only infrequently receives single requests from the PCIe link.</li> <li>• <b>Low</b>—configures a slightly larger amount of RX Buffer space for non-posted and posted request credits, but still dedicates most of the space for received completion header and data. Select this option for variations for which application logic generates many read requests and infrequently receives small bursts of requests from the PCIe link. This option is recommended for typical endpoint applications in which most of the PCIe traffic is generated by a DMA engine that is located in the endpoint application layer logic.</li> <li>• <b>Balanced</b>—configures approximately half the RX Buffer space to received requests and the other half of the RX Buffer space to received completions. Select this option for applications in which the received requests and received completions are roughly equal.</li> </ul>

Parameter	Value	Description
Reference clock frequency	100 MHz 125 MHz	<p>The <i>PCI Express Base Specification 3.0</i> requires a 100 MHz <math>\pm 300</math> ppm reference clock. The 125 MHz reference clock is provided as a convenience for systems that include a 125 MHz clock source. For more information about Gen3 operation, refer to 4.3.8 <i>Refclk Specifications for 8.0 GT/s</i> in the specification.</p> <p>For Gen3 operation, Altera recommends using a common reference clock (0 ppm) because when using separate reference clocks (non 0 ppm), the PCS occasionally must insert SKP symbols, potentially causes the PCIe link to go to recovery. Gen1 or Gen2 modes are not affected by this issue. Systems using the common reference clock (0 ppm) are not affected by this issue. The primary repercussion of this is a slight decrease in bandwidth. On Gen3 x8 systems, this bandwidth impact is negligible. If non 0 ppm mode is required, so that separate reference clocks are being used, please contact Altera for further information and guidance.</p>
Instantiate internal descriptor controller	On/Off	When you turn this option on, the descriptor controller is included in the Avalon-MM bridge. When you turn this option off, the descriptor controller should be included as a separate external component. Turn this option on, if you plan to use the Altera-provided descriptor controller in your design. Turn this option off if you plan to modify or replace the descriptor controller logic in your design.
Enable Avalon-MM CRA Slave hard IP status port	On/Off	Allows read and write access to bridge registers from the interconnect fabric using a specialized slave port. This option is required for <b>Requester/Completer</b> variants and optional for <b>Completer Only</b> variants. Enabling this option allows read and write access to bridge registers, except in the Completer-Only single dword variations.
Enable burst capabilities for RXM BAR2 port	On/Off	When you turn on this option, the BAR2 RX Avalon-MM masters is burst capable. If BAR2 is 32 bits and Burst capable, then BAR3 is not available for other use. If BAR2 is 64 bits, the BAR3 register holds the upper 32 bits of the address.
Enable configuration via the PCIe link	On/Off	<b>On</b> , the Quartus Prime software places the Endpoint in the location required for configuration via protocol (CvP). For more information about CvP, click the <i>Configuration via Protocol (CvP)</i> link below

Parameter	Value	Description
Use ATX PLL	On/Off	When you turn on this option, the Hard IP for PCI Express uses the ATX PLL instead of the CMU PLL. For other configurations, using the ATX PLL instead of the CMU PLL reduces the number of transceiver channels that are necessary. This option requires the use of the soft reset controller and does not support the CvP flow.
Enable Hard IP reset pulse at power-up when using the soft reset controller	On/Off	When you turn on this option, the soft reset controller generates a pulse at power up to reset the Hard IP. This pulse ensures that the Hard IP is reset after programming the device, regardless of the behavior of the dedicated PCI Express reset pin, <code>perstn</code> . This option is available for Gen2 and Gen3 designs that use a soft reset controller.

#### Related Information

##### [Stratix V Avalon-ST Interface for PCIe Solutions User Guide](#)

Explains how the **RX credit allocation** and the **Maximum payload RX Buffer credit allocation** and the **Maximum payload size** that you choose affect the allocation of flow control credits. You can set the **Maximum payload size** parameter on the **Device** tab.

## Base Address Register (BAR) Settings

The type and size of BARs available depend on port type.

Table 3-2: BAR Registers

Parameter	Value	Description
Type	<b>Disabled</b> <b>64-bit prefetchable memory</b> <b>32-bit non-prefetchable memory</b> <b>32-bit prefetchable memory</b> <b>I/O address space</b>	<p>If you select 64-bit prefetchable memory, 2 contiguous BARs are combined to form a 64-bit prefetchable BAR; you must set the higher numbered BAR to <b>Disabled</b>. A non-prefetchable 64-bit BAR is not supported because in a typical system, the Root Port Type 1 Configuration Space sets the maximum non-prefetchable memory window to 32 bits. The BARs can also be configured as separate 32-bit memories.</p> <p>Defining memory as prefetchable allows contiguous data to be fetched ahead. Prefetching memory is advantageous when the requestor may require more data from the same region than was originally requested. If you specify that a memory is prefetchable, it must have the following 2 attributes:</p> <ul style="list-style-type: none"><li>• Reads do not have side effects such as changing the value of the data read</li><li>• Write merging is allowed</li></ul>
Size	N/A	Platform Designer automatically calculates the required size after you connect your components.

## Device Identification Registers

Table 3-3: Device ID Registers

The following table lists the default values of the read-only Device ID registers. You can use the parameter editor to change the values of these registers. Refer to *Type 0 Configuration Space Registers* for the layout of the Device Identification registers.

Register Name	Range	Default Value	Description
<b>Vendor ID</b>	<b>16 bits</b>	0x00001172	Sets the read-only value of the <code>Vendor ID</code> register. This parameter cannot be set to 0xFFFF per the <i>PCI Express Specification</i> . Address offset: 0x000.
<b>Device ID</b>	16 bits	0x00000000	Sets the read-only value of the <code>Device ID</code> register. Address offset: 0x000.



Register Name	Range	Default Value	Description
<b>Revision ID</b>	8 bits	0x00000001	Sets the read-only value of the <code>Revision ID</code> register. Address offset: 0x008.
<b>Class code</b>	24 bits	0x00000000	Sets the read-only value of the <code>Class Code</code> register. Address offset: 0x008.
<b>Subsystem Vendor ID</b>	16 bits	0x00000000	Sets the read-only value of the <code>Subsystem Vendor ID</code> register in the PCI Type 0 Configuration Space. This parameter cannot be set to 0xFFFF per the <i>PCI Express Base Specification</i> . This value is assigned by PCI-SIG to the device manufacturer. Address offset: 0x02C.
<b>Subsystem Device ID</b>	16 bits	0x00000000	Sets the read-only value of the <code>Subsystem Device ID</code> register in the PCI Type 0 Configuration Space. Address offset: 0x02C

**Related Information**[PCI Express Base Specification 2.1 or 3.0](#)

## PCI Express and PCI Capabilities Parameters

This group of parameters defines various capability properties of the IP core. Some of these parameters are stored in the PCI Configuration Space - PCI Compatible Configuration Space. The byte offset indicates the parameter address.

### Device Capabilities

**Table 3-4: Capabilities Registers**

Parameter	Possible Values	Default Value	Description
<b>Maximum payload size</b>	128 bytes 256 bytes	128 bytes	Specifies the maximum payload size supported. This parameter sets the read-only value of the max payload size supported field of the Device Capabilities register (0x084[2:0]). Address: 0x084.

## Error Reporting

Table 3-5: Error Reporting

Parameter	Value	Default Value	Description
Advanced error reporting (AER)	On/Off	Off	When <b>On</b> , enables the Advanced Error Reporting (AER) capability.
ECRC checking	On/Off	Off	When <b>On</b> , enables ECRC checking. Sets the read-only value of the ECRC check capable bit in the Advanced Error Capabilities and Control Register. This parameter requires you to enable the AER capability.
ECRC generation	On/Off	Off	When <b>On</b> , enables ECRC generation capability. Sets the read-only value of the ECRC generation capable bit in the Advanced Error Capabilities and Control Register. This parameter requires you to enable the AER capability.
Enable ECRC forwarding on the Avalon-ST interface	On/Off	Off	When <b>On</b> , enables ECRC forwarding to the Application Layer. On the Avalon-ST RX path, the incoming TLP contains the ECRC dword <sup>(1)</sup> and the TD bit is set if an ECRC exists. On the transmit the TLP from the Application Layer must contain the ECRC dword and have the TD bit set.
Track RX completion buffer overflow on the Avalon-ST interface	On/Off	Off	When <b>On</b> , the core includes the <code>rxfx_cplbuf_ovf</code> output status signal to track the RX posted completion buffer overflow status

### Related Information

[PCI Express Base Specification Revision 2.1 or 3.0](#)

## Link Capabilities

Table 3-6: Link Capabilities

Parameter	Value	Description
Link port number	0x01	Sets the read-only value of the port number field in the Link Capabilities register.

Parameter	Value	Description
Slot clock configuration	On/Off	When <b>On</b> , indicates that the Endpoint or Root Port uses the same physical reference clock that the system provides on the connector. When <b>Off</b> , the IP core uses an independent clock regardless of the presence of a reference clock on the connector.

## MSI and MSI-X Capabilities

Table 3-7: MSI and MSI-X Capabilities

Parameter	Value	Description
MSI messages requested	1, 2, 4, 8, 16	Specifies the number of messages the Application Layer can request. Sets the value of the Multiple Message Capable field of the Message Control register.  Address: 0x050[31:16].
<b>MSI-X Capabilities</b>		
Implement MSI-X	On/Off	When <b>On</b> , adds the MSI-X functionality.
	Bit Range	
Table size	[10:0]	System software reads this field to determine the MSI-X Table size $\langle n \rangle$ , which is encoded as $\langle n-1 \rangle$ . For example, a returned value of 2047 indicates a table size of 2048. This field is read-only. Legal range is 0–2047 ( $2^{11}$ ).  Address offset: 0x068[26:16]
Table Offset	[31:0]	Points to the base of the MSI-X Table. The lower 3 bits of the table BAR indicator (BIR) are set to zero by software to form a 64-bit qword-aligned offset. This field is read-only.
Table BAR Indicator	[2:0]	Specifies which one of a function's BARs, located beginning at 0x10 in Configuration Space is used to map the MSI-X table into memory space. This field is read-only. Legal range is 0–5.
Pending Bit Array (PBA) Offset	[31:0]	Used as an offset from the address contained in one of the function's Base Address registers to point to the base of the MSI-X PBA. The lower 3 bits of the PBA BIR are set to zero by software to form a 32-bit qword-aligned offset. This field is read-only.

Parameter	Value	Description
<b>PBA BAR Indicator</b>	[2:0]	Specifies the function Base Address registers, located beginning at 0x10 in Configuration Space, that maps the MSI-X PBA into memory space. This field is read-only. Legal range is 0–5.

**Related Information**[PCI Express Base Specification Revision 2.1 or 3.0](#)

## Power Management

**Table 3-8: Power Management Parameters**

Parameter	Value	Description
<b>Endpoint L0s acceptable latency</b>	<b>Maximum of 64 ns</b> <b>Maximum of 128 ns</b> <b>Maximum of 256 ns</b> <b>Maximum of 512 ns</b> <b>Maximum of 1 us</b> <b>Maximum of 2 us</b> <b>Maximum of 4 us</b> <b>No limit</b>	<p>This design parameter specifies the maximum acceptable latency that the device can tolerate to exit the L0s state for any links between the device and the root complex. It sets the read-only value of the Endpoint L0s acceptable latency field of the Device Capabilities Register (0x084).</p> <p>This Endpoint does not support the L0s or L1 states. However, in a switched system there may be links connected to switches that have L0s and L1 enabled. This parameter is set to allow system configuration software to read the acceptable latencies for all devices in the system and the exit latencies for each link to determine which links can enable Active State Power Management (ASPM). This setting is disabled for Root Ports.</p> <p>The default value of this parameter is 64 ns. This is a safe setting for most designs.</p>



Parameter	Value	Description
<b>Endpoint L1 acceptable latency</b>	<b>Maximum of 1 us</b> <b>Maximum of 2 us</b> <b>Maximum of 4 us</b> <b>Maximum of 8 us</b> <b>Maximum of 16 us</b> <b>Maximum of 32 us</b> <b>Maximum of 64 ns</b> <b>No limit</b>	<p>This value indicates the acceptable latency that an Endpoint can withstand in the transition from the L1 to L0 state. It is an indirect measure of the Endpoint's internal buffering. It sets the read-only value of the Endpoint L1 acceptable latency field of the <code>Device Capabilities Register</code>.</p> <p>This Endpoint does not support the L0s or L1 states. However, a switched system may include links connected to switches that have L0s and L1 enabled. This parameter is set to allow system configuration software to read the acceptable latencies for all devices in the system and the exit latencies for each link to determine which links can enable Active State Power Management (ASPM). This setting is disabled for Root Ports.</p> <p>The default value of this parameter is 1 <math>\mu</math>s. This is a safe setting for most designs.</p>

## PCIe Address Space Settings

Table 3-9: PCIe Address Space Settings

Parameter	Value	Default Value	Description
<b>Address width of accessible PCIe Memory space</b>	<b>20–64</b>	<b>32</b>	Specifies the width of the TX Slave Module Avalon-MM address. This address is used unchanged as the PCIe address.

2018.07.31

UG-01154



Subscribe



Send Feedback

This chapter describes the top-level signals of the V-Series Hard IP for PCI Express using the Avalon-MM interface with DMA. The Avalon-MM DMA bridge includes high-performance, burst-capable Read DMA and Write DMA modules. The DMA Descriptor Controller that controls the Read DMA and Write DMA modules can be included in the Avalon-MM DMA bridge or separately instantiated. It uses 64-bit addressing, making address translation unnecessary. A separately instantiated Descriptor Controller manages the Read DMA and Write DMA modules. This variant is available for the following configurations:

- Gen1 x8
- Gen2 x4
- Gen2 x8
- Gen3 x4
- Gen3 x8

## V-Series DMA Avalon-MM DMA Interface to the Application Layer

This section describes the top-level interfaces in the PCIe variant when it includes the high-performance, burst-capable read data mover and write data mover modules.

Depending on the device, the interface to the Application Layer can be 128 or 256 bits.

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

Figure 4-1: Avalon-MM DMA Bridge with Internal Descriptor Controller

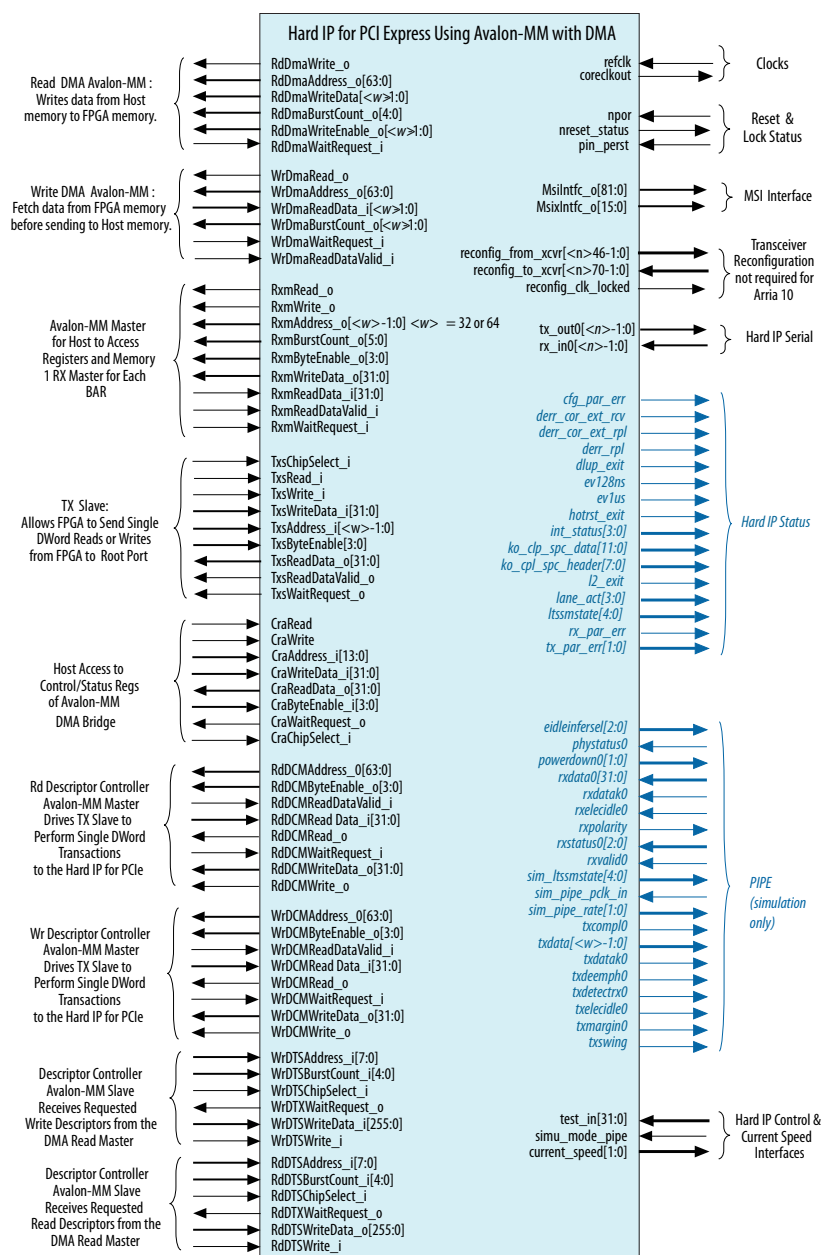
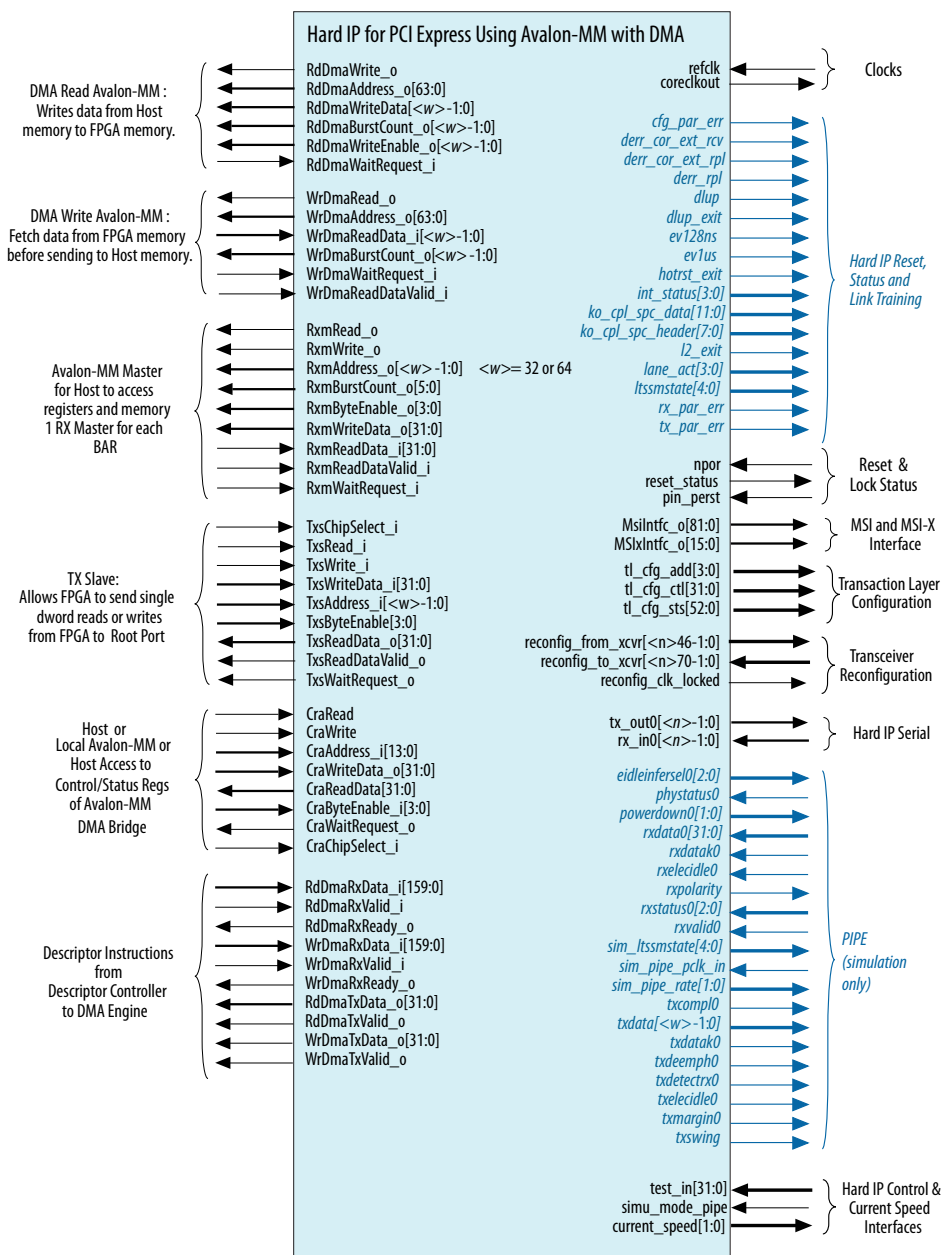


Figure 4-2: Avalon-MM DMA Bridge with Internal Descriptor Controller

Figure 4-3: Avalon-MM DMA Bridge with External Descriptor Controller



This section describes the interfaces that are required to implement the DMA. All other interfaces are described in the next section, *Avalon-MM Interface to the Application Layer*.

## Avalon-MM DMA Interfaces when Descriptor Controller Is Internally Instantiated

This configuration results from selecting both **Enable Avalon-MM DMA** and **Instantiate internal descriptor controller** in the component GUI.

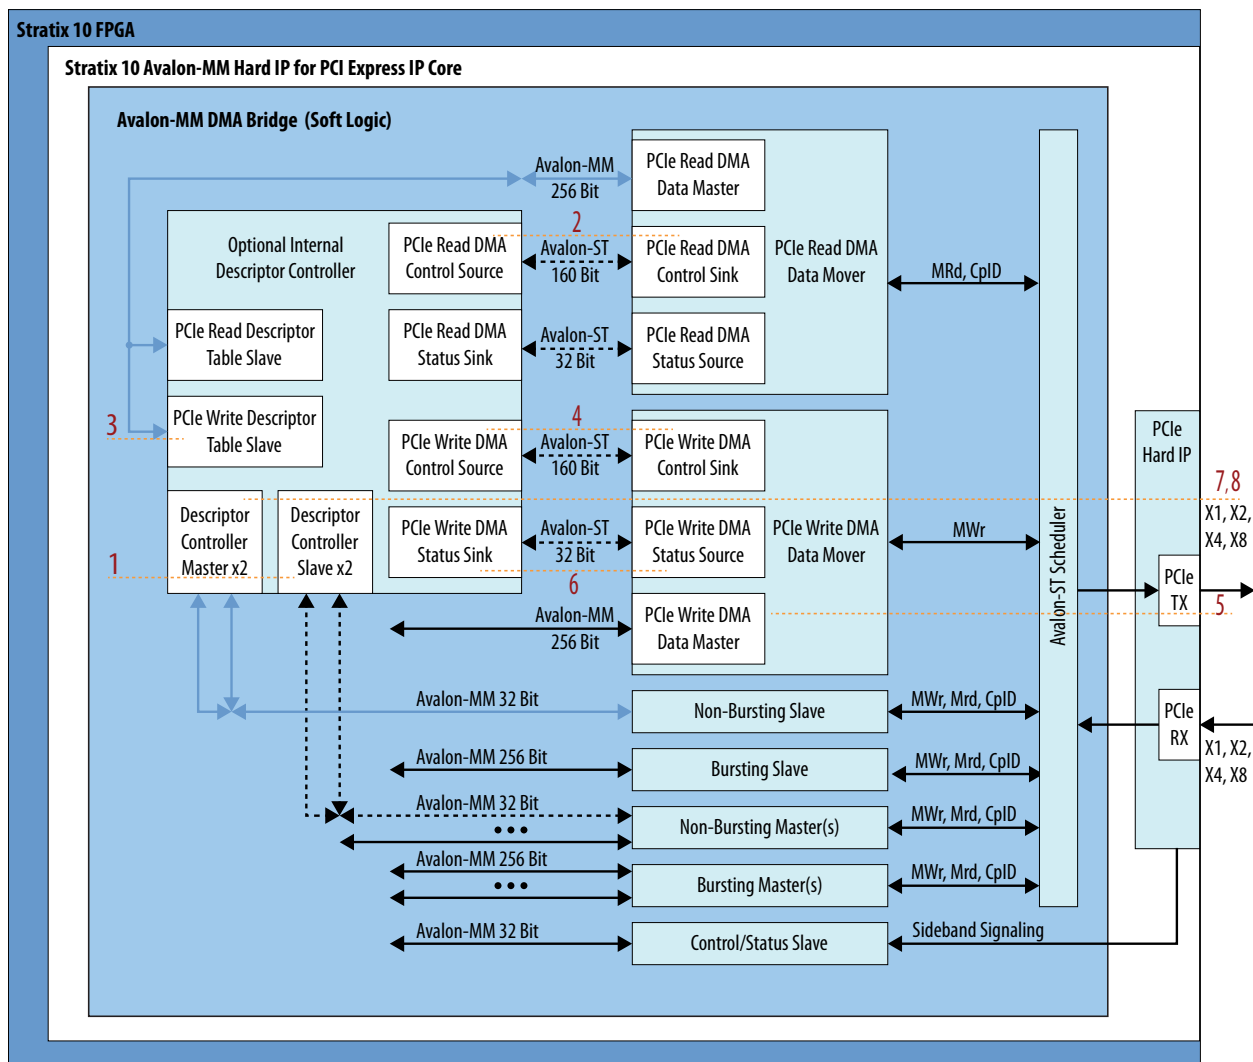


The following figure shows the Avalon-MM DMA Bridge, implemented in soft logic. It interfaces to the Hard IP for PCIe through Avalon-ST interfaces.

In the following figure, Avalon-ST connections and the connection from the BAR0 non-bursting master to the Descriptor Controller slaves are internal. Dashed black lines show these connections. Connections between the Descriptor Controller Masters and the non-bursting slave and the connections between the Read DMA Data Master and the Descriptor Table Slaves are made in the Platform Designer. Blue lines show these connections.

**Note:** In the following diagrams and text descriptions, the terms Read and Write are from the system memory perspective. Thus, a Read transaction reads data from the system memory and writes it to the local memory in Avalon-MM address space. A Write transaction writes the data that was read from the local memory in Avalon-MM address space to the system memory.

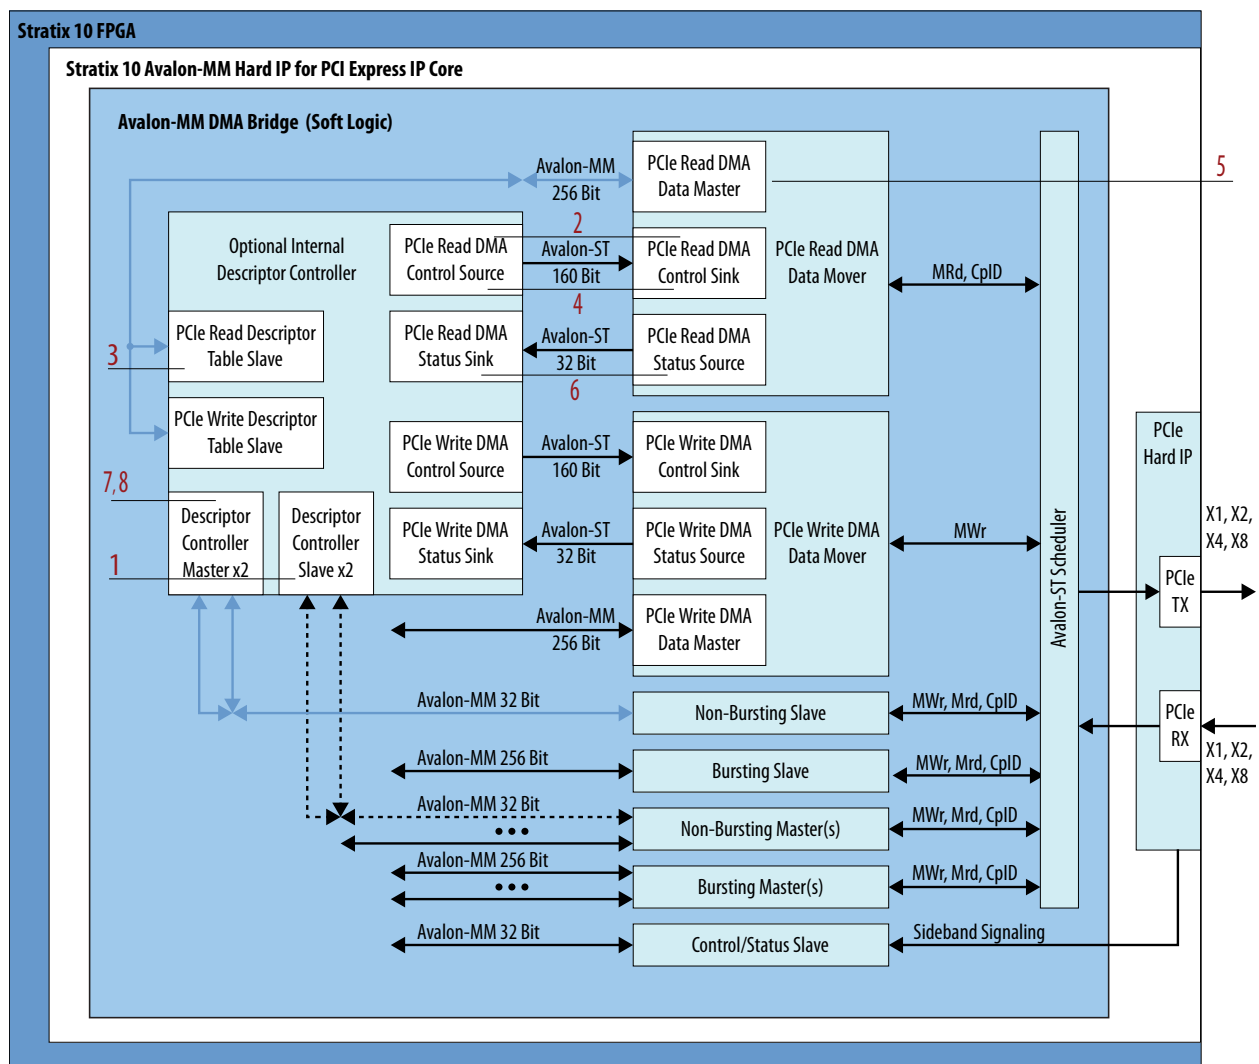
**Figure 4-4: Avalon-MM DMA Bridge Block Diagram with Optional Internal Descriptor Controller (Showing DMA Write Flow)**



The numbers in this figure describe the following steps in the DMA write flow:

1. The CPU writes registers in the Descriptor Controller Slave to start the DMA.
2. The Descriptor Controller instructs the Read Data Mover to fetch the descriptor table.
3. The Read Data Mover forwards the descriptor table to the PCIe Write Descriptor Table Slave.
4. The Descriptor Controller instructs the Write Data Mover to transfer data.
5. The Write Data Mover transfers data from FPGA to system memory.
6. The Write Data Mover notifies the Descriptor Controller of the completion of the data transfer using the done bit.
7. The Descriptor Controller Master updates the status of the descriptor table in system memory.
8. The Descriptor Controller Master sends an MSI interrupt to the host.

**Figure 4-5: Avalon-MM DMA Bridge Block Diagram with Optional Internal Descriptor Controller (Showing DMA Read Flow)**



The numbers in this figure describe the following steps in the DMA read flow:

1. The CPU writes registers in the Descriptor Controller Slave to start the DMA.
2. The Descriptor Controller instructs the Read Data Mover to fetch the descriptor table.
3. The Read Data Mover forwards the descriptor table to the PCIe Read Descriptor Table Slave.
4. The Descriptor Controller instructs the Read Data Mover to transfer data.
5. The Read Data Mover transfers data from system memory to FPGA.
6. The Read Data Mover notifies the Descriptor Controller of the completion of the data transfer using the `done` bit.
7. The Descriptor Controller Master updates the status of the descriptor table in system memory.
8. The Descriptor Controller Master sends an MSI interrupt to the host.

When the optional Descriptor Controller is included in the bridge, the Avalon-MM bridge includes the following Avalon interfaces to implement the DMA functionality:

- **PCIe Read DMA Data Master (`rd_dma`):** This is a 256-bit wide write only Avalon-MM master interface which supports bursts of up to 16 cycles with the `rd_dma*` prefix. The Read Data Mover uses this interface to write at high throughput the blocks of data that it has read from the PCIe system memory space. This interface writes descriptors to the Read and Write Descriptor table slaves and to any other Avalon-MM connected slaves interfaces.
- **PCIe Write DMA Data Master (`wr_dma`):** This read-only interface transfers blocks of data from the Avalon-MM domain to the PCIe system memory space at high throughput. It drives read transactions on its bursting Avalon-MM master interface. It also creates PCIe Memory Write (MWr) TLPs with data payload from Avalon-MM reads. It forwards the MWr TLPs to the Hard IP for transmission on the link. The Write Data Mover module decomposes the transfers into the required number of Avalon-MM burst read transactions and PCIe MWr TLPs. This is a bursting, 256-bit Avalon-MM interface with the `wr_dma` prefix.
- **PCIe Read Descriptor Table Slave (`rd_dts`):** This is a 256-bit Avalon-MM slave interface that supports write bursts of up to 16 cycles. The PCIe Read DMA Data Master writes descriptors to this table. This connection is made outside the DMA bridge because the Read Data Mover also typically connects to other Avalon-MM slaves. The prefix for this interface is `rd_dts`.
- **PCIe Write Descriptor Table Slave (`wr_dts`):** This is a 256-bit Avalon-MM slave interface that supports write bursts of up to 16 cycles. The PCIe Read DMA Data Master writes descriptors to this table. The PCIe Read DMA Data Master must connect to this interface outside the DMA bridge because the bursting master interface may also need to be connected to the destination of the PCIe Read Data Mover. The prefix for this interface is `wr_dts`.
- **Descriptor Controller Master (DCM):** This is a 32-bit, non-bursting Avalon-MM master interface with write-only capability. It controls the non-bursting Avalon-MM slave that transmits single DWORD DMA status information to the host. The prefixes for this interface are `wr_dcm` and `rd_dcm`.
- **Descriptor Controller Slave (DCS):** This is a 32-bit, non-bursting Avalon-MM slave interface with read and write access. The host accesses this interface through the BAR0 Non-Bursting Avalon-MM Master, to program the Descriptor Controller.

**Note:** This is not a top-level interface of the Avalon-MM Bridge. Because it connects to BAR0, you cannot use BAR0 to access any other Avalon-MM slave interface.

## Read Data Mover

The Read Data module sends memory read TLPs. It writes the completion data to an external Avalon-MM interface through the high throughput Read Master port. This data mover operates on descriptors the IP core receives from the DMA Descriptor Controller.

The Read DMA Avalon-MM Master interface performs the following functions:

### 1. Provides the Descriptor Table to the Descriptor Controller

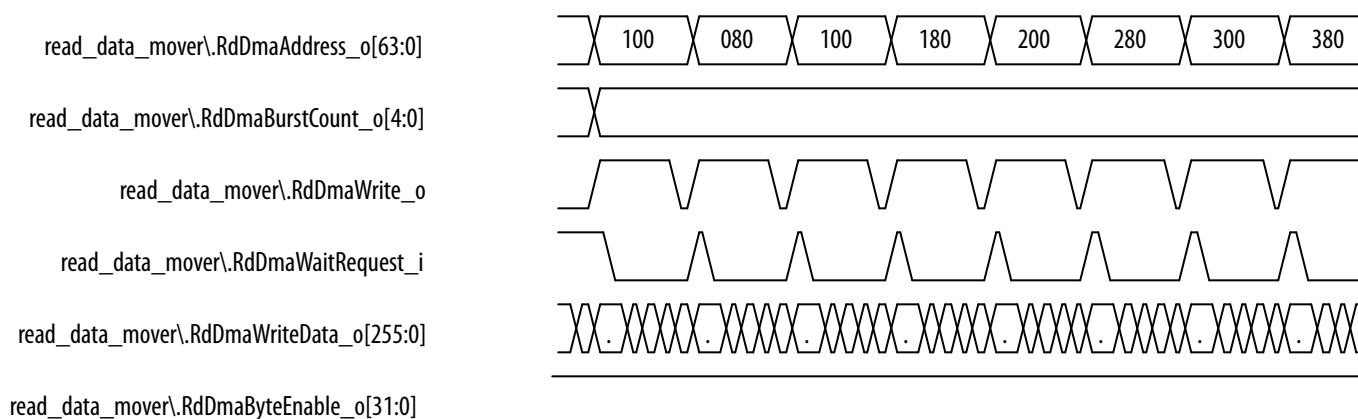
The Read Data Mover sends PCIe system memory read requests to fetch the descriptor table from PCIe system memory. This module then writes the returned descriptor entries in to the Descriptor Controller FIFO using this Avalon-MM interface.

### 2. Writes Data to Memory Located in Avalon-MM Space

After a DMA Read finishes fetching data from the source address in PCIe system memory, the Read Data Mover module writes the data to the destination address in Avalon-MM address space via this interface.

**Table 4-1: Read DMA 256-Bit Avalon-MM Master Interface**

Signal Name	Direction	Description
RdDmaWrite_o	Output	When asserted, indicates that the Read DMA module is ready to write read completion data to a memory component in the Avalon-MM address space.
RdDmaAddress_o[63:0]	Output	Specifies the write address in the Avalon-MM address space for the read completion data.
RdDmaWriteData_o[127 or 255:0]	Output	The read completion data to be written to the Avalon-MM address space.
RdDmaBurstCount_o[4:0] or [5:0]	Output	Specifies the burst count in 128- or 256-bit words. This bus is 5 bits for the 256-bit interface. It is 6 bits for the 128-bit interface.
RdDmaByteEnable_o[15 or 31:0]	Output	Specifies which DWORDs are valid.
RdDmaWaitRequest_i	Input	When asserted, indicates that the memory is not ready to receive data.

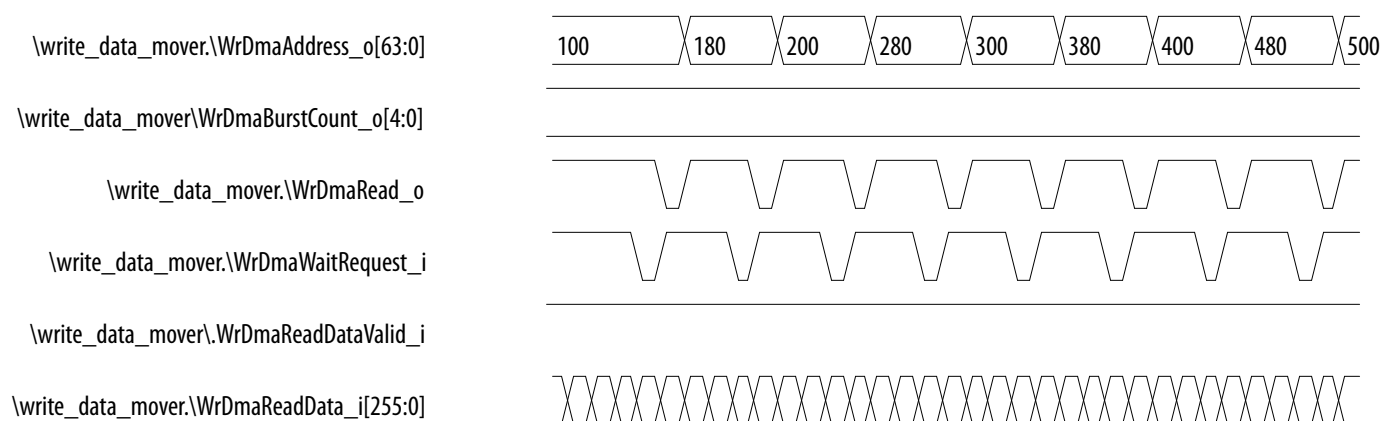
**Figure 4-6: Read DMA Avalon-MM Master Writes Data to FPGA Memory**

## Write DMA Avalon-MM Master Port

The Write Data Mover module fetches data from the Avalon-MM address space using this interface before issuing memory write requests to transfer data to PCIe system memory.

**Table 4-2: DMA Write 256-Bit Avalon-MM Master Interface**

Signal Name	Direction	Description
WrDMARead_o	Output	When asserted, indicates that the Write DMA module is reading data from a memory component in the Avalon-MM address space to write to the PCIe address space.
WrDmaAddress_o[63:0]	Output	Specifies the address for the data to be read from a memory component in the Avalon-MM address space .
WrDmaReadData_i[127 or 255:0]	Input	Specifies the completion data that the Write DMA module writes to the PCIe address space.
WrDmaBurstCount_o[4:0] or [5:0]	Output	Specifies the burst count in 128- or 256-bit words. This bus is 5 bits for the 256-bit interface. It is 6 bits for the 128-bit interface
WrDmaWaitRequest_i	Input	When asserted, indicates that the memory is not ready to be read.
WrDmaReadDataValid_i	Input	When asserted, indicates that WrDmaReadData_i is valid.

**Figure 4-7: Write DMA Avalon-MM Master Reads Data from FPGA Memory**

## RX Master Module

The RX Master module translates read and write TLPs received from the PCIe link to Avalon-MM requests for Platform Designer components connected to the interconnect. This module allows other PCIe components, including host software, to access other Avalon-MM slaves connected in the Platform Designer system.

If burst mode is not enabled, the RX Master module only supports 32-bit read or write request. All other requests received from the PCIe link are considered a violation of this device's programming model, and are therefore handled with the PCIe Completer Abort status. You can enable burst mode for BAR2 using 32-bit addressing or BAR2 and BAR3 using 64-bit addressing. When enabled, the module supports dword, burst read, or write requests. When the Descriptor Controller is internally instantiated, the RX Master for BAR0 is used internally and not available for other uses.

**Note:** When you enable burst mode for BAR2, Flush reads (i.e. reads with all byte enables set to 0) are not supported.

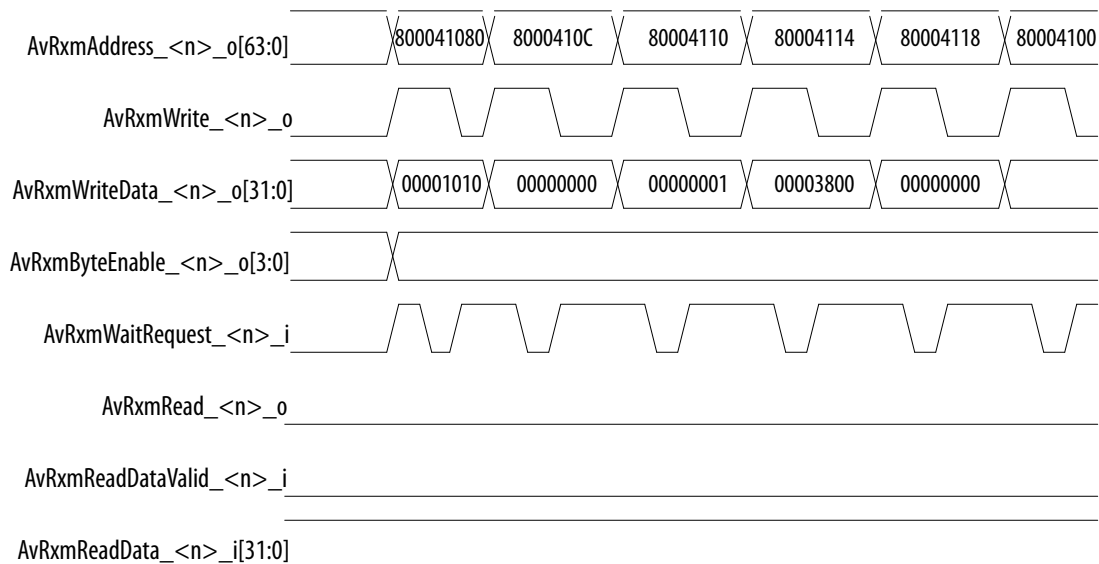
**Table 4-3: RX Master Control Interface Ports for BAR Access**

Each BAR has one corresponding RX Master Control interface. In this table,  $\langle n \rangle$  is the BAR number.

Signal Name	Direction	Description
RxmRead_ $\langle n \rangle$ _o	Output	When asserted, indicates an Avalon-MM read request.
RxmWrite_ $\langle n \rangle$ _o	Output	When asserted, indicates an Avalon-MM write request.
RxmAddress_ $\langle n \rangle$ _o[ $\langle w \rangle$ -1:0]	Output	Specifies the Avalon-MM byte address. Because all addresses are byte addresses, the meaningful bits of this address are [ $\langle w \rangle$ -1:2]. Bits 1 and 0 have a value of 0. $\langle w \rangle$ can be 32 or 64.
RxmBurstCount_ $\langle n \rangle$ _o[5:0]	Output	Specifies the burst count in dwords (32 bits). This optional signal is available for BAR2 when you turn on <b>Enable burst capabilities for RXM BAR2 ports</b> .

Signal Name	Direction	Description
RxmByteEnable_<n>_o[<w>-1:0]	Output	Specifies the valid bytes of data to be written. <w> has the following values: <ul style="list-style-type: none"> <li>4: for the non-bursting RX Master</li> <li>32: for the bursting 128-bit Avalon-MM interface</li> <li>64: for the bursting 256-bit Avalon-MM interface</li> </ul>
RxmDataWrite_<n>_o[<w>-1:0]	Output	Specifies the Avalon-MM write data. <w> has the following values: <ul style="list-style-type: none"> <li>32: for the non-bursting RX Master</li> <li>128: for the bursting 128-bit Avalon-MM interface</li> <li>256: for the bursting 256-bit Avalon-MM interface</li> </ul>
RxmReadData_<n>_i[<w>-1:0]	Input	Specifies the Avalon-MM read data. <w> has the following values: <ul style="list-style-type: none"> <li>32: for the non-bursting RX Master</li> <li>128: for the bursting 128-bit Avalon-MM interface</li> <li>256: for the bursting 256-bit Avalon-MM interface</li> </ul>
RxmReadDataValid_<n>_i[31:0]	Input	When asserted, indicates that RxmReadData_i[31:0] is valid.
RxmWaitRequest_<n>_i	Input	When asserted indicates that the control register access Avalon-MM slave port is not ready to respond.

Figure 4-8: RXM Master Writes to Memory in the Avalon-MM Address Space



## Non-Bursting Slave Module

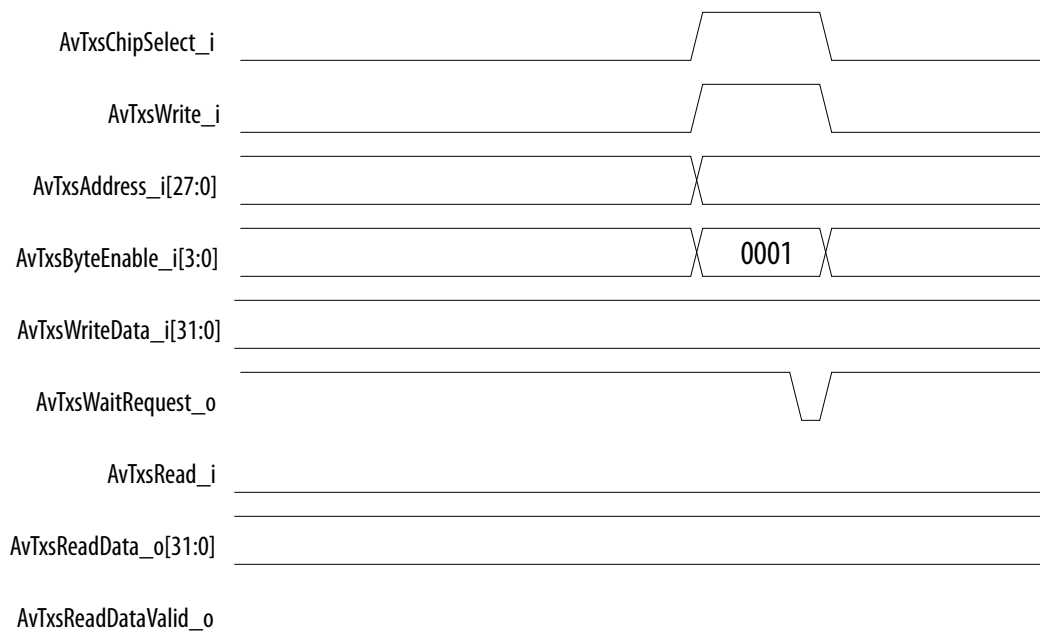
The TX Slave module translates Avalon-MM read and write requests to PCI Express TLPs.

The slave module supports a single outstanding non-bursting request. It typically sends status updates to the host. This is a 32-bit Avalon-MM slave interface.

Table 4-4: TX Slave Control

Signal Name	Direction	Description
TxsChipSelect_i	Input	When asserted, indicates that this slave interface is selected. When <code>txs_chipselect_i</code> is deasserted, <code>txs_read_i</code> and <code>txs_write_i</code> signals are ignored.
TxsRead_i	Input	When asserted, specifies a Avalon-MM Ignored when the chip select is deasserted.
TxsWrite_i	Input	When asserted, specifies a Avalon-MM Ignored when the chip select is deasserted.
TxsWriteData_i[31:0]	Input	Specifies the Avalon-MM data for a write command.
TxsAddress_i[<w>-1:0]	Input	Specifies the Avalon-MM byte address for the read or write command. The width of this address bus is specified by the parameter <b>Address width of accessible PCIe memory space</b> .
TxsByteEnable_i[3:0]	Input	Specifies the valid bytes for a write command.
TxsReadData_o[31:0]	Output	Drives the read completion data.
TxsReadDataValid_o	Output	When asserted, indicates that read data is valid.
TxsWaitRequest_o	Output	When asserted, indicates that the Avalon-MM slave port is not ready to respond to a read or write request.  The non-bursting Avalon-MM slave may assert <code>TxsWaitRequest_o</code> during idle cycles. An Avalon-MM master may initiate a transaction when <code>TxsWaitRequest_o</code> is asserted and wait for that signal to be deasserted.



**Figure 4-9: Non-Bursting Slave Interface Sends Status to Host**

## 32-Bit Control Register Access (CRA) Slave Signals

The CRA interface provides access to the control and status registers of the Avalon-MM bridge. This interface has the following properties:

- 32-bit data bus
- Supports a single transaction at a time
- Supports single-cycle transactions (no bursting)

**Note:** When Intel Stratix 10 is in Root Port mode, and the application logic issues a CfgWr or CfgRd via the CRA interface, it needs to fill the Tag field in the TLP Header with the value 0x10 to ensure that the corresponding Completion gets routed to the CRA interface correctly. If the application logic sets the Tag field to some other value, the Intel Stratix 10 Avalon-MM Hard IP for PCIe IP Core does not overwrite that value with the correct value.

**Table 4-5: Avalon-MM CRA Slave Interface Signals**

Signal Name	Direction	Description
CraRead_i	Input	Read enable.
CraWrite_i	Input	Write request.

Signal Name	Direction	Description
CraAddress_i[13:0]	Input	<p>An address space of 16 KB is allocated for the control registers. Avalon-MM slave addresses provide address resolution down to the width of the slave data bus. Because all addresses are byte addresses, this address logically goes down to bit 2. Bits 1 and 0 are 0. To read or write individual bytes of a DWORD, use byte enables. For example, to write bytes 0 and 1, set this signal to 4'b0011.</p> <p>An address space of 32 KB is allocated for the control registers. Avalon-MM slave addresses provide address resolution down to the width of the slave data bus. Because all addresses are byte addresses, this address logically goes down to bit 2. Bits 1 and 0 are 0. To read or write individual bytes of a DWORD, use byte enables. For example, to write bytes 0 and 1, set this signal to 4'b0011.</p>
CraWriteData_i[31:0]	Input	Write data. The current version of the CRA slave interface is read-only. Including this signal as part of the Avalon-MM interface, makes future enhancements possible.
CraReadData_o[31:0]	Output	Read data lines.
CraByteEnable_i[3:0]	Input	Byte enable.
CraWaitRequest_o	Output	Wait request to hold off additional requests.
CraChipSelect_i	Input	Chip select signal to this slave.
CraIrq_o	Output	Interrupt request. A port request for an Avalon-MM interrupt.

**Related Information**

[DMA Descriptor Controller Registers](#) on page 5-15

## Avalon-ST Descriptor Control Interface when Instantiated Separately

After fetching multiple descriptor entries from the Descriptor Table in the PCIe system memory, the Descriptor Controller uses its Avalon-ST Descriptor source interface to transfer 160-bit Descriptors to the Read or Write DMA Data Movers.

**Table 4-6: Avalon-ST Descriptor Sink Interface**

This interface sends instructions from Descriptor Controller to Read DMA Engine.

Signal Name	Direction	Description
RdAstRxData_i[159:0]	Input	Specifies the descriptors for the Read DMA module. Refer to <i>DMA Descriptor Format</i> table below for bit definitions.

Signal Name	Direction	Description
RdAstRxValid_i	Input	When asserted, indicates that the data is valid.
RdAstRxReady_o	Output	When asserted, indicates that the Read DMA read module is ready to receive a new descriptor.  The ready latency is 3 cycles. Consequently, interface can accept data 3 cycles after ready is asserted.

**Table 4-7: Avalon-ST Descriptor Sink Interface**

This interface sends instructions from Descriptor Controller to Write DMA Engine.

Signal Name	Direction	Description
WrAstRxData_i[159:0]	Input	Specifies the descriptors for the Write DMA module. Refer to <i>DMA Descriptor Format</i> table below for bit definitions.
WrAstRxValid_i	Input	When asserted, indicates that the data is valid.
WrAstRxReady_o	Output	When asserted, indicates that the Write DMA module engine is ready to receive a new descriptor. The ready latency for this signal is 3 cycles. Consequently, interface can accept data 3 cycles after ready is asserted.

## Avalon-ST Descriptor Status Interface

When DMA module completes the processing for one Descriptor Instruction, it returns DMA Status to the Descriptor Controller via the following interfaces.

**Table 4-8: Read DMA Status Interface from Read DMA Engine to Descriptor Controller**

Signal Name	Direction	Description
RdAstTxData_o[31:0]	Output	Drives status information to the Descriptor Controller component. Refer to <i>DMA Status Bus</i> table below for more information
RdAstTxValid_o	Output	When asserted, indicates that the data is valid.

**Table 4-9: Write DMA Status Interface from Write DMA Engine to Descriptor Controller**

Signal Name	Direction	Description
WrAstTxData_o[31:0]	Output	Drives status information to the Descriptor Controller component. Refer to <i>DMA Status Bus</i> table below for more information about this bus.

Signal Name	Direction	Description
WrAstTxValid_o	Output	When asserted, indicates that the data is valid.

Table 4-10: DMA Descriptor Format

Bits	Name	Description
[31:0]	Source Low Address	Low-order 32 bits of the DMA source address. The address boundary must align to the 32 bits so that the 2 least significant bits are 2'b00. For the Read DMA module, the source address is the PCIe domain address. For the Write DMA module, the source address is the Avalon-MM domain address.
[63:32]	Source High Address	High-order 32 bits of the source address.
[95:64]	Destination Low Address	Low-order 32 bits of the DMA destination address. The address boundary must align to the 32 bits so that the 2 least significant bits have the value of 2'b00. For the Read DMA module, the destination address is the Avalon-MM domain address. For the Write DMA module the destination address is the PCIe domain address.
[127:96]	Destination High Address	High-order 32 bits of the destination address.
[145:128]	DMA Length	Specifies DMA length in dwords. The length must be greater than 0. The maximum length is 1 MB - 4 bytes.
[153:146]	DMA Descriptor ID	Specifies up to 128 descriptors.
[158:154]	Reserved	—
[159]	Immediate Write	When set to 1'b1, the Write DMA Engine performs an Immediate Write. The Immediate Write provides a fast mechanism to send a Write TLP upstream. The descriptor stores the 32-bit payload, replacing the Source Low Address field of the descriptor.

## Avalon-ST Descriptor Status Sources

Read DMA and Write DMA modules report status to the Descriptor Controller on the RdDmaTx-Data\_o[31:0] or WrDmaTxData\_o[31:0] bus when a descriptor completes successfully.

The following table shows the mappings of the triggering events to the DMA descriptor status bus:

Table 4-11: DMA Status Bus

Bits	Name	Description
[31:9]	Reserved	—

Bits	Name	Description
[8]	Done	When asserted, a single DMA descriptor has completed successfully.
[7:0]	Descriptor ID	The ID of the descriptor whose status is being reported.

## Descriptor Controller Interfaces when Instantiated Internally

The Descriptor Controller controls the Read DMA and Write DMA Data Movers. It provides a 32-bit Avalon-MM slave interface to control and manage data flow from PCIe system memory to Avalon-MM memory and in the reverse direction.

The Descriptor Controller includes two, 128-entry FIFOs to store the read and write descriptor tables. The Descriptor Controller forwards the descriptors to the Read DMA and Write DMA Data Movers.

The Data Movers send completion status to the Read Descriptor Controller and Write Descriptor Controller. The Descriptor Controller forwards status and MSI to the host using the TX slave port.

### Read Descriptor Controller Avalon-MM Master interface

The Read Descriptor Controller Avalon-MM master interface drives the non-bursting Avalon-MM slave interface. The Read Descriptor Controller uses this interface to write descriptor status to the PCIe domain and possibly to MSI when MSI messages are enabled. This Avalon-MM master interface is only available for variants with the internally instantiated Descriptor Controller.

By default MSI interrupts are enabled. You specify the **Number of MSI messages requested** on the **MSI** tab of the parameter editor. The MSI Capability Structure is defined in *Section 6.8.1 MSI Capability Structure of the PCI Local Bus Specification*.

**Table 4-12: Read Descriptor Controller Avalon-MM Master interface**

Signal Name	Direction	Description
RdDCMAddress_o[63:0]	Output	Specifies the descriptor status table or MSI address.
RdDCMByteEnable_o[3:0]	Output	Specifies which data bytes are valid.
RdDCMReadDataValid_i	Input	When asserted, indicates that the read data is valid.
RdDCMReadData_i[31:0]	Input	Specifies the read data of the descriptor status table entry addressed.
RdDCMRead_o	Output	When asserted, indicates a read transaction. Currently, this is a write-only interface so that this signal never asserts.
RdDCMWaitRequest_i	Input	When asserted, indicates that the connected Avalon-MM slave interface is busy and cannot accept a transaction.
RdDCMWriteData_o[31:0]	Output	Specifies the descriptor status or MSI data..



Signal Name	Direction	Description
RdDCMWrite_o	Output	When asserted, indicates a write transaction.

## Write Descriptor Controller Avalon-MM Master Interface

The Avalon-MM Descriptor Controller Master interface is a 32-bit single-DWORD master with wait request support. The Write Descriptor Controller uses this interface to write status back to the PCI-Express domain and possibly MSI when MSI messages are enabled. This Avalon-MM master interface is only available for the internally instantiated Descriptor Controller.

By default MSI interrupts are enabled. You specify the **Number of MSI messages requested** on the **MSI** tab of the parameter editor. The MSI Capability Structure is defined in *Section 6.8.1 MSI Capability Structure of the PCI Local Bus Specification*.

**Table 4-13: Write Descriptor Controller Avalon-MM Master interface**

Signal Name	Direction	Description
WrDCMAddress_o[63:0]	Output	Specifies the descriptor status table or MSI address.
WrDCMByteEnable_o[3:0]	Output	Specifies which data bytes are valid.
WrDCMReadDataValid_i	Input	When asserted, indicates that the read data is valid.
WrRdDCMReadData_i[31:0]	Output	Specifies the read data for the descriptor status table entry addressed.
WrDCMRead_o	Output	When asserted, indicates a read transaction.
WrDCMWaitRequest_i	Input	When asserted, indicates that the Avalon-MM slave device is not ready to respond.
WrDCMWriteData_o[31:0]	Output	Specifies the descriptor status table or MSI address.
WrDCMWrite_o	Output	When asserted, indicates a write transaction.

## Read Descriptor Table Avalon-MM Slave Interface

This interface is available when you select the internal Descriptor Controller. It receives the Read DMA descriptors which are fetched by the Read Data Mover. Connect the interface to the Read DMA Avalon-MM master interface.

**Table 4-14: Read Descriptor Table Avalon-MM Slave Interface**

Signal Name	Direction	Description
RdDTSAddress_i[7:0]	Input	Specifies the descriptor table address.

Signal Name	Direction	Description
RdDTSBurstCount_i[4:0] or [5:0]	Input	Specifies the burst count of the transaction in words.
RdDTSChipSelect_i	Input	When asserted, indicates that the read targets this slave interface.
RdDTSWriteData_i[255:0] or [127:0]	Input	Specifies the descriptor.
RdDTSWrite_i	Input	When asserted, indicates a write transaction.
RdDTSWaitRequest_o	Output	When asserted, indicates that the Avalon-MM slave device is not ready to respond.

### Write Descriptor Table Avalon-MM Slave Interface

This interface is available when you select the internal Descriptor Controller. This interface receives the Write DMA descriptors which are fetched by Read Data Mover. Connect the interface to the Read DMA Avalon-MM master interface.

**Table 4-15: Write Descriptor Table Avalon-MM Slave Interface**

Signal Name	Direction	Description
WrDTSAddress_i[7:0]	Input	Specifies the descriptor table address.
WrDTSBurstCount_i[4:0]	Input	Specifies the burst count of the transaction in words.
WrDTSChipSelect_i	Input	When asserted, indicates that the write is for this slave interface.
WrDTSWaitRequest_o	Output	When asserted, indicates that this interface is busy and is not ready to respond.
WrDTSWriteData_i[255:0] or [127:0]	Input	Drives the descriptor table entry data.
WrDTSWrite_i	Input	When asserted, indicates a write transaction.

## Clock Signals

**Table 4-16: Clock Signals**

Signal	Direction	Description
refclk	Input	<p>Reference clock for the IP core. It must have the frequency specified under the <b>System Settings</b> heading in the parameter editor. This is a dedicated free running input clock to the dedicated REFCLK pin.</p> <p>If your design meets the following criteria:</p> <ul style="list-style-type: none"><li>• Enables CvP</li><li>• Includes an additional transceiver PHY connected to the same Transceiver Reconfiguration Controller</li></ul> <p>then you must connect refclk to the mgmt_clk_clk signal of the Transceiver Reconfiguration Controller and the additional transceiver PHY. In addition, if your design includes more than one Transceiver Reconfiguration Controller on the same side of the FPGA, they all must share the mgmt_clk_clk signal.</p>
coreclkout	Output	<p>This is a fixed frequency clock used by the Data Link and Transaction Layers. To meet PCI Express link bandwidth constraints, this clock has minimum frequency requirements as listed in <i>Application Layer Clock Frequency for All Combination of Link Width, Data Rate and Application Layer Interface Width</i> in the <i>Reset and Clocks</i> chapter .</p>

**Related Information**[Clocks](#) on page 6-5

## Reset, Status, and Link Training Signals

Refer to *Reset and Clocks* for more information about the reset sequence and a block diagram of the reset logic.



Table 4-17: Reset Signals

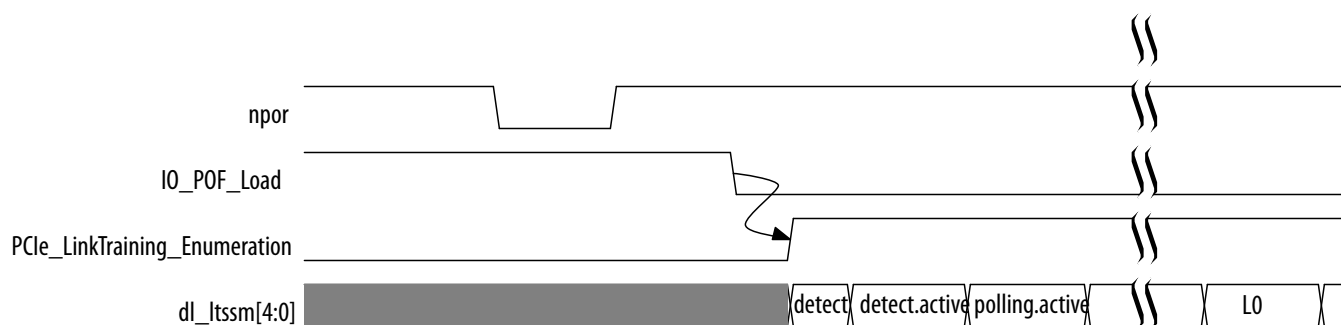
Signal	Direction	Description
<code>npor</code>	Input	<p>Active low reset signal. In the Intel hardware example designs, <code>npor</code> is the OR of <code>pin_perst</code> and <code>local_rstn</code> coming from the software Application Layer. If you do not drive a soft reset signal from the Application Layer, this signal must be derived from <code>pin_perst</code>. You cannot disable this signal. Resets the entire IP Core and transceiver. Asynchronous.</p> <p>In systems that use the hard reset controller, this signal is <i>edge, not level</i> sensitive; consequently, you cannot use a low value on this signal to hold custom logic in reset. For more information about the hard and soft reset controllers, refer to <i>Reset</i>.</p>
<code>nreset_status</code>	Output	<p>Active low reset signal. It is derived from <code>npor</code> or <code>pin_perstn</code>. You can use this signal to reset the Application Layer.</p>
<code>pin_perst</code>	Input	<p>Active low reset from the PCIe reset pin of the device. <code>pin_perst</code> resets the datapath and control registers. This signal is required for Configuration via Protocol (CvP). For more information about CvP refer to <i>Configuration via Protocol (CvP)</i>.</p> <p>V-Series have 1 or 2 instances of the Hard IP for PCI Express. Each instance has its own <code>pin_perst</code> signal. <i>You must connect the <code>pin_perst</code> of each Hard IP instance to the corresponding <code>nPERST</code> pin of the device.</i> These pins have the following locations:</p> <p>For Arria V devices:</p> <ul style="list-style-type: none"> <li><code>nPERSTL0</code>: bottom left Hard IP block and CvP blocks</li> <li><code>nPERSTL1</code>: top left Hard IP</li> </ul> <p>For Cyclone V devices:</p> <ul style="list-style-type: none"> <li><code>nPERSTL0</code>: top left Hard IP</li> <li><code>nPERSTL1</code>: bottom left Hard IP and CvP blocks</li> </ul> <p>For Arria V GZ and Stratix V devices:</p> <ul style="list-style-type: none"> <li><code>NPERSTL0</code>: bottom left Hard IP and CvP blocks</li> <li><code>NPERSTL1</code>: top left Hard IP block</li> <li><code>NPERSTR0</code>: bottom right Hard IP block</li> <li><code>NPERSTR1</code>: top right Hard IP block</li> </ul> <p>For maximum use of the V-Series device, Intel recommends that you use the bottom left Hard IP first. This is the only location that supports CvP over a PCIe link.</p> <p>Refer to the appropriate device pinout for correct pin assignment for more detailed information about these pins. The <i>PCI Express Card Electromechanical Specification 2.0</i> specifies this pin requires</p>



Signal	Direction	Description
		<p>3.3 V. You can drive this 3.3V signal to the <code>nPERST*</code> even if the <code>V<sub>VCCPGM</sub></code> of the bank is not 3.3V if the following 2 conditions are met:</p> <ul style="list-style-type: none"> <li>The input signal meets the <code>V<sub>IH</sub></code> and <code>V<sub>IL</sub></code> specification for LVTTTL.</li> </ul> <p>The input signal meets the overshoot specification for 100°C operation as specified by the <i>Maximum Allowed Overshoot and Undershoot Voltage</i> in the device datasheet.</p>

**Figure 4-10: Reset and Link Training Timing Relationships**

The following figure illustrates the timing relationship between `npor` and the LTSSM L0 state.



**Note:** To meet the 100 ms system configuration time, you must use the fast passive parallel configuration scheme with CvP and a 32-bit data width (FPP x32) or use the V-Series Hard IP for PCI Express in autonomous mode.

**Table 4-18: Status and Link Training Signals**

Signal	Direction	Description
<code>cfg_par_err</code>	Output	<p>Indicates that a parity error in a TLP routed to the internal Configuration Space. This error is also logged in the Vendor Specific Extended Capability internal error register. You must reset the Hard IP if this error occurs.</p> <p>The signal is not available for Arria V and Cyclone V devices.</p>

Signal	Direction	Description
derr_cor_ext_rcv	Output	Indicates a corrected error in the RX buffer. This signal is for debug only. It is not valid until the RX buffer is filled with data. This is a pulse, not a level, signal. Internally, the pulse is generated with the 500 MHz clock. A pulse extender extends the signal so that the FPGA fabric running at 250 MHz can capture it. Because the error was corrected by the IP core, no Application Layer intervention is required. <sup>(9)</sup>
derr_cor_ext_rpl	Output	Indicates a corrected ECC error in the retry buffer. This signal is for debug only. Because the error was corrected by the IP core, no Application Layer intervention is required. <sup>(9)</sup>
derr_rpl	Output	Indicates an uncorrectable error in the retry buffer. This signal is for debug only. <sup>(9)</sup>  The signal is not available for Arria V and Cyclone V devices.
dlup_exit	Output	This signal is asserted low for one <code>pld_clk</code> cycle when the IP core exits the DLCMSM DL_Up state, indicating that the Data Link Layer has lost communication with the other end of the PCIe link and left the Up state. When this pulse is asserted, the Application Layer should generate an internal reset signal that is asserted for at least 32 cycles.
ev128ns	Output	Asserted every 128 ns to create a time base aligned activity.
ev1us	Output	Asserted every 1µs to create a time base aligned activity.
hotrst_exit	Output	Hot reset exit. This signal is asserted for 1 clock cycle when the LTSSM exits the hot reset state. This signal should cause the Application Layer to be reset. This signal is active low. When this pulse is asserted, the Application Layer should generate an internal reset signal that is asserted for at least 32 cycles.
int_status[3:0]	Output	These signals drive legacy interrupts to the Application Layer as follows: <ul style="list-style-type: none"> <li>int_status[0]: interrupt signal A</li> <li>int_status[1]: interrupt signal B</li> <li>int_status[2]: interrupt signal C</li> <li>int_status[3]: interrupt signal D</li> </ul>

<sup>(9)</sup> Intel does not rigorously test or verify debug signals. Only use debug signals to observe behavior. Do not use debug signals to drive custom logic.

Signal	Direction	Description
ko_cpl_spc_data[11:0]	Output	The Application Layer can use this signal to build circuitry to prevent RX buffer overflow for completion data. Endpoints must advertise infinite space for completion data; however, RX buffer space is finite. ko_cpl_spc_data is a static signal that reflects the total number of 16 byte completion data units that can be stored in the completion RX buffer.
ko_cpl_spc_header[7:0]	Output	The Application Layer can use this signal to build circuitry to prevent RX buffer overflow for completion headers. Endpoints must advertise infinite space for completion headers; however, RX buffer space is finite. ko_cpl_spc_header is a static signal that indicates the total number of completion headers that can be stored in the RX buffer.
l2_exit	Output	L2 exit. This signal is active low and otherwise remains high. It is asserted for one cycle (changing value from 1 to 0 and back to 1) after the LTSSM transitions from l2.idle to detect. When this pulse is asserted, the Application Layer should generate an internal reset signal that is asserted for at least 32 cycles.
lane_act[3:0]	Output	Lane Active Mode: This signal indicates the number of lanes that configured during link training. The following encodings are defined: <ul style="list-style-type: none"> <li>• 4'b0001: 1 lane</li> <li>• 4'b0010: 2 lanes</li> <li>• 4'b0100: 4 lanes</li> <li>• 4'b1000: 8 lanes</li> </ul>
ltssmstate[4:0]	Output	LTSSM state: The LTSSM state machine encoding defines the following states: <ul style="list-style-type: none"> <li>• 00000: Detect.Quiet</li> <li>• 00001: Detect.Active</li> <li>• 00010: Polling.Active</li> <li>• 00011: Polling.Compliance</li> <li>• 00100: Polling.Configuration</li> <li>• 00101: Polling.Speed</li> <li>• 00110: Config.Linkwidthstart</li> <li>• 00111: Config.Linkaccept</li> <li>• 01000: Config.Lanenumaccept</li> <li>• 01001: Config.Lanenumwait</li> <li>• 01010: Config.Complete</li> <li>• 01011: Config.Idle</li> <li>• 01100: Recovery.Rcvlock</li> </ul>

Signal	Direction	Description
		<ul style="list-style-type: none"> <li>• 01101: Recovery.Rcvconfig</li> <li>• 01110: Recovery.Idle</li> <li>• 01111: L0</li> <li>• 10000: Disable</li> <li>• 10001: Loopback.Entry</li> <li>• 10010: Loopback.Active</li> <li>• 10011: Loopback.Exit</li> <li>• 10100: Hot.Reset</li> <li>• 10101: L0s</li> <li>• 11001: L2.transmit.Wake</li> <li>• 11010: Speed.Recovery</li> <li>• 11011: Recovery.Equalization, Phase 0</li> <li>• 11100: Recovery.Equalization, Phase 1</li> <li>• 11101: Recovery.Equalization, Phase 2</li> <li>• 11110: Recovery.Equalization, Phase 3</li> <li>• 11111: Recovery.Equalization, Done</li> </ul>
rx_par_err	Output	<p>When asserted for a single cycle, indicates that a parity error was detected in a TLP at the input of the RX buffer. This error is logged as an uncorrectable internal error in the VSEC registers. For more information, refer to <i>Uncorrectable Internal Error Status Register</i>. You must reset the Hard IP if this error occurs because parity errors can leave the Hard IP in an unknown state.</p> <p>The signal is not available for Arria V and Cyclone V devices.</p>
tx_par_err[1:0]	Output	<p>When asserted for a single cycle, indicates a parity error during TX TLP transmission. These errors are logged in the VSEC register. The following encodings are defined:</p> <ul style="list-style-type: none"> <li>• 2'b10: A parity error was detected by the TX Transaction Layer. The TLP is nullified and logged as an uncorrectable internal error in the VSEC registers. For more information, refer to <i>Uncorrectable Internal Error Status Register</i>.</li> <li>• 2'b01: Some time later, the parity error is detected by the TX Data Link Layer which drives 2'b01 to indicate the error. Reset the IP core when this error is detected. Contact Intel technical support if resetting becomes unworkable.</li> </ul> <p>Note that not all simulation models assert the Transaction Layer error bit in conjunction with the Data Link Layer error bit.</p> <p>The signal is not available for Arria V and Cyclone V devices.</p>



## Related Information

- [PCI Express Card Electromechanical Specification 2.0](#)
- [Configuration via Protocol \(CvP\) Implementation in V-Series FPGAs User Guide](#)

## MSI Interrupts for Endpoints

The MSI interrupt notifies the host when a DMA operation has completed. After the host receives this interrupt, it can poll the DMA read or write status table to determine which entry or entries have the `done` bit set. This mechanism allows host software to avoid continuous polling of the status table `done` bits. Use this interface to receive information required to generate MSI or MSI-X interrupts to the Root Port via the TX Slave interface.

**Table 4-19: MSI Interrupt**

Signal	Direction	Description
MSIIntfc_o[81:0]	Output	This bus provides the following MSI address, data, and enabled signals: <ul style="list-style-type: none"><li>• MSIIntfc_o[81]: Master enable</li><li>• MSIIntfc_o[80]: MSI enable</li><li>• MSIIntfc_o[79:64]: MSI data</li><li>• MSIIntfc_o[63:0]: MSI address</li></ul>
MSIXIntfc_o[15:0]	Output	Provides for system software control of MSI-X as defined in Section 6.8.2.3 <i>Message Control for MSI-X</i> in the <i>PCI Local Bus Specification, Rev. 3.0</i> . The following fields are defined: <ul style="list-style-type: none"><li>• MSIXIntfc_o[15]: Enable</li><li>• MSIXIntfc_o[14]: Mask</li><li>• MSIXIntfc_o[13:11]: Reserved</li><li>• MSIXIntfc_o[10:0]: Table size</li></ul>
MSIControl_o[15:0]	Output	Provides system software control of the MSI messages as defined in Section 6.8.1.3 <i>Message Control for MSI</i> in the <i>PCI Local Bus Specification, Rev. 3.0</i> . The following fields are defined: <ul style="list-style-type: none"><li>• MSIControl_o[15:9]: Reserved</li><li>• MSIControl_o[8]: Per-Vector Masking Capable</li><li>• MSIControl_o[7]: 64-Bit Address Capable</li><li>• MSIControl_o[6:4]: Multiple Message Enable</li><li>• MSIControl_o[3:1]: MSI Message Capable</li><li>• MSIControl_o[0]: MSI Enable</li></ul>

## Physical Layer Interface Signals

Intel provides an integrated solution with the Transaction, Data Link and Physical Layers. The IP Parameter Editor generates a SERDES variation file, `<variation>_serdes.v` or `.vhd`, in addition to the Hard IP variation file, `<variation>.v` or `.vhd`. The SERDES entity is included in the library files for PCI Express.

### Transceiver Reconfiguration

Dynamic reconfiguration compensates for variations due to process, voltage and temperature (PVT). Among the analog settings that you can reconfigure are  $V_{OD}$ , pre-emphasis, and equalization.

You can use the Intel Transceiver Reconfiguration Controller to dynamically reconfigure analog settings. For more information about instantiating the Intel Transceiver Reconfiguration Controller IP core refer to *Hard IP Reconfiguration*.

**Table 4-20: Transceiver Control Signals**

In this table,  $\langle n \rangle$  is the number of interfaces required.

Signal Name	Direction	Description
<code>reconfig_from_xcvr[<math>\langle n \rangle</math>46)-1:0]</code>	Output	Reconfiguration signals to the Transceiver Reconfiguration Controller.
<code>reconfig_to_xcvr[<math>\langle n \rangle</math>70)-1:0]</code>	Input	Reconfiguration signals from the Transceiver Reconfiguration Controller.
<code>reconfig_clk_locked</code>	Output	When asserted, indicates that the PLL that provides the fixed clock required for transceiver initialization is locked. The Application Layer should be held in reset until <code>reconfig_clk_locked</code> is asserted.

The following table shows the number of logical reconfiguration and physical interfaces required for various configurations. The Quartus Prime Fitter merges logical interfaces to minimize the number of physical interfaces configured in the hardware. Typically, one logical interface is required for each channel and one for each PLL. The  $\times 8$  variants require an extra channel for PCS clock routing and control. The  $\times 8$  variants use channel 4 for clocking.

**Table 4-21: Number of Logical and Physical Reconfiguration Interfaces**

Variant	Logical Interfaces
Gen2 $\times 4$	5
Gen1 and Gen2 $\times 8$	10
Gen3 $\times 4$	6

Variant	Logical Interfaces
Gen3 ×8	11

For more information about the Transceiver Reconfiguration Controller, refer to the *Transceiver Reconfiguration Controller* chapter in the *Intel Transceiver PHY IP Core User Guide*.

#### Related Information

[Intel Transceiver PHY IP Core User Guide](#)

## Serial Data Signals

This differential, serial interface is the physical link between a Root Port and an Endpoint.

The PCIe IP Core supports 1, 2, 4, or 8 lanes. Each lane includes a TX and RX differential pair. Data is striped across all available lanes.

The PCIe IP Core supports 1, 2, or 4 lanes. Each lane includes a TX and RX differential pair. Data is striped across all available lanes.

**Table 4-22: 1-Bit Interface Signals**

In the following table  $\langle n \rangle$  is the number of lanes.

Signal	Direction	Description
tx_out[ $\langle n \rangle - 1 : 0$ ]	Output	Transmit output. These signals are the serial outputs of lanes $\langle n \rangle - 1 - 0$ .
rx_in[ $\langle n \rangle - 1 : 0$ ]	Input	Receive input. These signals are the serial inputs of lanes $\langle n \rangle - 1 - 0$ .

Refer to *Pin-out Files for Intel Devices* for pin-out tables for all Intel devices in .pdf, .txt, and .xls formats.

Transceiver channels are arranged in groups of six. For GX devices, the lowest six channels on the left side of the device are labeled GXB\_L0, the next group is GXB\_L1, and so on. Channels on the right side of the device are labeled GXB\_R0, GXB\_R1, and so on. Be sure to connect the Hard IP for PCI Express on the left side of the device to appropriate channels on the left side of the device, as specified in the *Pin-out Files for Intel Devices*.

#### Related Information

[Pin-out Files for Intel Devices](#)

## Physical Layout of Hard IP In Arria V GX/GX/SX/ST Devices

Arria V devices include one or two Hard IP for PCI Express IP cores. The following figures illustrate the placement of the PCIe IP core, transceiver banks, and channels.

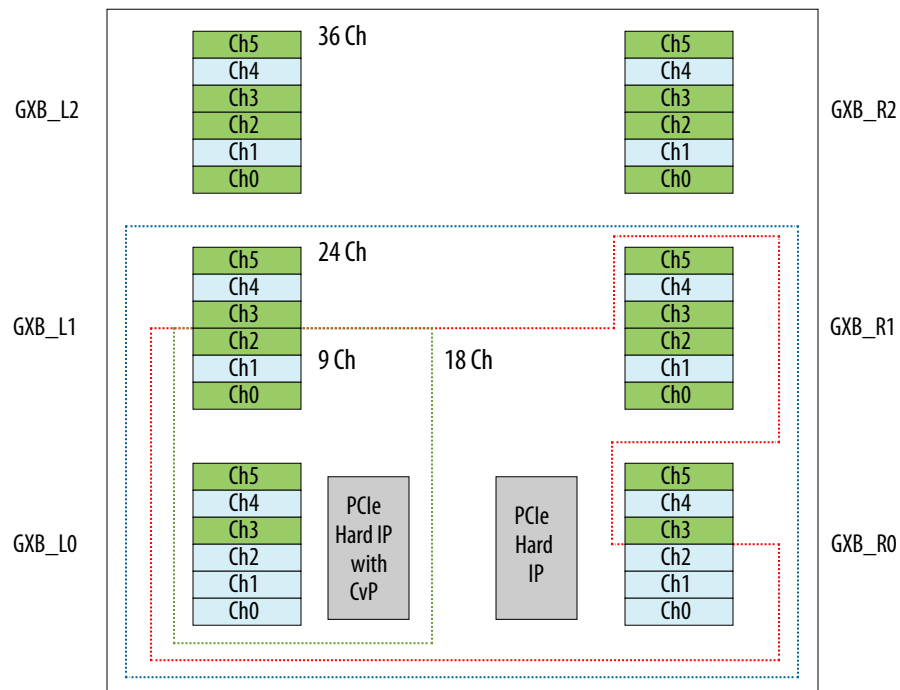
Transceiver channels are arranged in groups of six. For GX devices, the lowest six channels on the left side of the device are labeled GXB\_L0, the next group is GXB\_L1, and so on. Channels on the right side of the device are labeled GXB\_R0, GXB\_R1, and so on. Be sure to connect the Hard IP for PCI Express on the



left side of the device to appropriate channels on the left side of the device, as specified in the *Pin-out Files for Intel Devices*.

Arria V devices include one or two Hard IP for PCI Express IP Cores. The following figures illustrates the placement of the Hard IP for PCIe IP cores, transceiver banks and channels for the largest Arria V devices. Note that the bottom left IP core includes the CvP functionality. Devices with a single Hard IP for PCIe IP Core only include the bottom left core.

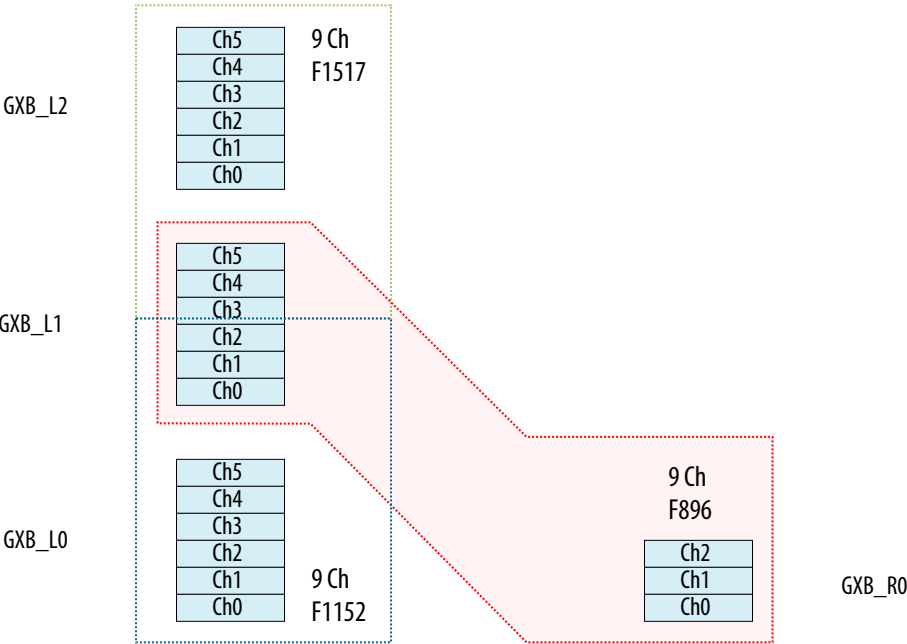
**Figure 4-11: Transceiver Bank and Hard IP for PCI Express IP Core Locations in Arria GX and GT Devices**



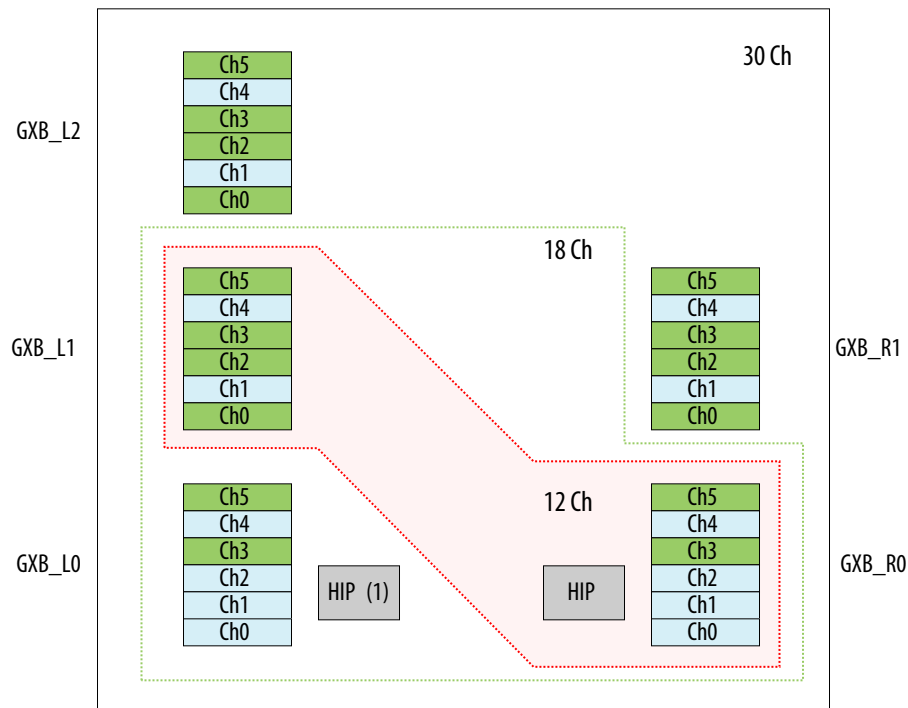
**Notes:**

1. Green blocks are 10-Gbps channels.
2. Blue blocks are 6-Gbps channels.

Figure 4-12: Transceiver Bank and Hard IP for PCI Express IP Core Locations in Arria SX Devices



Note: Blue blocks are 6 Gbps channels.

**Figure 4-13: Transceiver Bank and Hard IP for PCI Express IP Core Locations in Arria ST Devices****Notes:**

1. PCIe HIP availability varies with device variants.
2. Green blocks are 10-Gbps channels.
3. Blue blocks are 6-Gbps channels. With the exception of Ch0 to Ch2 in GXB\_L0 and GXB\_R0, the 6-Gbps channels can be used for TX-only or RX-only 10-Gbps channels.

Refer to *Transceiver Architecture in Arria V Devices* for comprehensive information on the number of Hard IP for PCIe IP cores available in various Arria V packages.

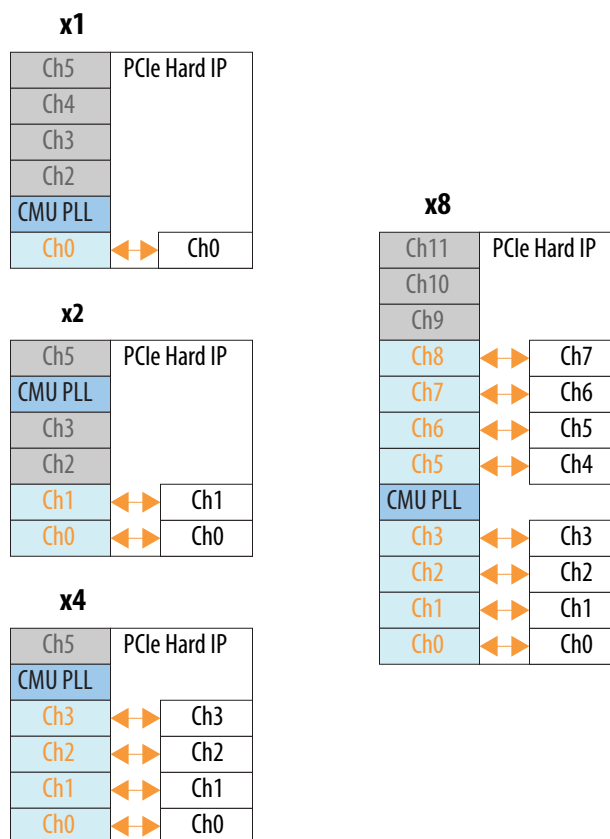
**Related Information**

- [Transceiver Architecture in Arria V Devices](#)
- [Pin-Out Files for Intel Devices](#)

## Channel Placement in Arria V Devices

**Figure 4-14: Arria V Gen1 and Gen2 Channel Placement Using the CMU PLL**

In the following figures the channels shaded in blue provide the transmit CMU PLL generating the high-speed serial clock.



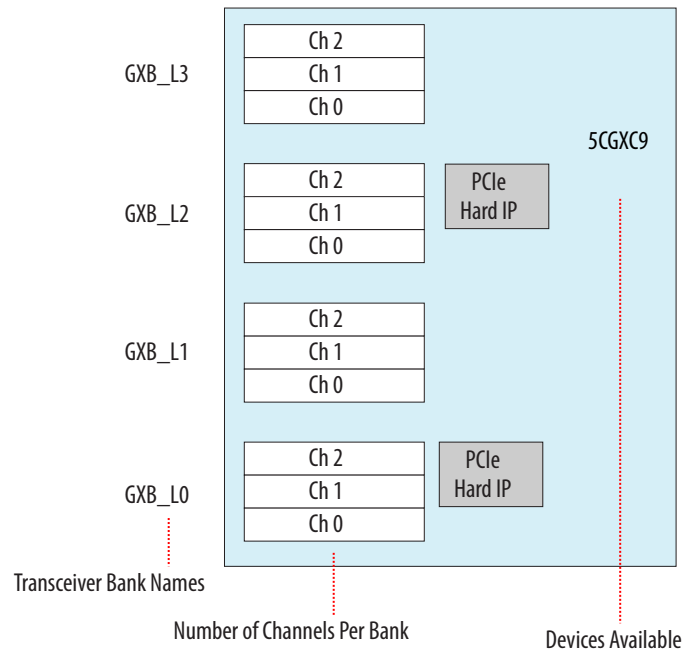
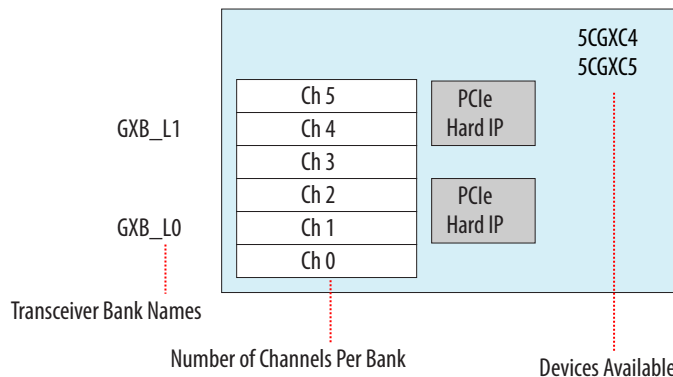
You can assign other protocols to unused channels if the data rate and clock specification exactly match the PCIe configuration.

## Physical Layout of Hard IP in Cyclone V Devices

Cyclone V devices include one or two Hard IP for PCI Express IP cores. The following figures illustrate the placement of the PCIe IP cores, transceiver banks, and channels. Note that the bottom left IP core includes the CvP functionality. The other Hard IP blocks do not include the CvP functionality. Transceiver banks include six channels. Within a bank, channels are arranged in 3-packs. GXB\_L0 contains channels 0–2, GXB\_L1 includes channels 3–5, and so on.

**Figure 4-15: Cyclone V GX/GT/ST/ST Devices with 9 or 12 Transceiver Channels and 2 PCIe Cores**

In the following figure, the x1 Hard IP for PCI Express uses channel 0 and channel 1 of GXB\_L0 and channel 0 and channel 1 of GXB\_L2.

**Figure 4-16: Cyclone V GX/GT/ST/ST Devices with 6 Transceiver Channels and 2 PCIe Cores**

For more comprehensive information about Cyclone V transceivers, refer to the *Transceiver Banks* section in the *Transceiver Architecture in Cyclone V Devices*.

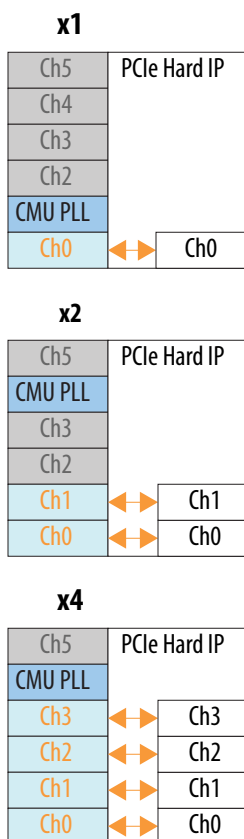
#### Related Information

- [Transceiver Architecture in Cyclone V Devices](#)
- [Pin-Out Files for Intel Devices](#)

## Channel Placement in Cyclone V Devices

**Figure 4-17: Cyclone V Gen1 and Gen2 Channel Placement Using the CMU PLL**

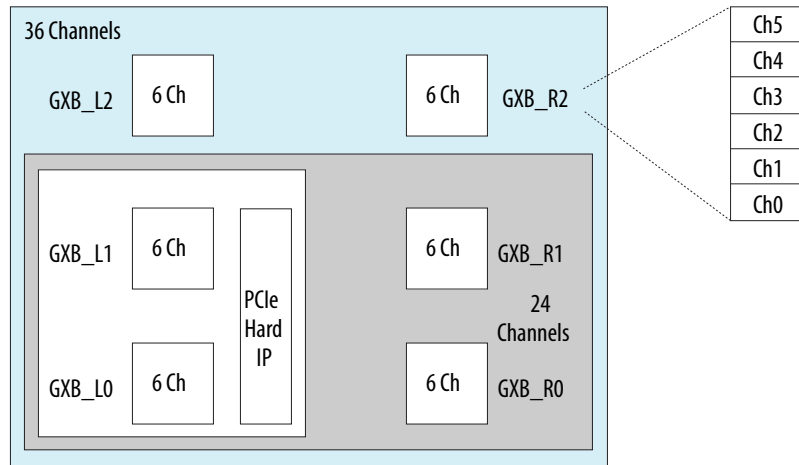
In the following figures the channels shaded in blue provide the transmit CMU PLL generating the high-speed serial clock.



You can assign other protocols to unused channels if the data rate and clock specification exactly match the PCIe configuration.

## Physical Layout of Hard IP in Arria V GZ Devices

Arria V GZ devices include one Hard IP for PCI Express IP core. The following figures illustrate the placement of the PCIe IP core, transceiver banks, and channels.

**Figure 4-18: Physical Layout of Hard IP in Arria V GZ Devices****Notes:**

1. 12-channel devices use banks L0 and L1.
2. All channels capable of backplane support up to 12.5 Gbps.

Refer to *Transceiver Architecture in Arria V Devices* for comprehensive information on the number of Hard IP for PCIe IP cores available in various Arria V GZ packages.

Refer to *Channel Utilization for Data and Clock Routing in Arria V GZ and Stratix V Devices* for additional information about channel and PLL utilization.

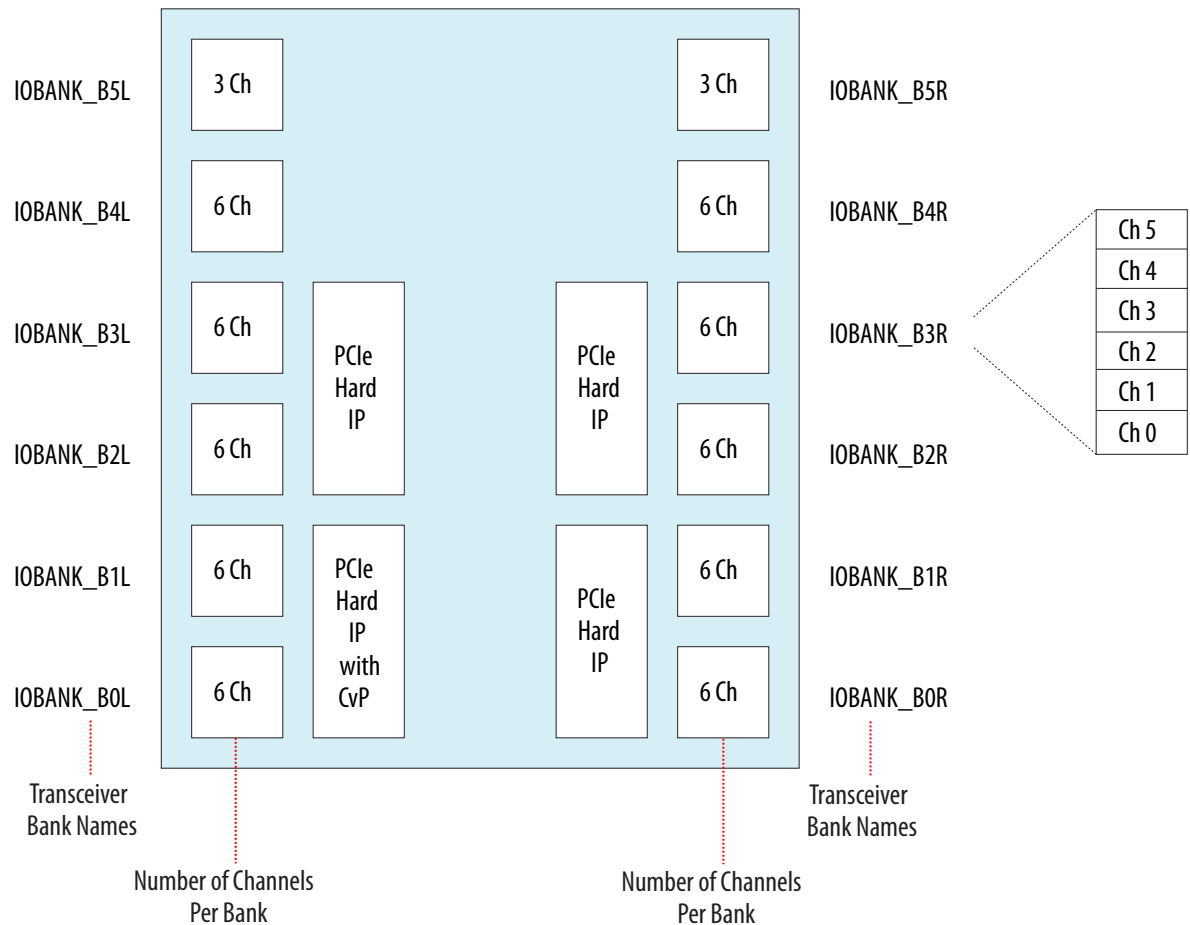
**Related Information**

- [Channel Placement in Arria V GZ and Stratix V GX/GT/GS Devices](#) on page 4-36
- [Transceiver Architecture in Arria V Devices](#)
- [Pin-Out Files for Intel Devices](#)

**Physical Layout of Hard IP in Stratix V GX/GT/GS Devices**

Stratix V devices include one, two, or four Hard IP for PCI Express IP cores. The following figures illustrate the placement of the PCIe IP cores, transceiver banks, and channels for the largest Stratix V devices. Note that the bottom left hard IP block includes the CvP functionality for flip chip packages. For other package types, the CvP functionality is in the bottom right block. All other Hard IP blocks do not include the CvP functionality.

Figure 4-19: Stratix V GX/GT/GS Devices with Four PCIe Hard IP Blocks



Smaller devices include the following PCIe Hard IP Cores:

- One Hard IP for PCIe IP core - bottom left IP core with CvP, located at GX banks L0 and L1
- Two Hard IP for PCIe IP cores - bottom left IP core with CvP and bottom right IP Core, located at banks L0 and L1, and banks R0 and R1

Refer to *Stratix V GX/GT Channel and PCIe Hard IP (HIP) Layout* for comprehensive information on the number of Hard IP for PCIe IP cores available in various Stratix V packages.

#### Related Information

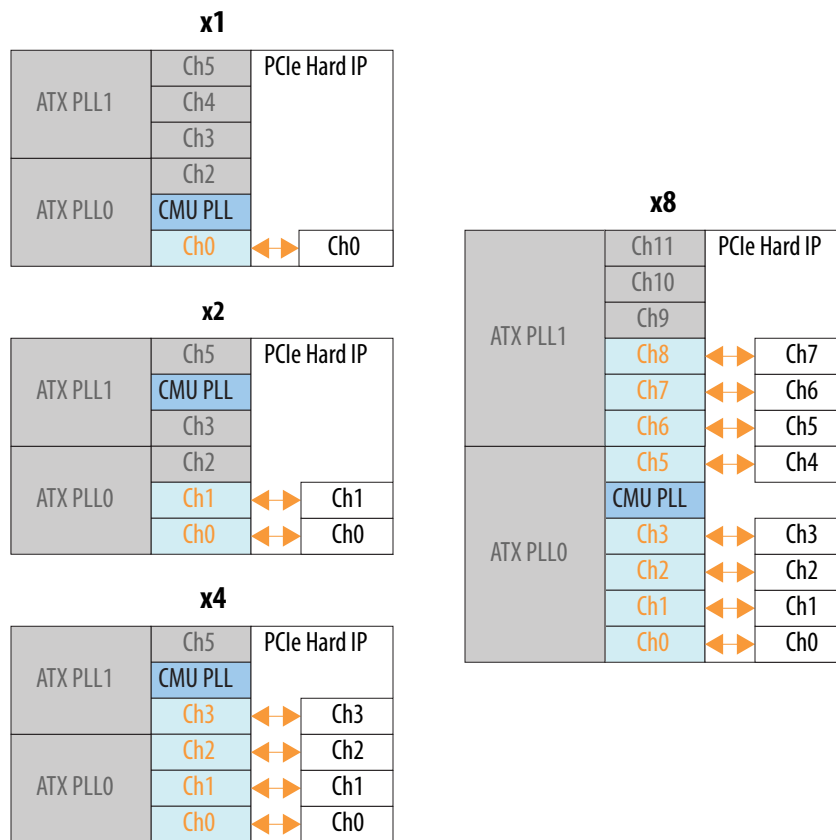
- [Transceiver Architecture in Stratix V Devices](#)
- [Pin-Out Files for Intel Devices](#)



## Channel Placement in Arria V GZ and Stratix V GX/GT/GS Devices

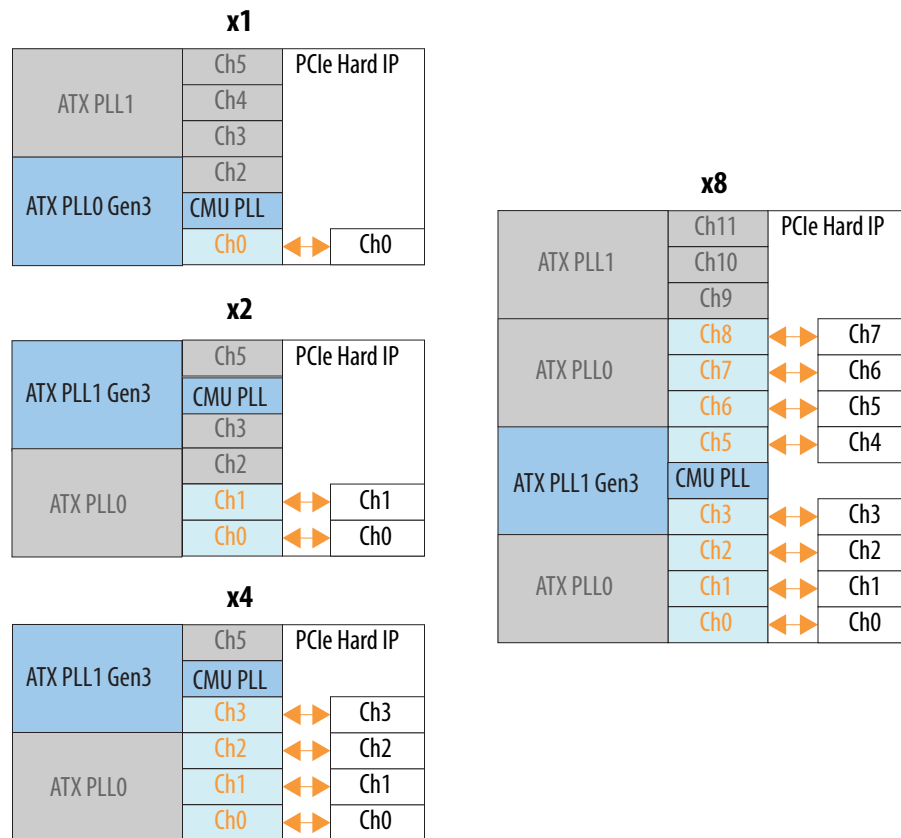
**Figure 4-20: Arria V GZ and Stratix V GX/GT/GS Gen1 and Gen2 Channel Placement Using the CMU PLL**

In the following figures the channels shaded in blue provide the transmit CMU PLL generating the high-speed serial clock.



**Figure 4-21: Arria V GZ and Stratix V GX/GT/GS Gen3 Channel Placement Using the CMU and ATX PLLs**

Gen3 requires two PLLs to facilitate rate switching between the Gen1, Gen2, and Gen3 data rates. Channels shaded in blue provide the transmit CMU PLL generating the high-speed serial clock. The ATX PLL shaded in blue is the ATX PLL used in these configurations.

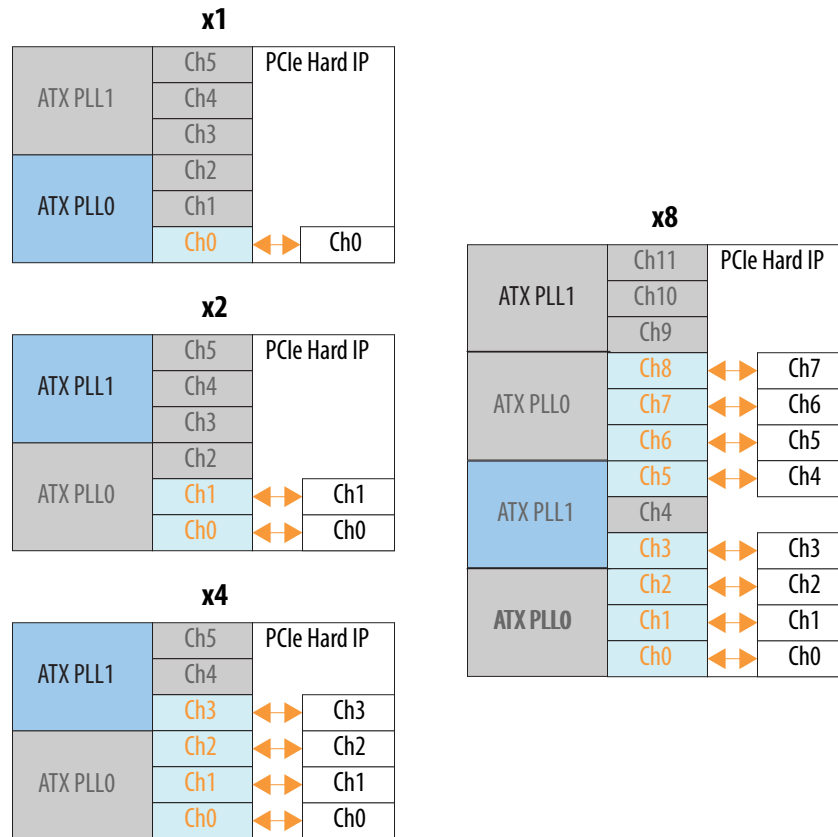


**Figure 4-22: Arria V GZ and Stratix V GX/GT/GS Gen1 and Gen2 Channel Placement Using the ATX PLL**

Selecting the ATX PLL has the following advantages over selecting the CMU PLL:

- The ATX PLL saves one channel in Gen1 and Gen2  $\times 1$ ,  $\times 2$ , and  $\times 4$  configurations.
- The ATX PLL has better jitter performance than the CMU PLL.

**Note:** You must use the soft reset controller when you select the ATX PLL and you cannot use CvP.



## PIPE Interface Signals

These PIPE signals are available for Gen1, Gen2, and Gen3 variants so that you can simulate using either the serial or the PIPE interface. Note that Intel Arria 10 and Intel Cyclone 10 GX devices do not support the Gen3 PIPE interface. Simulation is faster using the PIPE interface because the PIPE simulation bypasses the SERDES model. By default, the PIPE interface is 8 bits for Gen1 and Gen2 and 32 bits for Gen3. You can use the PIPE interface for simulation even though your actual design includes a serial interface to the internal transceivers. However, it is not possible to use the Hard IP PIPE interface in hardware, including probing these signals using Signal Tap Embedded Logic Analyzer. These signals are not top-level signals of the Hard IP. They are listed here to assist in debugging link training issues.

These PIPE signals are available for Gen1 and Gen2 variants so that you can simulate using either the serial or the PIPE interface. Simulation is faster using the PIPE interface because the PIPE simulation bypasses the SERDES model. The PIPE interface is 8 bits for Gen1 and Gen2. You can use the PIPE interface for simulation even though your actual design includes a serial interface to the internal

transceivers. However, it is not possible to use the hard IP PIPE interface in hardware, including probing these signals using Signal Tap Embedded Logic Analyzer. These signals are not top-level signals of the Hard IP. They are listed here to assist in debugging link training issues.

**Note:** The Intel Root Port BFM bypasses Gen3 Phase 2 and Phase 3 Equalization. However, Gen3 variants can perform Phase 2 and Phase 3 equalization if instructed by a third-party BFM.

**Table 4-23: PIPE Interface Signals**

In the following table, signals that include lane number 0 also exist for lanes 1-7. These signals are for simulation only. For Quartus Prime software compilation, these pipe signals can be left floating. In Platform Designer, the signals that are part of the PIPE interface have the prefix, *hip\_pipe*. The signals which are included to simulate the PIPE interface have the prefix, *hip\_pipe\_sim\_pipe*.

Signal	Direction	Description
<code>eidleinferse0[2:0]</code>	Output	Electrical idle entry inference mechanism selection. The following encodings are defined: <ul style="list-style-type: none"> <li>3'b0xx: Electrical Idle Inference not required in current LTSSM state</li> <li>3'b100: Absence of COM/SKP Ordered Set the in 128 us window for Gen1 or Gen2</li> <li>3'b101: Absence of TS1/TS2 Ordered Set in a 1280 UI interval for Gen1 or Gen2</li> <li>3'b110: Absence of Electrical Idle Exit in 2000 UI interval for Gen1 and 16000 UI interval for Gen2</li> <li>3'b111: Absence of Electrical idle exit in 128 us window for Gen1</li> </ul>
<code>phystatus0</code>	Input	PHY status $\langle n \rangle$ . This signal communicates completion of several PHY requests.
<code>powerdown0[1:0]</code>	Output	Power down $\langle n \rangle$ . This signal requests the PHY to change its power state to the specified state (P0, P0s, P1, or P2).
<code>rxdata0[31:0]</code>	Input	Receive data. This bus receives data on lane $\langle n \rangle$ .
<code>rxdatak0[3:0]</code>	Input	Data/Control bits for the symbols of receive data. Bit 0 corresponds to the lowest-order byte of <code>rxdata</code> , and so on. A value of 0 indicates a data byte. A value of 1 indicates a control byte. For Gen1 and Gen2 only.
<code>rxelecidle0</code>	Input	Receive electrical idle $\langle n \rangle$ . When asserted, indicates detection of an electrical idle.
<code>rxpolarity0</code>	Output	Receive polarity $\langle n \rangle$ . This signal instructs the PHY layer to invert the polarity of the 8B/10B receiver decoding block.

Signal	Direction	Description
<code>rxstatus0[2:0]</code>	Input	Receive status $\langle n \rangle$ . This signal encodes receive status and error codes for the receive data stream and receiver detection.
<code>rxvalid0</code>	Input	Receive valid $\langle n \rangle$ . This symbol indicates symbol lock and valid data on <code>rxdata<math>\langle n \rangle</math></code> and <code>rxdatak<math>\langle n \rangle</math></code> .
<code>sim_pipe_ ltssmstate0[4:0]</code>	Input and Output	<p>LTSSM state: The LTSSM state machine encoding defines the following states:</p> <ul style="list-style-type: none"> <li>• 5'b00000: Detect.Quiet</li> <li>• 5'b 00001: Detect.Active</li> <li>• 5'b00010: Polling.Active</li> <li>• 5'b 00011: Polling.Compliance</li> <li>• 5'b 00100: Polling.Configuration</li> <li>• 5'b00101: Polling.Speed</li> <li>• 5'b00110: Config.LinkwidthsStart</li> <li>• 5'b 00111: Config.Linkaccept</li> <li>• 5'b 01000: Config.Lanenumaccept</li> <li>• 5'b01001: Config.Lanenumwait</li> <li>• 5'b01010: Config.Complete</li> <li>• 5'b 01011: Config.Idle</li> <li>• 5'b01100: Recovery.Rcvlock</li> <li>• 5'b01101: Recovery.Rcvconfig</li> <li>• 5'b01110: Recovery.Idle</li> <li>• 5'b 01111: L0</li> <li>• 5'b10000: Disable</li> <li>• 5'b10001: Loopback.Entry</li> <li>• 5'b10010: Loopback.Active</li> <li>• 5'b10011: Loopback.Exit</li> <li>• 5'b10100: Hot.Reset</li> <li>• 5'b10101: L0s</li> <li>• 5'b11001: L2.transmit.Wake</li> <li>• 5'b11010: Recovery.Speed</li> <li>• 5'b11011: Recovery.Equalization, Phase 0</li> <li>• 5'b11100: Recovery.Equalization, Phase 1</li> <li>• 5'b11101: Recovery.Equalization, Phase 2</li> <li>• 5'b11110: Recovery.Equalization, Phase 3</li> <li>• 5'b11111: Recovery.Equalization, Done</li> </ul>
<code>sim_pipe_pclk_in</code>	Input	This clock is used for PIPE simulation only, and is derived from the <code>refclk</code> . It is the PIPE interface clock used for PIPE mode simulation.



Signal	Direction	Description
sim_pipe_rate[1:0]	Input	Specifies the data rate. The 2-bit encodings have the following meanings: <ul style="list-style-type: none"> <li>2'b00: Gen1 rate (2.5 Gbps)</li> <li>2'b01: Gen2 rate (5.0 Gbps)</li> <li>2'b1X: Gen3 rate (8.0 Gbps)</li> </ul>
txcompl0	Output	Transmit compliance <n>. This signal forces the running disparity to negative in compliance mode (negative COM character).
txdata0[31:0]	Output	Transmit data. This bus transmits data on lane <n>.
txdatak0[3:0]	Output	Transmit data control <n>. This signal serves as the control bit for txdata <n>. Bit 0 corresponds to the lowest-order byte of rxdata, and so on. A value of 0 indicates a data byte. A value of 1 indicates a control byte. For Gen1 and Gen2 only.
txdataskip0	Output	For Gen3 operation. Allows the MAC to instruct the TX interface to ignore the TX data interface for one clock cycle. The following encodings are defined: <ul style="list-style-type: none"> <li>1'b0: TX data is invalid</li> <li>1'b1: TX data is valid</li> </ul>
txdeemph0	Output	Transmit de-emphasis selection. The value for this signal is set based on the indication received from the other end of the link during the Training Sequences (TS). You do not need to change this value.
txdetectrx0	Output	Transmit detect receive <n>. This signal tells the PHY layer to start a receive detection operation or to begin loopback.
txelecidle0	Output	Transmit electrical idle <n>. This signal forces the TX output to electrical idle.
txmargin0[2:0]	Output	Transmit V <sub>OD</sub> margin selection. The value for this signal is based on the value from the Link Control 2 Register. Available for simulation only.
txswing0	Output	When asserted, indicates full swing for the transmitter voltage. When deasserted indicates half swing.



Signal	Direction	Description
txsynchd0[1:0]	Output	For Gen3 operation, specifies the block type. The following encodings are defined: <ul style="list-style-type: none"> <li>2'b01: Ordered Set Block</li> <li>2'b10: Data Block</li> </ul>

## Test Signals

**Table 4-24: Test Interface Signals**

The `test_in` bus provides run-time control and monitoring of the internal state of the IP core.

Signal	Direction	Description
test_in[31:0]	Input	<p>The bits of the <code>test_in</code> bus have the following definitions:</p> <ul style="list-style-type: none"> <li>[0]: Simulation mode. This signal can be set to 1 to accelerate initialization by reducing the value of many initialization counters.</li> <li>[1]: Reserved. Must be set to 1'b0</li> <li>[2]: Descramble mode disable. This signal must be set to 1 during initialization in order to disable data scrambling. You can use this bit in simulation for Gen1 and Gen2 Endpoints and Root Ports to observe descrambled data on the link. Descrambled data cannot be used in open systems because the link partner typically scrambles the data.</li> <li>[4:3]: Reserved. Must be set to 2'b01.</li> <li>[5]: Compliance test mode. Disable/force compliance mode. When set, prevents the LTSSM from entering compliance mode. Toggling this bit controls the entry and exit from the compliance state, enabling the transmission of Gen1, Gen2 and Gen3 compliance patterns.</li> <li>[6]: Forces entry to compliance mode when a timeout is reached in the polling.active state and not all lanes have detected their exit condition.</li> <li>[7]: Disable low power state negotiation. Intel recommends setting this bit.</li> <li>[31:8]: Reserved. Set to all 0s.</li> </ul> <p>For more information about using the <code>test_in</code> to debug, refer to the Knowledge Base Solution <i>How can I observe the Hard IP for PCI Express PIPE interface signals for Arria V GZ and Stratix V devices?</i> in the <i>Related Links</i> below.</p>
simu_pipe_mode	Input	When 1'b1, counter values are reduced to speed simulation.

Signal	Direction	Description
currentspeed[1:0]	Output	Indicates the current speed of the PCIe link. The following encodings are defined: <ul style="list-style-type: none"><li>• 2b'00: Undefined</li><li>• 2b'01: Gen1</li><li>• 2b'10: Gen2</li><li>• 2b'11: Gen3</li></ul>

**Related Information**

[How can I observe the Hard IP for PCI Express PIPE interface signals for Arria V GZ and Stratix V devices?](#)



2018.07.31

UG-01154



Subscribe



Send Feedback

## Correspondence between Configuration Space Registers and the PCIe Specification

**Table 5-1: Address Map of Hard IP Configuration Space Registers**

For the Type 0 and Type 1 Configuration Space Headers, the first line of each entry lists Type 0 values and the second line lists Type 1 values when the values differ.

Byte Address	Hard IP Configuration Space Register	Corresponding Section in PCIe Specification
0x000:0x03C	PCI Header Type 0 Configuration Registers	Type 0 Configuration Space Header
0x000:0x03C	PCI Header Type 1 Configuration Registers	Type 1 Configuration Space Header The Type 1 Configuration Space is not available for the Avalon-MM with DMA interface
0x040:0x04C	Reserved	N/A
0x050:0x05C	MSI Capability Structure	MSI Capability Structure
0x068:0x070	MSI-X Capability Structure	MSI-X Capability Structure
0x070:0x074	Reserved	N/A
0x078:0x07C	Power Management Capability Structure	PCI Power Management Capability Structure
0x080:0x0BC	PCI Express Capability Structure	PCI Express Capability Structure
0x0C0:0x0FC	Reserved	N/A
0x100:0x16C	Virtual Channel Capability Structure	Virtual Channel Capability

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

**ALTERA**  
now part of Intel

Byte Address	Hard IP Configuration Space Register	Corresponding Section in PCIe Specification
0x170:0x17C	Reserved	N/A
0x180:0x1FC	Virtual channel arbitration table	VC Arbitration Table
0x200:0x23C	Port VC0 arbitration table	Port Arbitration Table
0x240:0x27C	Port VC1 arbitration table	Port Arbitration Table
0x280:0x2BC	Port VC2 arbitration table	Port Arbitration Table
0x2C0:0x2FC	Port VC3 arbitration table	Port Arbitration Table
0x300:0x33C	Port VC4 arbitration table	Port Arbitration Table
0x340:0x37C	Port VC5 arbitration table	Port Arbitration Table
0x380:0x3BC	Port VC6 arbitration table	Port Arbitration Table
0x3C0:0x3FC	Port VC7 arbitration table	Port Arbitration Table
0x400:0x7FC	Reserved	PCIe spec corresponding section name
0x800:0x834	Advanced Error Reporting AER (optional)	Advanced Error Reporting Capability
0x838:0xFFF	Reserved	N/A
Overview of Configuration Space Register Fields		
0x000	Device ID, Vendor ID	Type 0 Configuration Space Header Type 1 Configuration Space Header The Type 1 Configuration Space is not available for the Avalon-MM with DMA interface
0x004	Status, Command	Type 0 Configuration Space Header Type 1 Configuration Space Header The Type 1 Configuration Space is not available for the Avalon-MM with DMA interface

Byte Address	Hard IP Configuration Space Register	Corresponding Section in PCIe Specification
0x008	Class Code, Revision ID	Type 0 Configuration Space Header Type 1 Configuration Space Header The Type 1 Configuration Space is not available for the Avalon-MM with DMA interface
0x00C	BIST, Header Type, Primary Latency Timer, Cache Line Size	Type 0 Configuration Space Header Type 1 Configuration Space Header The Type 1 Configuration Space is not available for the Avalon-MM with DMA interface
0x010	Base Address 0	Base Address Registers
0x014	Base Address 1	Base Address Registers
0x018	Base Address 2 Secondary Latency Timer, Subordinate Bus Number, Secondary Bus Number, Primary Bus Number	Base Address Registers Secondary Latency Timer, Type 1 Configuration Space Header, Primary Bus Number
0x01C	Base Address 3 Secondary Status, I/O Limit, I/O Base	Base Address Registers Secondary Status Register ,Type 1 Configuration Space Header
0x020	Base Address 4 Memory Limit, Memory Base	Base Address Registers Type 1 Configuration Space Header
0x024	Base Address 5 Prefetchable Memory Limit, Prefetchable Memory Base	Base Address Registers Prefetchable Memory Limit, Prefetchable Memory Base
0x028	Reserved Prefetchable Base Upper 32 Bits	N/A Type 1 Configuration Space Header
0x02C	Subsystem ID, Subsystem Vendor ID Prefetchable Limit Upper 32 Bits	Type 0 Configuration Space Header Type 1 Configuration Space Header
0x030	I/O Limit Upper 16 Bits, I/O Base Upper 16 Bits	Type 0 Configuration Space Header Type 1 Configuration Space Header

Byte Address	Hard IP Configuration Space Register	Corresponding Section in PCIe Specification
0x034	Reserved, Capabilities PTR	Type 0 Configuration Space Header Type 1 Configuration Space Header
0x038	Reserved	N/A
0x03C	Interrupt Pin, Interrupt Line Bridge Control, Interrupt Pin, Interrupt Line	Type 0 Configuration Space Header Type 1 Configuration Space Header
0x050	MSI-Message Control Next Cap Ptr Capability ID	MSI and MSI-X Capability Structures
0x054	Message Address	MSI and MSI-X Capability Structures
0x058	Message Upper Address	MSI and MSI-X Capability Structures
0x05C	Reserved Message Data	MSI and MSI-X Capability Structures
0x068	MSI-X Message Control Next Cap Ptr Capability ID	MSI and MSI-X Capability Structures
0x06C	MSI-X Table Offset BIR	MSI and MSI-X Capability Structures
0x070	Pending Bit Array (PBA) Offset BIR	MSI and MSI-X Capability Structures
0x078	Capabilities Register Next Cap PTR Cap ID	PCI Power Management Capability Structure
0x07C	Data PM Control/Status Bridge Extensions Power Management Status & Control	PCI Power Management Capability Structure
0x080	PCI Express Capabilities Register Next Cap Ptr PCI Express Cap ID	PCI Express Capability Structure
0x084	Device Capabilities Register	PCI Express Capability Structure
0x088	Device Status Register Device Control Register	PCI Express Capability Structure
0x08C	Link Capabilities Register	PCI Express Capability Structure
0x090	Link Status Register Link Control Register	PCI Express Capability Structure
0x094	Slot Capabilities Register	PCI Express Capability Structure

Byte Address	Hard IP Configuration Space Register	Corresponding Section in PCIe Specification
0x098	Slot Status Register Slot Control Register	PCI Express Capability Structure
0x09C	Root Capabilities Register Root Control Register	PCI Express Capability Structure
0x0A0	Root Status Register	PCI Express Capability Structure
0x0A4	Device Capabilities 2 Register	PCI Express Capability Structure
0x0A8	Device Status 2 Register Device Control 2 Register	PCI Express Capability Structure
0x0AC	Link Capabilities 2 Register	PCI Express Capability Structure
0x0B0	Link Status 2 Register Link Control 2 Register	PCI Express Capability Structure
0x0B4:0x0BC	Reserved	PCI Express Capability Structure
0x800	Advanced Error Reporting Enhanced Capability Header	Advanced Error Reporting Enhanced Capability Header
0x804	Uncorrectable Error Status Register	Uncorrectable Error Status Register
0x808	Uncorrectable Error Mask Register	Uncorrectable Error Mask Register
0x80C	Uncorrectable Error Severity Register	Uncorrectable Error Severity Register
0x810	Correctable Error Status Register	Correctable Error Status Register
0x814	Correctable Error Mask Register	Correctable Error Mask Register
0x818	Advanced Error Capabilities and Control Register	Advanced Error Capabilities and Control Register
0x81C	Header Log Register	Header Log Register
0x82C	Root Error Command	Root Error Command Register
0x830	Root Error Status	Root Error Status Register
0x834	Error Source Identification Register Correctable Error Source ID Register	Error Source Identification Register

## Related Information

[PCI Express Base Specification 2.1 or 3.0](#)

## Type 0 Configuration Space Registers

Figure 5-1: Type 0 Configuration Space Registers - Byte Address Offsets and Layout

Endpoints store configuration data in the Type 0 Configuration Space. The [Correspondence between Configuration Space Registers and the PCIe Specification](#) on page 5-1 lists the appropriate section of the *PCI Express Base Specification* that describes these registers.

	31	24	23	16	15	8	7	0	
0x000	Device ID					Vendor ID			
0x004	Status					Command			
0x008	Class Code						Revision ID		
0x00C	0x00	Header Type			0x00		Cache Line Size		
0x010	BAR Registers								
0x014	BAR Registers								
0x018	BAR Registers								
0x01C	BAR Registers								
0x020	BAR Registers								
0x024	BAR Registers								
0x028	Reserved								
0x02C	Subsystem Device ID					Subsystem Vendor ID			
0x030	Expansion ROM Base Address								
0x034	Reserved						Capabilities Pointer		
0x038	Reserved								
0x03C	0x00				Interrupt Pin		Interrupt Line		

## Type 1 Configuration Space Registers

Figure 5-2: Type 1 Configuration Space Registers (Root Ports)

	31	24	23	16	15	8	7	0	
0x0000	Device ID					Vendor ID			
0x0004	Status					Command			
0x0008	Class Code						Revision ID		
0x000C	BIST		Header Type		Primary Latency Timer		Cache Line Size		
0x0010	BAR Registers								
0x0014	BAR Registers								
0x0018	Secondary Latency Timer		Subordinate Bus Number		Secondary Bus Number		Primary Bus Number		
0x001C	Secondary Status				I/O Limit		I/O Base		
0x0020	Memory Limit				Memory Base				
0x0024	Prefetchable Memory Limit				Prefetchable Memory Base				
0x0028	Prefetchable Base Upper 32 Bits								
0x002C	Prefetchable Limit Upper 32 Bits								
0x0030	I/O Limit Upper 16 Bits				I/O Base Upper 16 Bits				
0x0034	Reserved						Capabilities Pointer		
0x0038	Expansion ROM Base Address								
0x003C	Bridge Control				Interrupt Pin		Interrupt Line		

**Note:** Avalon-MM DMA for PCIe does not support Type 1 configuration space registers.

## PCI Express Capability Structures

The layout of the most basic Capability Structures are provided below. Refer to the *PCI Express Base Specification* for more information about these registers.

Figure 5-3: MSI Capability Structure

	31	24 23	16 15	8 7	0
0x050	Message Control Configuration MSI Control Status Register Field Descriptions			Next Cap Ptr	Capability ID
0x054	Message Address				
0x058	Message Upper Address				
0x05C	Reserved			Message Data	

**Note:** Refer to the *Advanced Error Reporting Capability* section for more details about the PCI Express AER Extended Capability Structure.

## Related Information

- [PCI Express Base Specification 3.0](#)
- [PCI Local Bus Specification](#)

## Intel-Defined VSEC Registers

**Figure 5-4: VSEC Registers**

This extended capability structure supports Configuration via Protocol (CvP) programming and detailed internal error reporting.

	31	20 19	16 15	8 7	0
0x200	Next Capability Offset		Version	Intel-Defined VSEC Capability Header	
0x204	VSEC Length		VSEC Revision	VSEC ID Intel-Defined, Vendor-Specific Header	
0x208	Intel Marker				
0x20C	JTAG Silicon ID DW0 JTAG Silicon ID				
0x210	JTAG Silicon ID DW1 JTAG Silicon ID				
0x214	JTAG Silicon ID DW2 JTAG Silicon ID				
0x218	JTAG Silicon ID DW3 JTAG Silicon ID				
0x21C	CvP Status			User Device or Board Type ID	
0x220	CvP Mode Control				
0x224	CvP Data2 Register				
0x228	CvP Data Register				
0x22C	CvP Programming Control Register				
0x230	Reserved				
0x234	Uncorrectable Internal Error Status Register				
0x238	Uncorrectable Internal Error Mask Register				
0x23C	Correctable Internal Error Status Register				
0x240	Correctable Internal Error Mask Register				

**Table 5-2: Intel-Defined VSEC Capability Register, 0x200**

The Intel-Defined Vendor Specific Extended Capability. This extended capability structure supports Configuration via Protocol (CvP) programming and detailed internal error reporting.

Bits	Register Description	Value	Access
[15:0]	PCI Express Extended Capability ID. Intel-defined value for VSEC Capability ID.	0x000B	RO
[19:16]	Version. Intel-defined value for VSEC version.	0x1	RO
[31:20]	Next Capability Offset. Starting address of the next Capability Structure implemented, if any.	Variable	RO



**Table 5-3: Intel-Defined Vendor Specific Header**

You can specify these values when you instantiate the Hard IP. These registers are read-only at run-time.

Bits	Register Description	Value	Access
[15:0]	VSEC ID. A user configurable VSEC ID.	User entered	RO
[19:16]	VSEC Revision. A user configurable VSEC revision.	Variable	RO
[31:20]	VSEC Length. Total length of this structure in bytes.	0x044	RO

**Table 5-4: Intel Marker Register**

Bits	Register Description	Value	Access
[31:0]	Intel Marker. This read only register is an additional marker. If you use the standard Intel Programmer software to configure the device with CvP, this marker provides a value that the programming software reads to ensure that it is operating with the correct VSEC.	A Device Value	RO

**Table 5-5: JTAG Silicon ID Register**

Bits	Register Description	Value	Access
[127:96]	JTAG Silicon ID DW3	Application Specific	RO
[95:64]	JTAG Silicon ID DW2	Application Specific	RO
[63:32]	JTAG Silicon ID DW1	Application Specific	RO
[31:0]	JTAG Silicon ID DW0. This is the JTAG Silicon ID that CvP programming software reads to determine that the correct SRAM object file (.sof) is being used.	Application Specific	RO

**Table 5-6: User Device or Board Type ID Register**

Bits	Register Description	Value	Access
[15:0]	Configurable device or board type ID to specify to CvP the correct .sof.	Variable	RO

## CvP Registers

**Table 5-7: CvP Status**

The CvP Status register allows software to monitor the CvP status signals.

Bits	Register Description	Reset Value	Access
[31:26]	Reserved	0x00	RO
[25]	PLD_CORE_READY. From FPGA fabric. This status bit is provided for debug.	Variable	RO
[24]	PLD_CLK_IN_USE. From clock switch module to fabric. This status bit is provided for debug.	Variable	RO
[23]	CVP_CONFIG_DONE. Indicates that the FPGA control block has completed the device configuration via CvP and there were no errors.	Variable	RO
[22]	Reserved	Variable	RO
[21]	USERMODE. Indicates if the configurable FPGA fabric is in user mode.	Variable	RO
[20]	CVP_EN. Indicates if the FPGA control block has enabled CvP mode.	Variable	RO
[19]	CVP_CONFIG_ERROR. Reflects the value of this signal from the FPGA control block, checked by software to determine if there was an error during configuration.	Variable	RO
[18]	CVP_CONFIG_READY. Reflects the value of this signal from the FPGA control block, checked by software during programming algorithm.	Variable	RO
[17:0]	Reserved	Variable	RO

**Table 5-8: CvP Mode Control**

The CvP Mode Control register provides global control of the CvP operation.

Bits	Register Description	Reset Value	Access
[31:16]	Reserved.	0x0000	RO
[15:8]	CVP_NUMCLKS.  This is the number of clocks to send for every CvP data write. Set this field to one of the values below depending on your configuration image: <ul style="list-style-type: none"> <li>0x01 for uncompressed and unencrypted images</li> <li>0x04 for uncompressed and encrypted images</li> <li>0x08 for all compressed images</li> </ul>	0x00	RW
[7:3]	Reserved.	0x0	RO

Bits	Register Description	Reset Value	Access
[2]	CVP_FULLCONFIG. Request that the FPGA control block reconfigure the entire FPGA including the V-Series Hard IP for PCI Express, bring the PCIe link down.	1'b0	RW
[1]	<p>HIP_CLK_SEL. Selects between PMA and fabric clock when USER_MODE = 1 and PLD_CORE_READY = 1. The following encodings are defined:</p> <ul style="list-style-type: none"> <li>1: Selects internal clock from PMA which is required for CVP_MODE.</li> <li>0: Selects the clock from soft logic fabric. This setting should only be used when the fabric is configured in USER_MODE with a configuration file that connects the correct clock.</li> </ul> <p>To ensure that there is no clock switching during CvP, you should only change this value when the Hard IP for PCI Express has been idle for 10 <math>\mu</math>s and wait 10 <math>\mu</math>s after changing this value before resuming activity.</p>	1'b0	RW
[0]	<p>CVP_MODE. Controls whether the IP core is in CVP_MODE or normal mode. The following encodings are defined:</p> <ul style="list-style-type: none"> <li>1: CVP_MODE is active. Signals to the FPGA control block active and all TLPs are routed to the Configuration Space. This CVP_MODE cannot be enabled if CVP_EN = 0.</li> <li>0: The IP core is in normal mode and TLPs are routed to the FPGA fabric.</li> </ul>	1'b0	RW

**Table 5-9: CvP Data Registers**

The following table defines the CvP Data registers. For 64-bit data, the optional CvP Data2 stores the upper 32 bits of data. Programming software should write the configuration data to these registers. If you Every write to these register sets the data output to the FPGA control block and generates  $\langle n \rangle$  clock cycles to the FPGA control block as specified by the CVP\_NUM\_CLKS field in the CvP Mode Control register. Software must ensure that all bytes in the memory write dword are enabled. You can access this register using configuration writes, alternatively, when in CvP mode, these registers can also be written by a memory write to any address defined by a memory space BAR for this device. Using memory writes should allow for higher throughput than configuration writes.

Bits	Register Description	Reset Value	Access
[31:0]	Upper 32 bits of configuration data to be transferred to the FPGA control block to configure the device. You can choose 32- or 64-bit data.	0x00000000	RW
[31:0]	Lower 32 bits of configuration data to be transferred to the FPGA control block to configure the device.	0x00000000	RW

**Table 5-10: CvP Programming Control Register**

This register is written by the programming software to control CvP programming.

Bits	Register Description	Reset Value	Access
[31:2]	Reserved.	0x0000	RO
[1]	START_XFER. Sets the CvP output to the FPGA control block indicating the start of a transfer.	1'b0	RW
[0]	CVP_CONFIG. When asserted, instructs that the FPGA control block begin a transfer via CvP.	1'b0	RW

## Advanced Error Reporting Capability

### Uncorrectable Internal Error Mask Register

**Table 5-11: Uncorrectable Internal Error Mask Register**

The `Uncorrectable Internal Error Mask` register controls which errors are forwarded as internal uncorrectable errors. With the exception of the configuration error detected in CvP mode, all of the errors are severe and may place the device or PCIe link in an inconsistent state. The configuration error detected in CvP mode may be correctable depending on the design of the programming software. The access code *RWS* stands for Read Write Sticky meaning the value is retained after a soft reset of the IP core.

Bits	Register Description	Reset Value	Access
[31:12]	Reserved.	1b'0	RO
[11]	Mask for RX buffer posted and completion overflow error.	1b'0	RWS
[10]	Reserved	1b'1	RO
[9]	Mask for parity error detected on Configuration Space to TX bus interface.	1b'1	RWS
[8]	Mask for parity error detected on the TX to Configuration Space bus interface.	1b'1	RWS
[7]	Mask for parity error detected at TX Transaction Layer error.	1b'1	RWS
[6]	Reserved	1b'1	RO
[5]	Mask for configuration errors detected in CvP mode.	1b'0	RWS
[4]	Mask for data parity errors detected during TX Data Link LCRC generation.	1b'1	RWS

Bits	Register Description	Reset Value	Access
[3]	Mask for data parity errors detected on the RX to Configuration Space Bus interface.	1b'1	RWS
[2]	Mask for data parity error detected at the input to the RX Buffer.	1b'1	RWS
[1]	Mask for the retry buffer uncorrectable ECC error.	1b'1	RWS
[0]	Mask for the RX buffer uncorrectable ECC error.	1b'1	RWS

## Uncorrectable Internal Error Status Register

**Table 5-12: Uncorrectable Internal Error Status Register**

This register reports the status of the internally checked errors that are uncorrectable. When specific errors are enabled by the `Uncorrectable Internal Error Mask` register, they are handled as Uncorrectable Internal Errors as defined in the *PCI Express Base Specification 3.0*. This register is for debug only. It should only be used to observe behavior, not to drive custom logic. The access code RW1CS represents Read Write 1 to Clear Sticky.

Bits	Register Description	Reset Value	Access
[31:12]	Reserved.	0	RO
[11]	When set, indicates an RX buffer overflow condition in a posted request or Completion	0	RW1CS
[10]	Reserved.	0	RO
[9]	When set, indicates a parity error was detected on the Configuration Space to TX bus interface	0	RW1CS
[8]	When set, indicates a parity error was detected on the TX to Configuration Space bus interface	0	RW1CS
[7]	When set, indicates a parity error was detected in a TX TLP and the TLP is not sent.	0	RW1CS
[6]	When set, indicates that the Application Layer has detected an uncorrectable internal error.	0	RW1CS
[5]	When set, indicates a configuration error has been detected in CVP mode which is reported as uncorrectable. This bit is set whenever a <code>CVP_CONFIG_ERROR</code> rises while in <code>CVP_MODE</code> .	0	RW1CS

Bits	Register Description	Reset Value	Access
[4]	When set, indicates a parity error was detected by the TX Data Link Layer.	0	RW1CS
[3]	When set, indicates a parity error has been detected on the RX to Configuration Space bus interface.	0	RW1CS
[2]	When set, indicates a parity error was detected at input to the RX Buffer.	0	RW1CS
[1]	When set, indicates a retry buffer uncorrectable ECC error.	0	RW1CS
[0]	When set, indicates a RX buffer uncorrectable ECC error.	0	RW1CS

**Related Information**

[PCI Express Base Specification 2.1 or 3.0](#)

## Correctable Internal Error Mask Register

**Table 5-13: Correctable Internal Error Mask Register**

The `Correctable Internal Error Mask` register controls which errors are forwarded as Internal Correctable Errors. This register is for debug only.

Bits	Register Description	Reset Value	Access
[31:8]	Reserved.	0	RO
[7]	Reserved.	1	RO
[6]	Mask for Corrected Internal Error reported by the Application Layer.	1	RWS
[5]	Mask for configuration error detected in CvP mode.	1	RWS
[4:2]	Reserved.	0	RO
[1]	Mask for retry buffer correctable ECC error.	1	RWS
[0]	Mask for RX Buffer correctable ECC error.	1	RWS

## Correctable Internal Error Status Register

**Table 5-14: Correctable Internal Error Status Register**

The `Correctable Internal Error Status` register reports the status of the internally checked errors that are correctable. When these specific errors are enabled by the `Correctable Internal Error Mask` register, they are forwarded as Correctable Internal Errors as defined in the *PCI Express Base Specification 3.0*. This register is for debug only. Only use this register to observe behavior, not to drive logic custom logic.

Bits	Register Description	Reset Value	Access
[31:7]	Reserved.	0	RO
[6]	Corrected Internal Error reported by the Application Layer.	0	RW1CS
[5]	When set, indicates a configuration error has been detected in CvP mode which is reported as correctable. This bit is set whenever a <code>CVP_CONFIG_ERROR</code> occurs while in <code>CVP_MODE</code> .	0	RW1CS
[4:2]	Reserved.	0	RO
[1]	When set, the retry buffer correctable ECC error status indicates an error.	0	RW1CS
[0]	When set, the RX buffer correctable ECC error status indicates an error.	0	RW1CS

### Related Information

[PCI Express Base Specification 2.1 or 3.0](#)

## DMA Descriptor Controller Registers

The DMA Descriptor Controller manages Read and Write DMA operations. The DMA Descriptor Controller is available for use with Endpoint variations. The Descriptor Controller supports up to 128 descriptors each for Read and Write Data Movers. Read and Write are from the perspective of the FPGA. A read is from PCIe address space to the FPGA Avalon-MM address space. A write is to PCIe address space from the FPGA Avalon-MM space.

You program the Descriptor Controller internal registers with the location and size of the descriptor table residing in the PCIe address space. The DMA Descriptor Controller instructs the Read Data Mover to copy the table to its own internal FIFO. When the DMA Descriptor Controller is instantiated as a separate component, it drives table entries on the `RdDmaRxData_i[159:0]` and `WrDmaRxData_i[159:0]` buses. When the DMA Descriptor Controller is embedded inside the Avalon-MM DMA bridge, it drives this information on internal buses.

The Read Data Mover transfers data from the PCIe address space to Avalon-MM address space. It issues memory read TLPs on the PCIe link. It writes the data returned to a location in the Avalon-MM address

space. The source address is the address for the data in the PCIe address space. The destination address is in the Avalon-MM address space.

The Write Data Mover reads data from the Avalon-MM address space and writes to the PCIe address space. It issues memory write TLPs on the PCIe link. The source address is in the Avalon-MM address space. The destination address is in the PCIe address space.

The DMA Descriptor Controller records the completion status for read and write descriptors in separate status tables. Each table has 128 consecutive DWORD entries that correspond to the 128 descriptors. The actual descriptors are stored immediately after the status entries at offset 0x200 from the values programmed into the `RC_Read_Descriptor_Base` and `RC_Write_Descriptor_Base` registers. The status and descriptor table must be located on a 32-byte boundary in Root Complex memory.

The Descriptor Controller writes a 1 to the `done` bit of the status DWORD to indicate successful completion. The Descriptor Controller also sends an MSI interrupt for the final descriptor. After receiving this MSI, host software can poll the `done` bit to determine status. The status table precedes the descriptor table in memory. The Descriptor Controller does not write the `done` bit nor send an MSI as each descriptor completes. It only writes the `done` bit or sends an MSI for the descriptor whose ID is stored in the `RD_DMA_LAST_PTR` or `WR_DMA_LAST_PTR` registers. The Descriptor Controller supports out-of-order completions. Consequently, it is possible for the `done` bit to be set before all descriptors have completed.

**Note:** The following example clarifies this programming model. If 128 descriptors are specified and all of them execute, then descriptor ID 127 is written to the `RD_DMA_LAST_PTR` or `WR_DMA_LAST_PTR` register to start the DMA. The DMA Descriptor Controller only writes the `done` bit when descriptor 127 completes. To get intermediate status updates, host software should write multiple IDs into the last pointer register. For example, to get an intermediate status update when half of the 128 read descriptors have completed, host software should complete the following sequence:

1. Program the `RD_DMA_LAST_PTR = 63`.
2. Program the `RD_DMA_LAST_PTR = 127`.
3. Poll the status DWORD for read descriptor 63.
4. Poll the status DWORD for read descriptor 127.

Many commercial system Root Ports return out-of-order Read Completions based on optimized accesses to host memory channels. Consequently, the `done` status stored for descriptor  $\langle n \rangle$  does not necessarily mean that descriptors  $\langle n-1 \rangle$  and  $\langle n-2 \rangle$  have also completed. You must request the completion status for every descriptor by writing the descriptor ID for every descriptor to `RD_DMA_LAST_PTR` or `WR_DMA_LAST_PTR`.

**Note:** Because the DMA Descriptor Controller uses FIFOs to store descriptor table entries, you cannot reprogram the DMA Descriptor Controller once it begins the transfers specified in the descriptor table.

## Read DMA Descriptor Controller Registers

The following table describes the registers in the internal DMA Descriptor Controller. When the DMA Descriptor Controller is externally instantiated, these registers are accessed through a BAR. The offsets must be added to the base address for the read controller. When the Descriptor Controller is internally instantiated these registers are accessed through BAR0. The read controller is at offset 0x0000.



Address Offset	Register	Access	Description
0x0000	RC Read Status and Descriptor Base (Low)	R/W	Specifies the lower 32-bits of the base address of the read status and descriptor table in the Root Complex memory. This address must be on a 32-byte boundary. Software must program this register after programming the upper 32 bits at offset 0x4. To change the RC Read Status and Descriptor Base (Low) base address, all descriptors specified by the RD_TABLE_SIZE must be exhausted.
0x0004	RC Read Status and Descriptor Base (High)	R/W	Specifies the upper 32-bits of the base address of the read status and descriptor table in the Root Complex memory. Software must program this register before programming the lower 32 bits of this register.
0x0008	EP Read Descriptor FIFO Base (Low)	RW	Specifies the lower 32 bits of the base address of the read descriptor FIFO in Endpoint memory. The Read DMA fetches the descriptors from Root Complex memory. The address must be the Avalon-MM address of the Descriptor Controller's Read Descriptor Table Avalon-MM Slave Port as seen by the Read DMA Avalon-MM Master Port. You must program this register after programming the upper 32 bits at offset 0xC.
0x000C	EP Read Descriptor FIFO Base (High)	RW	Specifies the upper 32 bits of the base address of the read descriptor table in Endpoint Avalon-MM memory. The Read DMA fetches the descriptors from Root Complex memory and writes the descriptors to the FIFO at this location. This must be the Avalon-MM address of the descriptor controller's Read Descriptor Table Avalon-MM Slave Port as seen by the Read DMA Avalon-MM Master Port. You must program this register before programming the lower 32 bits of this register.

Address Offset	Register	Access	Description
0x0010	RD_DMA_LAST_PTR	RW	<p>When read, returns the ID of the last descriptor requested. If no DMA request is outstanding or the DMA is in reset, returns a value 0xFF.</p> <p>When written, specifies the ID of the last descriptor requested. The difference between the value read and the value written is the number of descriptors to be processed.</p> <p>For example, if the value reads 4, the last descriptor requested is 4. To specify 5 more descriptors, software should write a 9 into the RD_DMA_LAST_PTR register. The DMA executes 5 more descriptors.</p> <p>To have the read DMA record the done bit of every descriptor, program this register to transfer one descriptor at a time.</p> <p>The descriptor ID loops back to 0 after reaching RD_TABLE_SIZE. For example, if the RD_DMA_LAST_PTR value read is 126 and you want to execute three more descriptors, software must write 127, and then 1 into the RD_DMA_LAST_PTR register.</p>
0x0014	RD_TABLE_SIZE	RW	<p>Specifies the size of the Read descriptor table. Set to the number of descriptors - 1. By default, RD_TABLE_SIZE is set to 127. This value specifies the last Descriptor ID. To change the RC Read Status and Descriptor Base (Low) base address, all descriptors specified by the RD_TABLE_SIZE must be exhausted.</p>
0x0018	RD_CONTROL	RW	<p>[31:1] Reserved.</p> <p>[0] Done. When set, the Descriptor Controller writes the Done bit for each descriptor in the status table. When not set the Descriptor Controller writes the Done for the final descriptor, as specified by RD_DMA_LAST_PTR. In both cases, the Descriptor Controller sends a MSI to the host after the completion of the last descriptor along with the status update for the last descriptor.</p>

## Write DMA Descriptor Controller Registers

The following table describes the registers in the internal DMA Descriptor Controller. When the DMA Descriptor Controller is externally instantiated, these registers are accessed through a BAR. The offsets must be added to the base address for the write controller. When the Descriptor Controller is internally instantiated these registers are accessed through BAR0. The write controller is at offset 0x0100.

Address Offset	Register	Access	Description
0x0100	RC Write Status and Descriptor Base (Low)	R/W	Specifies the lower 32-bits of the base address of the write status and descriptor table in the Root Complex memory. This address must be on a 32-byte boundary. Software must program this register after programming the upper 32-bit register at offset 0x104. To change the RC Write Status and Descriptor Base (Low) base address, all descriptors specified by the <code>WR_TABLE_SIZE</code> must be exhausted.
0x0104	RC Write Status and Descriptor Base (High)	R/W	Specifies the upper 32-bits of the base address of the write status and descriptor table in the Root Complex memory. Software must program this register before programming the lower 32-bit register at offset 0x100.
0x0108	EP Write Status and Descriptor FIFO Base (Low)	RW	Specifies the lower 32 bits of the base address of the write descriptor table in Endpoint memory. The Write Descriptor Controller requests descriptors from Root Complex memory and writes the descriptors to this location. The address is the Avalon-MM address of the Descriptor Controller's Write Descriptor Table Avalon-MM Slave Port as seen by the Read DMA Avalon-MM Master Port. Software must program this register after programming the upper 32-bit register at offset 0x10C.
0x010C	EP Write Status and Descriptor FIFO Base (High)	RW	Specifies the upper 32 bits of the base address of the write descriptor table in Endpoint memory. The read DMA fetches the table from Root Complex memory and writes the table to this location. Software must program this register before programming the lower 32-bit register at offset 0x108.

Address Offset	Register	Access	Description
0x0110	WR_DMA_LAST_PTR	RW	<p>When read, returns the ID of the last descriptor requested. If no DMA request is outstanding or the DMA is in reset, returns a value 0xFF.</p> <p>When written, specifies the ID of the last descriptor requested. The difference between the value read and the value written is the number of descriptors to be processed.</p> <p>For example, if the value reads 4, the last descriptor requested is 4. To specify 5 more descriptors, software should write a 9 into the WR_DMA_LAST_PTR register. The DMA executes 5 more descriptors.</p> <p>To have the write DMA record the done bit of every descriptor, program this register to transfer one descriptor at a time.</p> <p>The Descriptor ID loops back to 0 after reaching WR_TABLE_SIZE.</p> <p>For example, if the WR_DMA_LAST_PTR value read is 126 and you want to execute three more descriptors, software must write 127, and then 1 into the WR_DMA_LAST_PTR register.</p>
0x0114	WR_TABLE_SIZE	RW	<p>Specifies the size of the Read descriptor table. Set to the number of descriptors - 1. By default, WR_TABLE_SIZE is set to 127. This value specifies the last Descriptor ID. To change the RC Write Status and Descriptor Base (Low) base address, all descriptors specified by the WR_TABLE_SIZE must be exhausted.</p>



Address Offset	Register	Access	Description
0x0118	WR_CONTROL	RW	<p>[31:1] Reserved.</p> <p>[0]Done. When set, the Descriptor Controller writes the Done bit for each descriptor in the status table. The Descriptor Controller sends a single MSI interrupt after the final descriptor completes. When not set the Descriptor Controller generates Done for the final descriptor, as specified by WR_DMA_LAST_PTR. In both cases, the Descriptor Controller sends an MSI to the host after the completion of the last descriptor along with the status update for the last descriptor.</p>

## Read DMA and Write DMA Descriptor Format

Read and write descriptors are stored in separate descriptor tables in PCIe system memory. Each table can store up to 128 descriptors. Each descriptor is 8 DWORDs, or 32 bytes. The Read DMA and Write DMA descriptor tables start at a 0x200 byte offset from the addresses programmed into the Read Status and Descriptor Base and Write Status and Descriptor Base address registers.

Programming RD\_DMA\_LAST\_PTR or WR\_DMA\_LAST\_PTR registers triggers the Read or Write Descriptor Controller descriptor table fetch process. Consequently, writing these registers must be the last step in setting up DMA transfers.

**Note:** Because the DMA Descriptor Controller uses FIFOs to store descriptor table entries, you cannot reprogram the DMA Descriptor Controller once it begins the transfers specified in the descriptor table.

**Table 5-15: Read Descriptor Format**

You must also use this format for the Read and Write Data Movers on their Avalon-ST when you use your own DMA Controller.

Address Offset	Register Name	Description
0x00	RD_LOW_SRC_ADDR	Lower DWORD of the read DMA source address. Specifies the address in PCIe system memory from which the Read Data Mover fetches data.
0x04	RD_HIGH_SRC_ADDR	Upper DWORD of the read DMA source address. Specifies the address in PCIe system memory from which the Read Data Mover fetches data.
0x08	RD_CTRL_LOW_DEST_ADDR	Lower DWORD of the read DMA destination address. Specifies the address in the Avalon-MM domain to which the Read Data Mover writes data.

Address Offset	Register Name	Description
0x0C	RD_CTRL_HIGH_DEST_ADDR	Upper DWORD of the read DMA destination address. Specifies the address in the Avalon-MM domain to which the Read Data Mover writes data.
0x10	CONTROL	Specifies the following information: <ul style="list-style-type: none"> <li>[31:25] Reserved must be 0.</li> <li>[24:18] ID. Specifies the <code>Descriptor ID</code>. <code>Descriptor ID 0</code> is at the beginning of the table. For descriptor tables of the maximum size, <code>Descriptor ID 127</code> is at the end of the table.</li> <li>[17:0] SIZE. The transfer size in DWORDs. Must be non-zero. The maximum transfer size is (1 MB - 4 bytes). If the specified transfer size is less than the maximum, the transfer size is the actual size entered.</li> </ul>
0x14 - 0x1C	Reserved	N/A

Table 5-16: Write Descriptor Format

Address Offset	Register Name	Description
0x00	WR_LOW_SRC_ADDR	Lower DWORD of the write DMA source address. Specifies the address in the AvalonMM domain from which the Write Data Mover fetches data.
0x04	WR_HIGH_SRC_ADDR	Upper DWORD of the write DMA source address. Specifies the address in the AvalonMM domain from which the Write Data Mover fetches data.
0x08	WR_CTRL_LOW_DEST_ADDR	Lower DWORD of the Write Data Mover destination address. Specifies the address in PCIe system memory to which the Write DMA writes data.
0x0C	WR_CTRL_HIGH_DEST_ADDR	Upper DWORD of the write DMA destination address. Specifies the address in PCIe system memory to which the Write Data Mover writes data.

Address Offset	Register Name	Description
0x10	CONTROL	<p>Specifies the following information:</p> <ul style="list-style-type: none"><li>• [31:25]: Reserved must be 0.</li><li>• [24:18]:ID: Specifies the <code>Descriptor ID</code>. <code>Descriptor ID 0</code> is at the beginning of the table. <code>Descriptor ID</code> is at the end of the table.</li><li>• [17:0] :SIZE: The transfer size in DWORDs. Must be non-zero. The maximum transfer size is (1 MB - 4 bytes). If the specified transfer size is less than the maximum, the transfer size is the actual size entered.</li></ul>
0x14 - 0x1C	Reserved	N/A

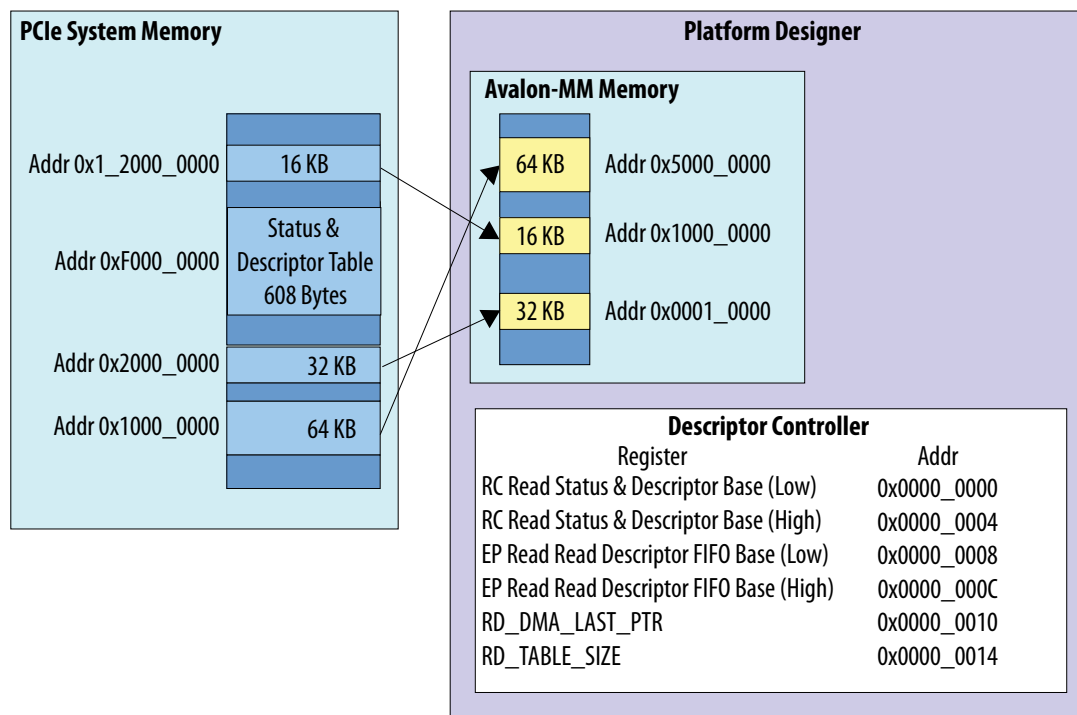
## Read DMA Example

This example moves three data blocks from the system memory to the Avalon-MM address space. Host software running on an embedded CPU allocates the memory and creates the descriptor table in system memory.

This example uses the addresses in the Platform Designer design example, `ep_g3x8_avmm256_integrated.qsys`, available in the `<install_dir>/ip/altera/altera_pcie/altera_pcie_<dev>_ed/example_design/<dev>` directory.

The following figures illustrate the location and size of the data blocks in the PCIe and Avalon-MM address spaces and the descriptor table format. In this example, the value of `RD_TABLE_SIZE` is 127.

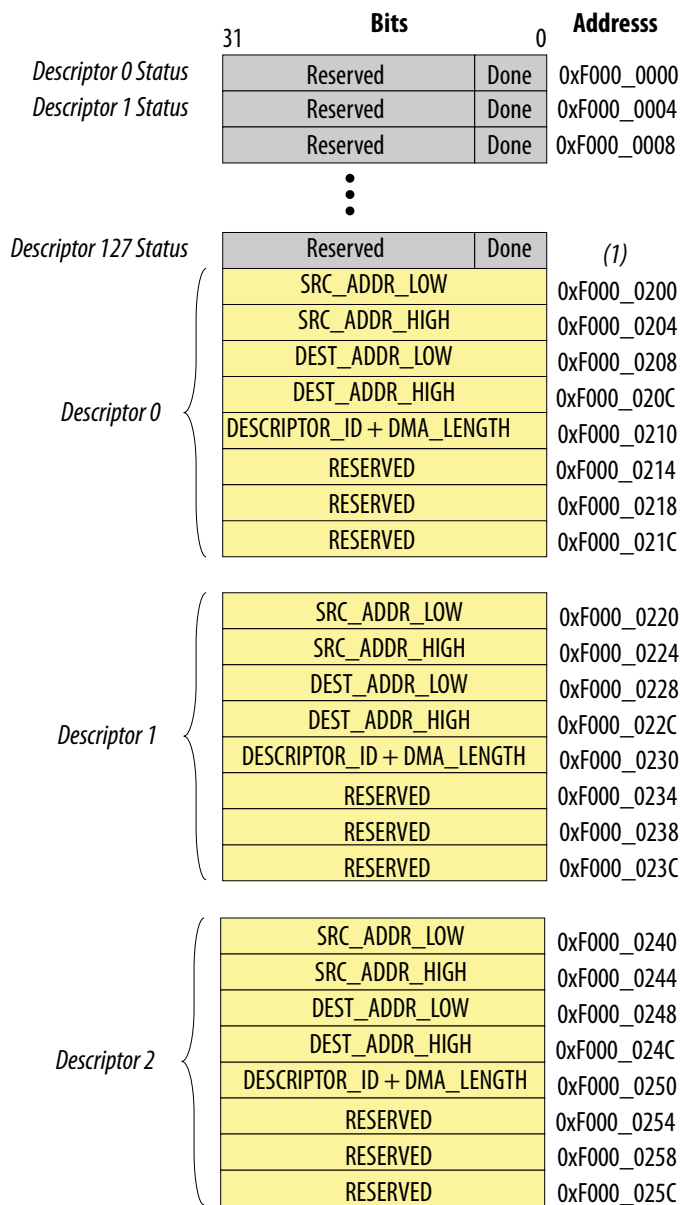
**Figure 5-5: Data Blocks to Transfer from PCIe to Avalon-MM Address Space Using Read DMA**



Assume the descriptor table includes 128 entries. The status table precedes a variable number of descriptors in memory. The Read and Write Status and Descriptor Tables are at the address specified in the RC Read Descriptor Base Register and RC Write Descriptor Base Register, respectively.



Figure 5-6: Descriptor Table Format



Note:

1. Software automatically adds 0x200 to the base address of the status table to determine the address of the first read descriptor.

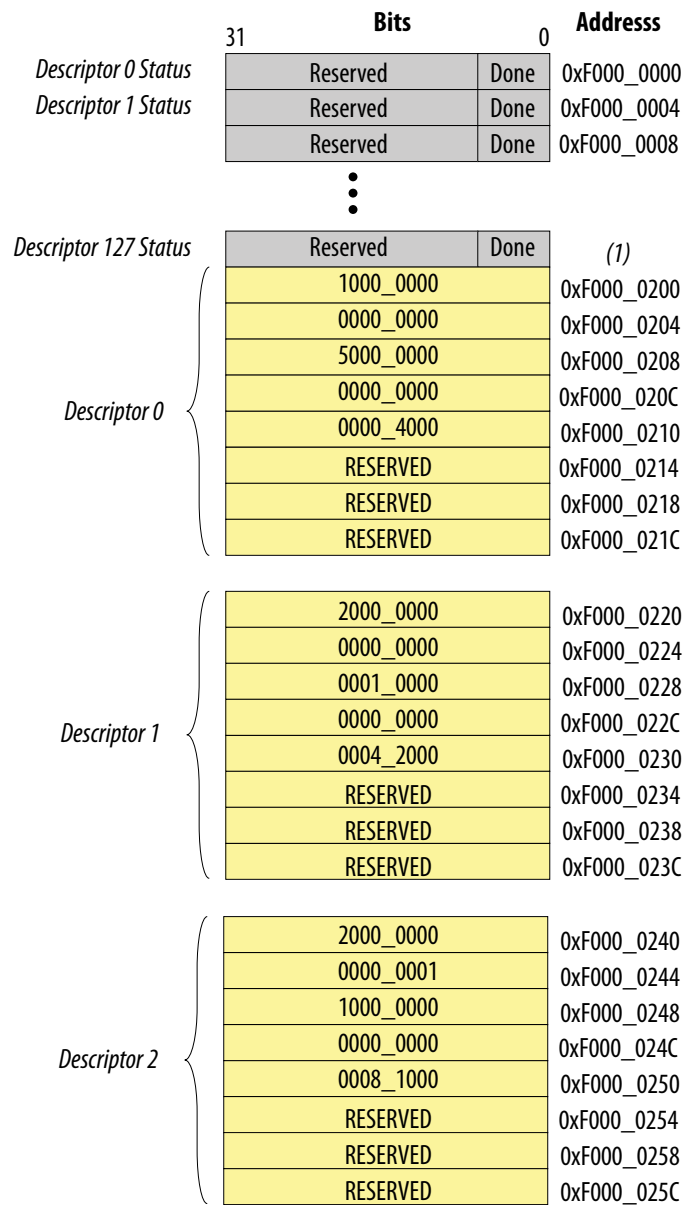
1. Calculate the memory allocation required:
  - a. Each entry in the status table is 4 bytes. The 128 entries require 512 bytes of memory.
  - b. Each descriptor is 32 bytes. The three descriptors require 96 bytes of memory.
 The total memory allocation for the status and descriptor tables is 608 bytes.
2. Allocate 608 bytes of memory in the PCI Express address space.

The start address of the allocated memory in this example is 0xF000\_0000. Program this address into the Root Complex Read Status and Descriptor Base Address Registers.

3. Create the descriptor table in the PCI Express address space. Because the status table is stored before the descriptors, the first descriptor is stored at 0xF000\_0000 + 0x200 = 0xF000\_0200.
  - a. Program 0x0000\_0000 into the source address 0xF000\_0204 in descriptor 0.  
This is the upper 32 bits of the source address.
  - b. Program 0x1000\_0000 into the source address 0xF000\_0200 in descriptor 0.  
This is the lower 32 bits of the source address.
  - c. Program 0x0000\_0000 into the destination address 0xF000\_020C in descriptor 0.  
This is the upper 32 bits of the destination address.
  - d. Program 0x5000\_0000 into the destination address 0xF000\_0208 in descriptor 0.  
This is the lower 32 bits of the destination address.  
These four steps program the Avalon-MM destination address for the 64 KB block of memory into the Descriptor Table.
  - e. Program 0x0000\_4000 to 0xF000\_0210 to transfer 16K dwords (64 KB), of data for descriptor ID 0.
4. Repeat this procedure for the second data block:
  - a. Program 0x0000\_0000 to source address 0xF000\_0224.
  - b. Program 0x2000\_0000 to source address 0xF000\_0220.
  - c. Program 0x0000\_0000 to destination address 0xF000\_022C.
  - d. Program 0x0001\_0000 to destination address 0xF000\_0228.
  - e. Program 0x0004\_2000 to 0xF000\_0230 to transfer 8K dwords (32 KB) of data for descriptor ID 1.
5. Repeat this procedure for the third data block:
  - a. Program 0x0000\_0001 to source address 0xF000\_0244.
  - b. Program 0x2000\_0000 to source address 0xF000\_0240.
  - c. Program 0x0000\_0000 to destination address 0xF000\_024C.
  - d. Program 0x1000\_0000 to destination address 0xF000\_0248.
  - e. Program 0x0008\_1000 to 0xF000\_0250 to transfer 4K dwords (16 KB) of data for descriptor ID 2.

The following figure shows the values in the Descriptor Table after programming completes.

Figure 5-7: Descriptor Table Format



Note:

1. The DMA Descriptor Controller automatically adds 0x200 to the base address of the status table to determine the address of the first read descriptor.

6. Program the DMA Descriptor Controller with the address of the status and descriptor table in the PCI Express system memory address space. When the DMA Descriptor Controller is internal, these registers are accessed through combined BAR0 and BAR1 because this example uses 64-bit addresses. The DMA read control registers start at offset 0x0000. The Write DMA control registers start at offset 0x0100
  - a. Program 0x0000\_0000 to offset 0x0000\_0004.

This is the upper 32 bits of the PCIe system memory where the status and descriptor table is stored.

- b. Program 0xF000\_0000 to offset 0x0000\_0000.

This is the lower 32 bits of the address in PCIe memory that stores the status and descriptor tables. The Read DMA automatically adds an offset of 0x200 to this value to start the copy of the descriptors which follow the status table in memory.

7. Program the DMA Descriptor Controller with the on-chip FIFO address. This is the address to which the Descriptor Controller copies the status and descriptor table.

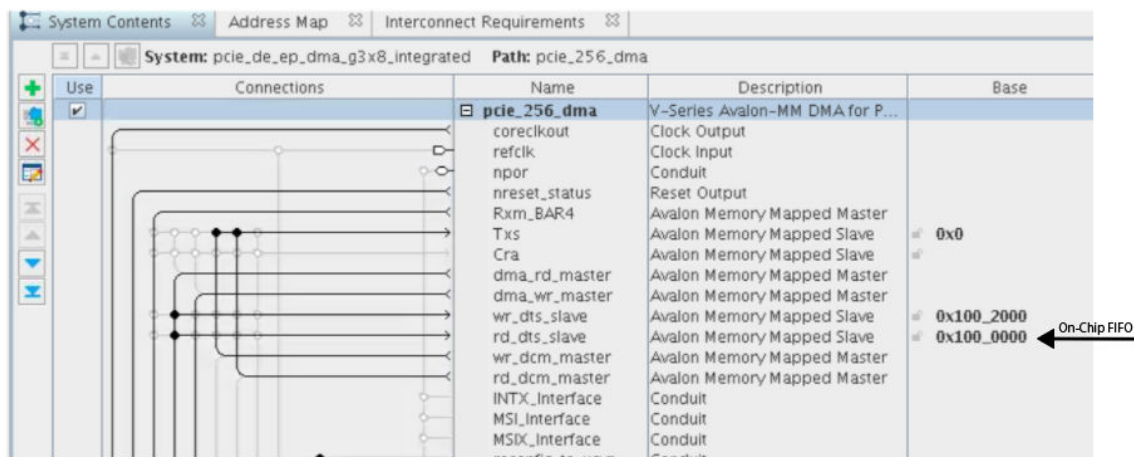
- a. Program 0x0000\_0000 to offset 0x0000\_000C

This is the upper 32 bits of the on-chip FIFO address in the Avalon-MM address domain.

- b. Program 0x0100\_0000 to offset 0x0000\_0008.

This is the lower 32 bits of the on-chip FIFO address. This is address of the internal on-chip FIFO that is a part of the Descriptor Controller as seen by the RX Master.

**Figure 5-8: Address of the On-Chip FIFO**



8. Program the Descriptor Controller RD\_DMA\_LAST\_PTR register.

This step starts the DMA. It also specifies the status dword to be updated when the three descriptors complete.

- To update a single done bit for the final descriptor, program 0x2 to offset 0x0000\_0010. The Descriptor Controller processes all three descriptors and writes the done bit to 0xF000\_0008 of the status table.
- To update the done bits for all three descriptors, program address 0x0000\_0010 RD\_DMA\_LAST\_PTR three times with the values 0, 1, and 2. The Descriptor Controller sets the done bits for addresses 0xF000\_0000, 0xF000\_0004, and 0xF000\_0008. If the system returns Read Completions out-of-order, the Descriptor Controller may complete descriptors out of order. In such systems, you must use this method of requesting done status for each descriptor. Software must check for done status for every descriptor.

## Software Program for Simultaneous Read and Write DMA

Program the following steps to implement a simultaneous DMA transfer:

1. Allocate PCIe system memory for the Read and Write DMA descriptor tables. If, for example, each table supports up to 128, eight-DWORD descriptors and 128, one-DWORD status entries for a total of 1152 DWORDs. Total memory for the Read and Write DMA descriptor tables is 2304 DWORDs.
2. Allocate PCIe system memory and initialize it with data for the Read Data Mover to read.
3. Allocate PCIe system memory for the Write Data Mover to write.
4. Create all the descriptors for the read DMA descriptor table. Assign the `DMA Descriptor IDs` sequentially, starting with 0 to a maximum of 127. For the read DMA, the source address is the memory space allocated in Step 2. The destination address is the Avalon-MM address that the Read Data Mover module writes. Specify the DMA length in DWORDs. Each descriptor transfers contiguous memory. Assuming a base address of 0, for the Read DMA, the following assignments illustrate construction of a read descriptor:
  - a. `RD_LOW_SRC_ADDR = 0x0000` (The base address for the read descriptor table in the PCIe system memory.)
  - b. `RD_HIGH_SRC_ADDR = 0x0004`
  - c. `RD_CTRL_LOW_DEST_ADDR = 0x0008`
  - d. `RD_CTRL_HIGH_DEST_ADDR = 0x000C`
  - e. `RD_DMA_LAST_PTR = 0x0010`

Writing the `RD_DMA_LAST_PTR` register starts operation.

5. For the Write DMA, the source address is the Avalon-MM address that the Write Data Mover module should read. The destination address is the PCIe system memory space allocated in Step 3. Specify the DMA size in DWORDs. Assuming a base address of 0x100, for the Write Data Mover, the following assignments illustrate construction of a write descriptor:
  - a. `WD_LOW_SRC_ADDR = 0x0100` (The base address for the write descriptor table in the PCIe system memory.)
  - b. `WD_HIGH_SRC_ADDR = 0x0104`
  - c. `WD_CTRL_LOW_DEST_ADDR = 0x0108`
  - d. `WD_CTRL_HIGH_DEST_ADDR = 0x010C`
  - e. `WD_DMA_LAST_PTR = 0x0110`

Writing the `WD_DMA_LAST_PTR` register starts operation.

6. To improve throughput, the Read DMA module copies the descriptor table to the Avalon-MM memory before beginning operation. Specify the memory address by writing to the `Descriptor Table Base (Low)` and `(High)` registers.
7. An MSI interrupt is sent for each `WD_DMA_LAST_PTR` or `RD_DMA_LAST_PTR` that completes. These completions result in updates to the `done` status bits. Host software can then read status bits to determine which DMA operations are complete.

**Note:** If the transfer size of the read DMA is greater than the maximum read request size, the Read DMA creates multiple read requests. For example, if maximum read request size is 512 bytes, the Read Data Mover breaks a 4 KB read request into 8 requests with 8 different tags. The Read Completions can come back in any order. The Read Data Mover's Avalon-MM master port writes the data received in the Read Completions to the correct locations in Avalon-MM memory, based on the tags in the same order as it receives the Completions. This order is not necessarily in increasing address order; The data mover does not include an internal reordering buffer. If system allows out

of order read completions, then status for the latest entry is latest only in number, but potentially earlier than other completions chronologically

## Control Register Access (CRA) Avalon-MM Slave Port

**Table 5-17: Configuration Space Register Descriptions**

The optional CRA Avalon-MM slave port provides host access to selected Configuration Space and status registers. These registers are read only. For registers that are less than 32 bits, the upper bits are unused.

Byte Offset	Register	Access	Description
14'h0000	cfg_dev_ctrl[15:0]	RO	cfg_devctrl[15:0] is device control for the PCI Express capability structure.
14'h0004	cfg_dev_ctrl2[15:0]	RO	cfg_dev2ctrl[15:0] is device control 2 for the PCI Express capability structure.
14'h0008	cfg_link_ctrl[15:0]	RO	<p>cfg_link_ctrl[15:0] is the primary Link Control of the PCI Express capability structure.</p> <p>For Gen2 or Gen3 operation, you must write a 1'b1 to Retrain Link bit (Bit[5] of the cfg_link_ctrl) of the Root Port to initiate retraining to a higher data rate after the initial link training to Gen1 L0 state. Retraining directs the LTSSM to the Recovery state. Retraining to a higher data rate is not automatic even if both devices on the link are capable of a higher data rate.</p>
14'h000C	cfg_link_ctrl2[15:0]	RO	<p>cfg_link_ctrl2[31:16] is the secondary Link Control register of the PCI Express capability structure for Gen2 operation.</p> <p>For Gen1 variants, the link bandwidth notification bit is always set to 0. For Gen2 variants, this bit is set to 1.</p>
14'h0010	cfg_prm_cmd[15:0]	RO	Base/Primary Command register for the PCI Configuration Space.
14'h0014	cfg_root_ctrl[7:0]	RO	Root control and status register of the PCI-Express capability. This register is only available in Root Port mode.
14'h0018	cfg_sec_ctrl[15:0]	RO	Secondary bus Control and Status register of the PCI-Express capability. This register is only available in Root Port mode.

Byte Offset	Register	Access	Description
14'h001C	cfg_secbus[7:0]	RO	Secondary bus number. Available in Root Port mode.
14'h0020	cfg_subbus[7:0]	RO	Subordinate bus number. Available in Root Port mode.
14'h0024	cfg_msi_addr_low[31:0]	RO	cfg_msi_add[31:0] is the MSI message address.
14'h0028	cfg_msi_addr_hi[63:32]	RO	cfg_msi_add[63:32] is the MSI upper message address.
14'h002C	cfg_io_bas[19:0]	RO	The IO base register of the Type1 Configuration Space. This register is only available in Root Port mode.
14'h0030	cfg_io_lim[19:0]	RO	The IO limit register of the Type1 Configuration Space. This register is only available in Root Port mode.
14'h0034	cfg_np_bas[11:0]	RO	The non-prefetchable memory base register of the Type1 Configuration Space. This register is only available in Root Port mode.
14'h0038	cfg_np_lim[11:0]	RO	The non-prefetchable memory limit register of the Type1 Configuration Space. This register is only available in Root Port mode.
14'h003C	cfg_pr_bas_low[31:0]	RO	The lower 32 bits of the prefetchable base register of the Type1 Configuration Space. This register is only available in Root Port mode.
14'h0040	cfg_pr_bas_hi[43:32]	RO	The upper 12 bits of the prefetchable base registers of the Type1 Configuration Space. This register is only available in Root Port mode.
14'h0044	cfg_pr_lim_low[31:0]	RO	The lower 32 bits of the prefetchable limit registers of the Type1 Configuration Space. This register is only available in Root Port mode.
14'h0048	cfg_pr_lim_hi[43:32]	RO	The upper 12 bits of the prefetchable limit registers of the Type1 Configuration Space. This register is only available in Root Port mode.

Byte Offset	Register	Access	Description
14'h004C	cfg_pmcsr[31:0]	RO	cfg_pmcsr[31:16] is Power Management Control and cfg_pmcsr[15:0] is the Power Management Status register.
14'h0050	cfg_msixcsr[15:0]	RO	MSI-X message control register.
14'h0054	cfg_msicsr[15:0]	RO	MSI message control.
14'h0058	cfg_tcvcmap[23:0]	RO	<p>Configuration traffic class (TC)/virtual channel (VC) mapping. The Application Layer uses this signal to generate a TLP mapped to the appropriate channel based on the traffic class of the packet.</p> <p>The following encodings are defined:</p> <ul style="list-style-type: none"> <li>cfg_tcvcmap[2:0]: Mapping for TC0 (always 0)</li> <li>cfg_tcvcmap[5:3]: Mapping for TC1.</li> <li>cfg_tcvcmap[8:6]: Mapping for TC2.</li> <li>cfg_tcvcmap[11:9]: Mapping for TC3.</li> <li>cfg_tcvcmap[14:12]: Mapping for TC4.</li> <li>cfg_tcvcmap[17:15]: Mapping for TC5.</li> <li>cfg_tcvcmap[20:18]: Mapping for TC6.</li> <li>cfg_tcvcmap[23:21]: Mapping for TC7.</li> </ul>
14'h005C	cfg_msi_data[15:0]	RO	cfg_msi_data[15:0] is message data for MSI.
14'h0060	cfg_busdev[12:0]	RO	<p>Bus/Device Number captured by or programmed in the Hard IP. The following fields are defined:</p> <ul style="list-style-type: none"> <li>cfg_busdev[12:5]: bus number</li> <li>cfg_busdev[4:0]: device number</li> </ul>
14'h0064	ltssm_reg[4:0]	RO	<p>Specifies the current LTSSM state. The LTSSM state machine encoding defines the following states:</p> <ul style="list-style-type: none"> <li>5'b: 00000: Detect.Quiet</li> <li>5'b: 00001: Detect.Active</li> <li>5'b: 00010: Polling.Active</li> <li>5'b: 00011: Polling.Compliance</li> <li>5'b: 00100: Polling.Configuration</li> <li>5'b: 00101: Polling.Speed</li> <li>5'b: 00110: config.Linkwidthstart</li> <li>5'b: 00111: Config.Linkaccept</li> <li>5'b: 01000: Config.Lanenumaccept</li> </ul>



Byte Offset	Register	Access	Description
			<ul style="list-style-type: none"> <li>5'b: 01001: Config.Lanenumwait</li> <li>5'b: 01010: Config.Complete</li> <li>5'b: 01011: Config.Idle</li> <li>5'b: 01100: Recovery.Rcvlock</li> <li>5'b: 01101: Recovery.Rcvconfig</li> <li>5'b: 01110: Recovery.Idle</li> <li>5'b: 01111: L0</li> <li>5'b: 10000: Disable</li> <li>5'b: 10001: Loopback.Entry</li> <li>5'b: 10010: Loopback.Active</li> <li>5'b: 10011: Loopback.Exit</li> <li>5'b: 10100: Hot.Reset</li> <li>5'b: 10101: LOs</li> <li>5'b: 11001: L2.transmit.Wake</li> <li>5'b: 11010: Speed.Recovery</li> <li>5'b: 11011: Recovery.Equalization, Phase 0</li> <li>5'b: 11100: Recovery.Equalization, Phase 1</li> <li>5'b: 11101: Recovery.Equalization, Phase 2</li> <li>5'b: 11110: recovery.Equalization, Phase 3</li> </ul>
14'h0068	current_speed_reg[1:0]	RO	<p>Indicates the current speed of the PCIe link. The following encodings are defined:</p> <ul style="list-style-type: none"> <li>2'b'00: Undefined</li> <li>2'b'01: Gen1</li> <li>2'b'10: Gen2</li> <li>2'b'11: Gen3</li> </ul>
14'h006C	lane_act_reg[3:0]	RO	<p>Lane Active Mode: This signal indicates the number of lanes that configured during link training. The following encodings are defined:</p> <ul style="list-style-type: none"> <li>4'b0001: 1 lane</li> <li>4'b0010: 2 lanes</li> <li>4'b0100: 4 lanes</li> <li>4'b1000: 8 lanes</li> </ul>

**Related Information**

- [PCI Express Base Specification 2.1 or 3.0](#)
- [PCI Local Bus Specification, Rev. 3.0](#)

2018.07.31

UG-01154



Subscribe



Send Feedback

V-Series Hard IP for PCI Express IP Core includes both a soft reset controller and a hard reset controller. Software selects the appropriate reset controller depending on the configuration you specify. Both reset controllers reset the IP core and provide sample reset logic in the example design. The figure below provides a simplified view of the logic that implements both reset controllers.

The `pin_perst` signal from the input pin of the FPGA resets the Hard IP for PCI Express IP Core. `app_rstn` which resets the Application Layer logic is derived from `reset_status` and `pld_clk_inuse`, which are outputs of the core. This reset controller is implemented in hardened logic. The figure below provides a simplified view of the logic that implements the reset controller.

**Table 6-1: Use of Hard and Soft Reset Controllers**

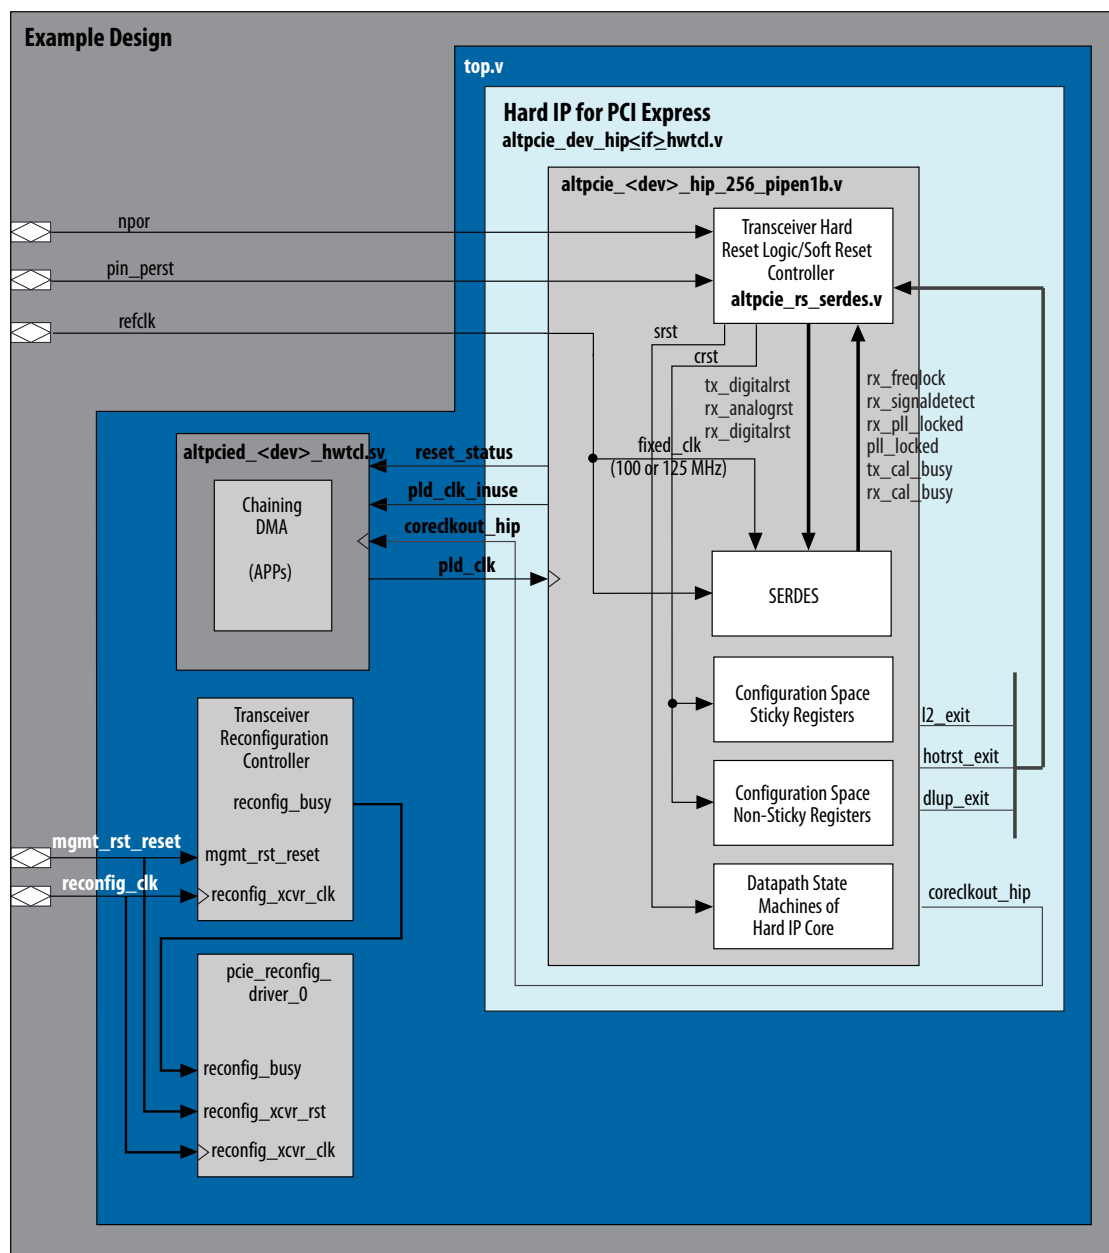
Reset Controller Used	Description
Hard Reset Controller	<code>pin_perst</code> from the input pin of the FPGA resets the Hard IP for PCI Express IP Core. <code>app_rstn</code> which resets the Application Layer logic is derived from <code>reset_status</code> and <code>pld_clk_inuse</code> , which are outputs of the core. This reset controller is supported for Gen 1 production devices.
Soft Reset Controller	Either <code>pin_perst</code> from the input pin of the FPGA or <code>npwr</code> which is derived from <code>pin_perst</code> or <code>local_rstn</code> can reset the Hard IP for PCI Express IP Core. Application Layer logic generates the optional <code>local_rstn</code> signal. <code>app_rstn</code> which resets the Application Layer logic is derived from <code>npwr</code> .  This reset controller is supported for Gen2 and Gen3 production devices.

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

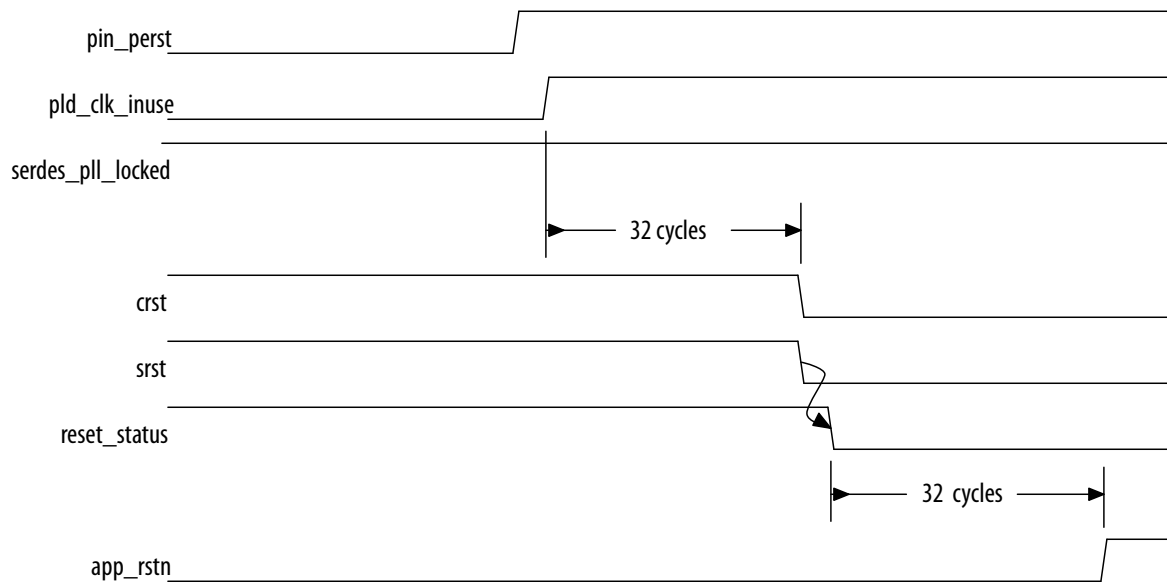
Figure 6-1: Reset Controller Block Diagram



## Reset Sequence for Hard IP for PCI Express IP Core and Application Layer

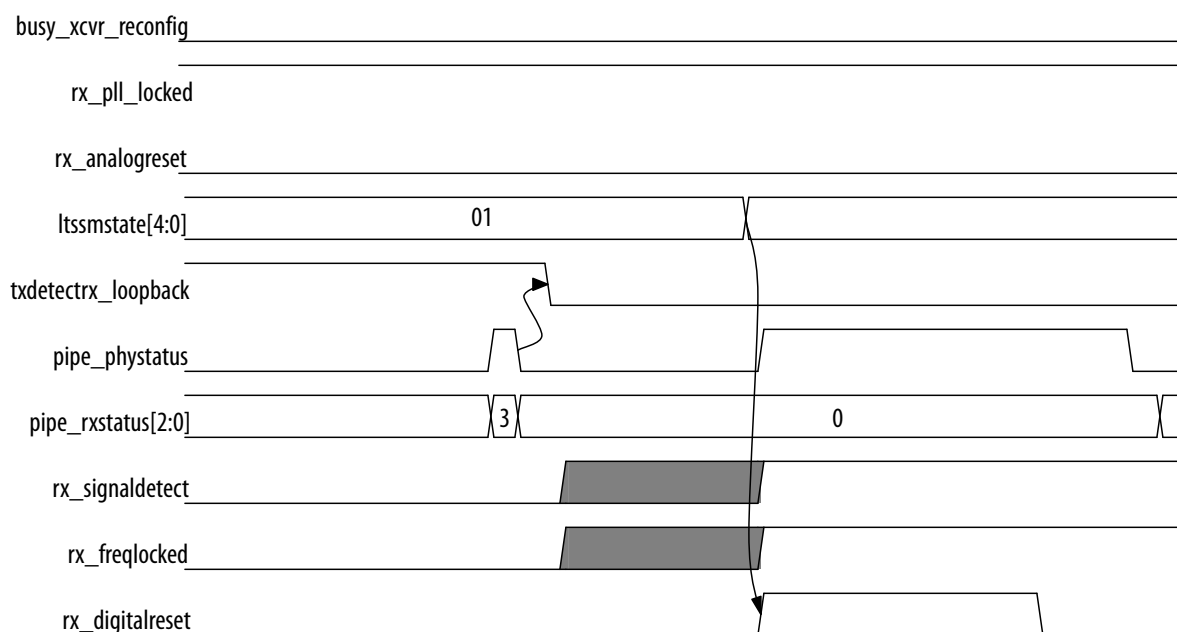
Figure 6-2: Hard IP for PCI Express and Application Logic Reset Sequence

Your Application Layer can instantiate a module with logic that implements the timing diagram shown below to generate `app_rstn`, which resets the Application Layer logic.



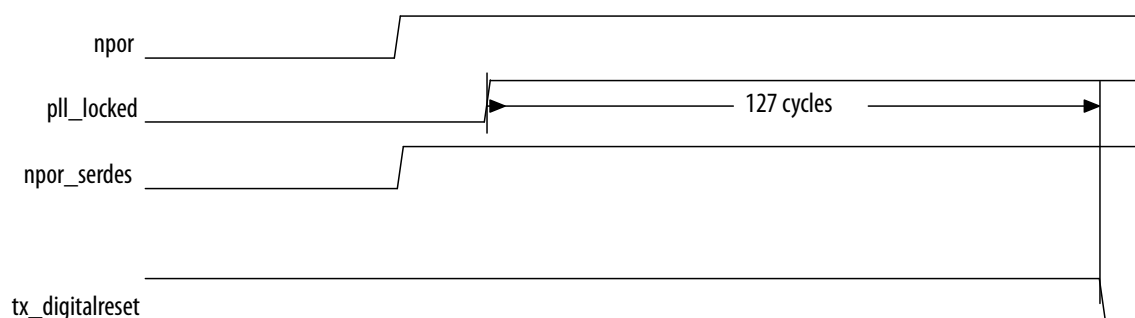
This reset sequence includes the following steps:

1. After `pin_perst` or `npwr` is released, the Hard IP reset controller waits for `pld_clk_inuse` to be asserted.
2. `crst` and `srst` are released 32 cycles after `pld_clk_inuse` is asserted.
3. The Hard IP for PCI Express deasserts the `reset_status` output to the Application Layer.
4. The `altpciied_<device>v_hwtcl.sv` deasserts `app_rstn` 32 `pld_clk` cycles after `reset_status` is released.

**Figure 6-3: RX Transceiver Reset Sequence**

The RX transceiver reset sequence includes the following steps:

1. After `rx_pll_locked` is asserted, the LTSSM state machine transitions from the Detect.Quiet to the Detect.Active state.
2. When the `pipe_phystatus` pulse is asserted and `pipe_rxstatus[2:0] = 3`, the receiver detect operation has completed.
3. The LTSSM state machine transitions from the Detect.Active state to the Polling.Active state.
4. The Hard IP for PCI Express asserts `rx_digitalreset`. The `rx_digitalreset` signal is deasserted after `rx_signaldetect` is stable for a minimum of 3 ms.

**Figure 6-4: TX Transceiver Reset Sequence**

The TX transceiver reset sequence includes the following steps:

1. After `npor` is deasserted, the IP core deasserts the `npor_serdes` input to the TX transceiver.
2. The SERDES reset controller waits for `pll_locked` to be stable for a minimum of 127 `pld_clk` cycles before deasserting `tx_digitalreset`.

For descriptions of the available reset signals, refer to *Reset Signals, Status, and Link Training Signals*.

## Clocks

The Hard IP contains a clock domain crossing (CDC) synchronizer at the interface between the PHY/MAC and the DLL layers. The synchronizer allows the Data Link and Transaction Layers to run at frequencies independent of the PHY/MAC. The CDC synchronizer provides more flexibility for the user clock interface. Depending on parameters you specify, the core selects the appropriate `coreclkout_hip`. You can use these parameters to enhance performance by running at a higher frequency for latency optimization or at a lower frequency to save power.

In accordance with the *PCI Express Base Specification*, you must provide a 100 MHz reference clock that is connected directly to the transceiver.

As a convenience, you may also use a 125 MHz input reference clock as input to the TX PLL.

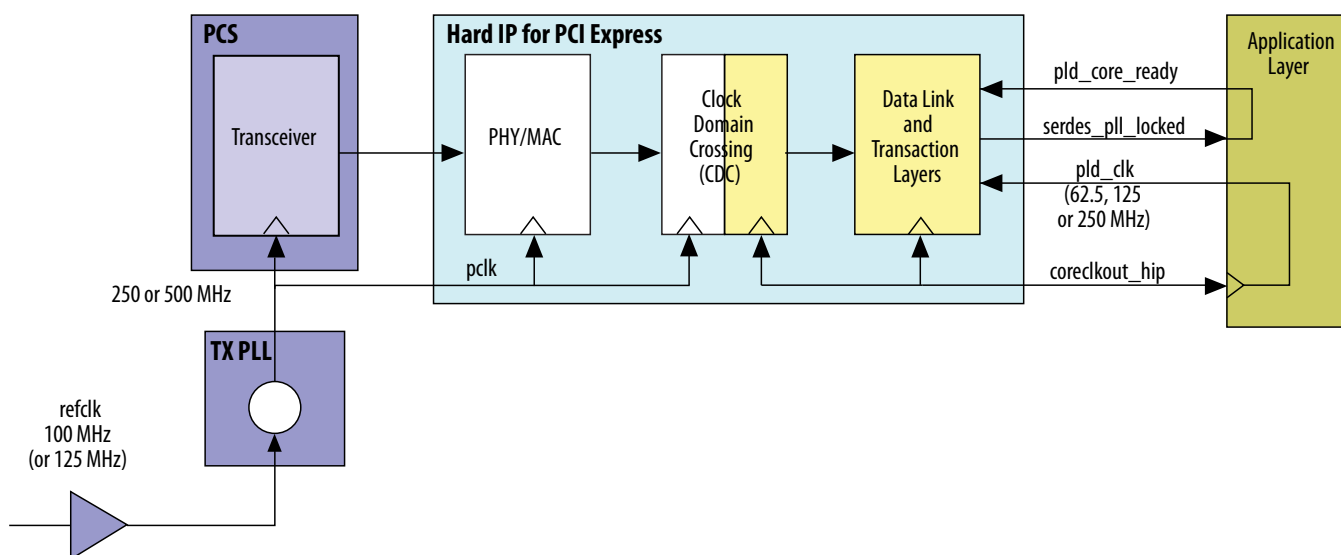
### Related Information

[PCI Express Base Specification 2.1 or 3.0](#)

## Clock Domains

**Figure 6-5: Clock Domains and Clock Generation for the Application Layer**

The following illustrates the clock domains when using `coreclkout_hip` to drive the Application Layer and the `p1d_clk` of the IP core. The Intel-provided example design connects `coreclkout_hip` to the `p1d_clk`. However, this connection is not mandatory. Inside the Hard IP for PCI Express, the blocks shown in white are in the `pclk` domain, while the blocks shown in yellow are in the `coreclkout_hip` domain.



As this figure indicates, the IP core includes the following clock domains: `pclk`, `coreclkout_hip` and `p1d_clk`.

## pclk

The transceiver derives pclk from the 100 MHz refclk signal that you must provide to the device.

The *PCI Express Base Specification* requires that the refclk signal frequency be 100 MHz  $\pm$ 300 PPM.

The transitions between Gen1, Gen2, and Gen3 should be glitchless. pclk can be turned off for most of the 1 ms timeout assigned for the PHY to change the clock rate; however, pclk should be stable before the 1 ms timeout expires.

The transitions between Gen1 and Gen2 should be glitchless. pclk can be turned off for most of the 1 ms timeout assigned for the PHY to change the clock rate; however, pclk should be stable before the 1 ms timeout expires.

**Table 6-2: pclk Clock Frequency**

Data Rate	Frequency
Gen1	250 MHz
Gen2	500 MHz

The CDC module implements the asynchronous clock domain crossing between the PHY/MAC pclk domain and the Data Link Layer coreclk domain. The transceiver pclk clock is connected directly to the Hard IP for PCI Express and does not connect to the FPGA fabric.

### Related Information

[PCI Express Base Specification 2.1 or 3.0](#)

## coreclkout\_hip

**Table 6-3: Application Layer Clock Frequency for All Combinations of Link Width, Data Rate and Application Layer Interface Widths**

The coreclkout\_hip signal is derived from pclk. The following table lists frequencies for coreclkout\_hip, which are a function of the link width, data rate, and the width of the Application Layer to Transaction Layer interface. The frequencies and widths specified in this table are maintained throughout operation. If the link downtrains to a lesser link width or changes to a different maximum link rate, it maintains the frequencies it was originally configured for as specified in this table. (The Hard IP throttles the interface to achieve a lower throughput.)

Link Width	Max Link Rate	Avalon Interface Width	coreclkout_hip
×8	Gen1	128	125 MHz
×4	Gen2	128	125 MHz
×8	Gen2	128	250 MHz
×8	Gen2	256	125 MHz

Link Width	Max Link Rate	Avalon Interface Width	coreclkout_hip
×2	Gen3	128	125 MHz
×4	Gen3	128	250 MHz
×4	Gen3	256	125 MHz
×8	Gen3	256	250 MHz

## p1d\_clk

coreclkout\_hip can drive the Application Layer clock along with the p1d\_clk input to the IP core. The p1d\_clk can optionally be sourced by a different clock than coreclkout\_hip. The p1d\_clk minimum frequency cannot be lower than the coreclkout\_hip frequency. Based on specific Application Layer constraints, a PLL can be used to derive the desired frequency.

**Note:** For Gen3, Intel recommends using a common reference clock (0 ppm) because when using separate reference clocks (non 0 ppm), the PCS occasionally must insert SKP symbols, potentially causing the PCIe link to go to recovery. Gen1 or Gen2 modes are not affected by this issue. Systems using the common reference clock (0 ppm) are not affected by this issue. The primary repercussion of this issue is a slight decrease in bandwidth. On Gen3 x8 systems, this bandwidth impact is negligible. If non 0 ppm mode is required, so that separate reference clocks are used, please contact Intel for further information and guidance.

## Clock Summary

Table 6-4: Clock Summary

Name	Frequency	Clock Domain
coreclkout_hip	62.5, 125 or 250 MHz	Avalon-ST interface between the Transaction and Application Layers.
p1d_clk	125 or 250 MHz	Application and Transaction Layers.
refclk	100 or 125 MHz	SERDES (transceiver). Dedicated free running input clock to the SERDES block.
reconfig_xcvr_clk	100 –125 MHz	Transceiver Reconfiguration Controller.



2018.07.31

UG-01154



Subscribe



Send Feedback

Each PCI Express compliant device must implement a basic level of error management and can optionally implement advanced error management. The IP core implements both basic and advanced error reporting. Error handling for a Root Port is more complex than that of an Endpoint.

**Table 7-1: Error Classification**

The *PCI Express Base Specification* defines three types of errors, outlined in the following table.

Type	Responsible Agent	Description
Correctable	Hardware	While correctable errors may affect system performance, data integrity is maintained.
Uncorrectable, non-fatal	Device software	Uncorrectable, non-fatal errors are defined as errors in which data is lost, but system integrity is maintained. For example, the fabric may lose a particular TLP, but it still works without problems.
Uncorrectable, fatal	System software	Errors generated by a loss of data and system failure are considered uncorrectable and fatal. Software must determine how to handle such errors: whether to reset the link or implement other means to minimize the problem.

## Related Information

[PCI Express Base Specification 2.1 and 3.0](#)

## Physical Layer Errors

**Table 7-2: Errors Detected by the Physical Layer**

The following table describes errors detected by the Physical Layer. Physical Layer error reporting is optional in the *PCI Express Base Specification*.

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

**ALTERA**  
now part of Intel

Error	Type	Description
Receive port error	Correctable	<p>This error has the following 3 potential causes:</p> <ul style="list-style-type: none"> <li>Physical coding sublayer error when a lane is in L0 state. These errors are reported to the Hard IP block via the per lane PIPE interface input receive status signals, <code>rxstatus&lt;lane_number&gt;[2:0]</code> using the following encodings: <ul style="list-style-type: none"> <li>3'b100: 8B/10B Decode Error</li> <li>3'b101: Elastic Buffer Overflow</li> <li>3'b110: Elastic Buffer Underflow</li> <li>3'b111: Disparity Error</li> </ul> </li> <li>Deskew error caused by overflow of the multilane deskew FIFO.</li> <li>Control symbol received in wrong lane.</li> </ul>

## Data Link Layer Errors

Table 7-3: Errors Detected by the Data Link Layer

Error	Type	Description
Bad TLP	Correctable	This error occurs when a LCRC verification fails or when a sequence number error occurs.
Bad DLLP	Correctable	This error occurs when a CRC verification fails.
Replay timer	Correctable	This error occurs when the replay timer times out.
Replay num rollover	Correctable	This error occurs when the replay number rolls over.
Data Link Layer protocol	Uncorrectable(fatal)	This error occurs when a sequence number specified by the Ack/Nak block in the Data Link Layer ( <code>AckNak_Seq_Num</code> ) does not correspond to an unacknowledged TLP.

## Transaction Layer Errors

**Table 7-4: Errors Detected by the Transaction Layer**

Error	Type	Description
Poisoned TLP received	Uncorrectable (non-fatal)	<p>This error occurs if a received Transaction Layer Packet has the EP poison bit set.</p> <p>The received TLP is passed to the Application Layer and the Application Layer logic must take appropriate action in response to the poisoned TLP. Refer to “2.7.2.2 Rules for Use of Data Poisoning” in the <i>PCI Express Base Specification</i> for more information about poisoned TLPs.</p>
ECRC check failed <sup>(1)</sup>	Uncorrectable (non-fatal)	<p>This error is caused by an ECRC check failing despite the fact that the TLP is not malformed and the LCRC check is valid.</p> <p>The Hard IP block handles this TLP automatically. If the TLP is a non-posted request, the Hard IP block generates a completion with completer abort status. In all cases the TLP is deleted in the Hard IP block and not presented to the Application Layer.</p>
Unsupported Request for Endpoints	Uncorrectable (non-fatal)	<p>This error occurs whenever a component receives any of the following Unsupported Requests:</p> <ul style="list-style-type: none"><li>• Type 0 Configuration Requests for a non-existing function.</li><li>• Completion transaction for which the Requester ID does not match the bus, device and function number.</li><li>• Unsupported message.</li><li>• A Type 1 Configuration Request TLP for the TLP from the PCIe link.</li><li>• A locked memory read (MEMRDLK) on native Endpoint.</li><li>• A locked completion transaction.</li><li>• A 64-bit memory transaction in which the 32 MSBs of an address are set to 0.</li><li>• A memory or I/O transaction for which there is no BAR match.</li><li>• A memory transaction when the Memory Space Enable bit (bit [1] of the PCI Command register at Configuration Space offset 0x4) is set to 0.</li><li>• A poisoned configuration write request (CfgWr0)</li></ul>

Error	Type	Description
		In all cases the TLP is deleted in the Hard IP block and not presented to the Application Layer. If the TLP is a non-posted request, the Hard IP block generates a completion with Unsupported Request status.
Completion timeout	Uncorrectable (non-fatal)	This error occurs when a request originating from the Application Layer does not generate a corresponding completion TLP within the established time. It is the responsibility of the Application Layer logic to provide the completion timeout mechanism. The completion timeout should be reported from the Transaction Layer using the <code>cpl_err[0]</code> signal.
Completer abort <sup>(1)</sup>	Uncorrectable (non-fatal)	The Application Layer reports this error using the <code>cpl_err[2]</code> signal when it aborts receipt of a TLP.
Unexpected completion	Uncorrectable (non-fatal)	<p>This error is caused by an unexpected completion transaction. The Hard IP block handles the following conditions:</p> <ul style="list-style-type: none"> <li>• The Requester ID in the completion packet does not match the Configured ID of the Endpoint.</li> <li>• The completion packet has an invalid tag number. (Typically, the tag used in the completion packet exceeds the number of tags specified.)</li> <li>• The completion packet has a tag that does not match an outstanding request.</li> <li>• The completion packet for a request that was to I/O or Configuration Space has a length greater than 1 dword.</li> <li>• The completion status is Configuration Retry Status (CRS) in response to a request that was not to Configuration Space.</li> </ul> <p>In all of the above cases, the TLP is not presented to the Application Layer; the Hard IP block deletes it.</p> <p>The Application Layer can detect and report other unexpected completion conditions using the <code>cpl_err[2]</code> signal. For example, the Application Layer can report cases where the total length of the received successful completions do not match the original read request length.</p>

Error	Type	Description
Receiver overflow <sup>(1)</sup>	Uncorrectable (fatal)	This error occurs when a component receives a TLP that violates the FC credits allocated for this type of TLP. In all cases the hard IP block deletes the TLP and it is not presented to the Application Layer.
Flow control protocol error (FCPE) <sup>(1)</sup>	Uncorrectable (fatal)	This error occurs when a component does not receive update flow control credits with the 200 $\mu$ s limit.
Malformed TLP	Uncorrectable (fatal)	<p>This error is caused by any of the following conditions:</p> <ul style="list-style-type: none"><li>• The data payload of a received TLP exceeds the maximum payload size.</li><li>• The <math>\text{TD}</math> field is asserted but no TLP digest exists, or a TLP digest exists but the <math>\text{TD}</math> bit of the PCI Express request header packet is not asserted.</li><li>• A TLP violates a byte enable rule. The Hard IP block checks for this violation, which is considered optional by the PCI Express specifications.</li><li>• A TLP in which the <code>type</code> and <code>length</code> fields do not correspond with the total length of the TLP.</li><li>• A TLP in which the combination of format and type is not specified by the PCI Express specification.</li><li>• A request specifies an address/length combination that causes a memory space access to exceed a 4 KB boundary. The Hard IP block checks for this violation, which is considered optional by the PCI Express specification.</li><li>• Messages, such as Assert_INTX, Power Management, Error Signaling, Unlock, and Set Power Slot Limit, must be transmitted across the default traffic class.</li></ul> <p>The Hard IP block deletes the malformed TLP; it is not presented to the Application Layer.</p>

Note:

1. Considered optional by the *PCI Express Base Specification Revision*.

## Error Reporting and Data Poisoning

How the Endpoint handles a particular error depends on the configuration registers of the device.

Refer to the *PCI Express Base Specification 3.0* for a description of the device signaling and logging for an Endpoint.

The Hard IP block implements data poisoning, a mechanism for indicating that the data associated with a transaction is corrupted. Poisoned TLPs have the error/poisoned bit of the header set to 1 and observe the following rules:

- Received poisoned TLPs are sent to the Application Layer and status bits are automatically updated in the Configuration Space.
- Received poisoned Configuration Write TLPs are not written in the Configuration Space.
- The Configuration Space never generates a poisoned TLP; the error/poisoned bit of the header is always set to 0.

Poisoned TLPs can also set the parity error bits in the PCI Configuration Space Status register.

**Table 7-5: Parity Error Conditions**

Status Bit	Conditions
Detected parity error (status register bit 15)	Set when any received TLP is poisoned.
Master data parity error (status register bit 8)	<p>This bit is set when the command register parity enable bit is set and one of the following conditions is true:</p> <ul style="list-style-type: none"><li>• The poisoned bit is set during the transmission of a Write Request TLP.</li><li>• The poisoned bit is set on a received completion TLP.</li></ul>

Poisoned packets received by the Hard IP block are passed to the Application Layer. Poisoned transmit TLPs are similarly sent to the link.

#### Related Information

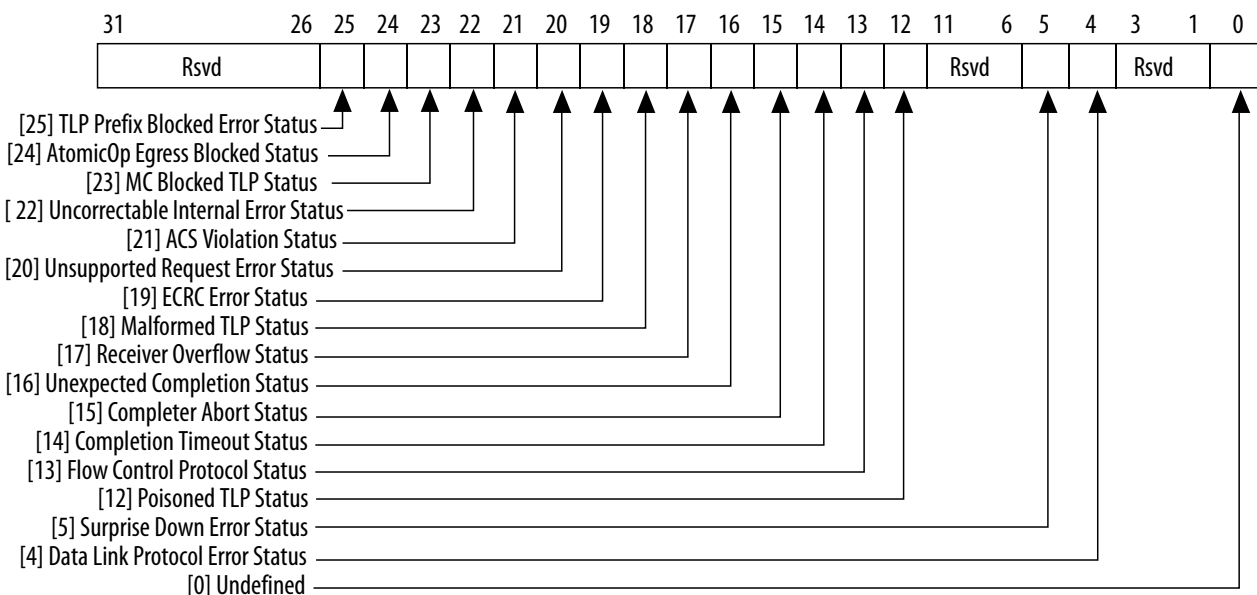
[PCI Express Base Specification 2.1 and 3.0](#)

## Uncorrectable and Correctable Error Status Bits

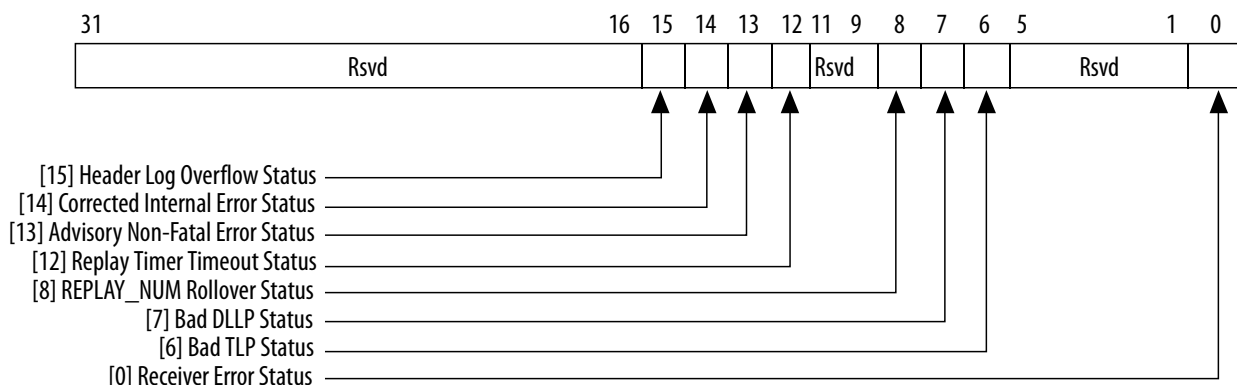
The following section is reprinted with the permission of PCI-SIG. Copyright 2010 PCI-SIG.

**Figure 7-1: Uncorrectable Error Status Register**

The default value of all the bits of this register is 0. An error status bit that is set indicates that the error condition it represents has been detected. Software may clear the error status by writing a 1 to the appropriate bit.

**Figure 7-2: Correctable Error Status Register**

The default value of all the bits of this register is 0. An error status bit that is set indicates that the error condition it represents has been detected. Software may clear the error status by writing a 1 to the appropriate bit.



2018.07.31 2018.07.31

UG-01154 UG-01154



Subscribe



Send Feedback

The V-Series Avalon-MM Hard IP for PCI Express implements the complete PCI Express protocol stack as defined in the *PCI Express Base Specification*. The protocol stack includes the following layers:

- *Transaction Layer*—The Transaction Layer contains the Configuration Space, which manages communication with the Application Layer, the RX and TX channels, the RX buffer, and flow control credits.
- *Data Link Layer*—The Data Link Layer, located between the Physical Layer and the Transaction Layer, manages packet transmission and maintains data integrity at the link level. Specifically, the Data Link Layer performs the following tasks:
  - Manages transmission and reception of Data Link Layer Packets (DLLPs)
  - Generates all transmission cyclical redundancy code (CRC) values and checks all CRCs during reception
  - Manages the retry buffer and retry mechanism according to received ACK/NAK Data Link Layer packets
  - Initializes the flow control mechanism for DLLPs and routes flow control credits to and from the Transaction Layer
- *Physical Layer*—The Physical Layer initializes the speed, lane numbering, and lane width of the PCI Express link according to packets received from the link and directives received from higher layers.

The following figure provides a high-level block diagram.

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered



Figure 8 9-1: V-Series Avalon-MM DMA for PCI Express

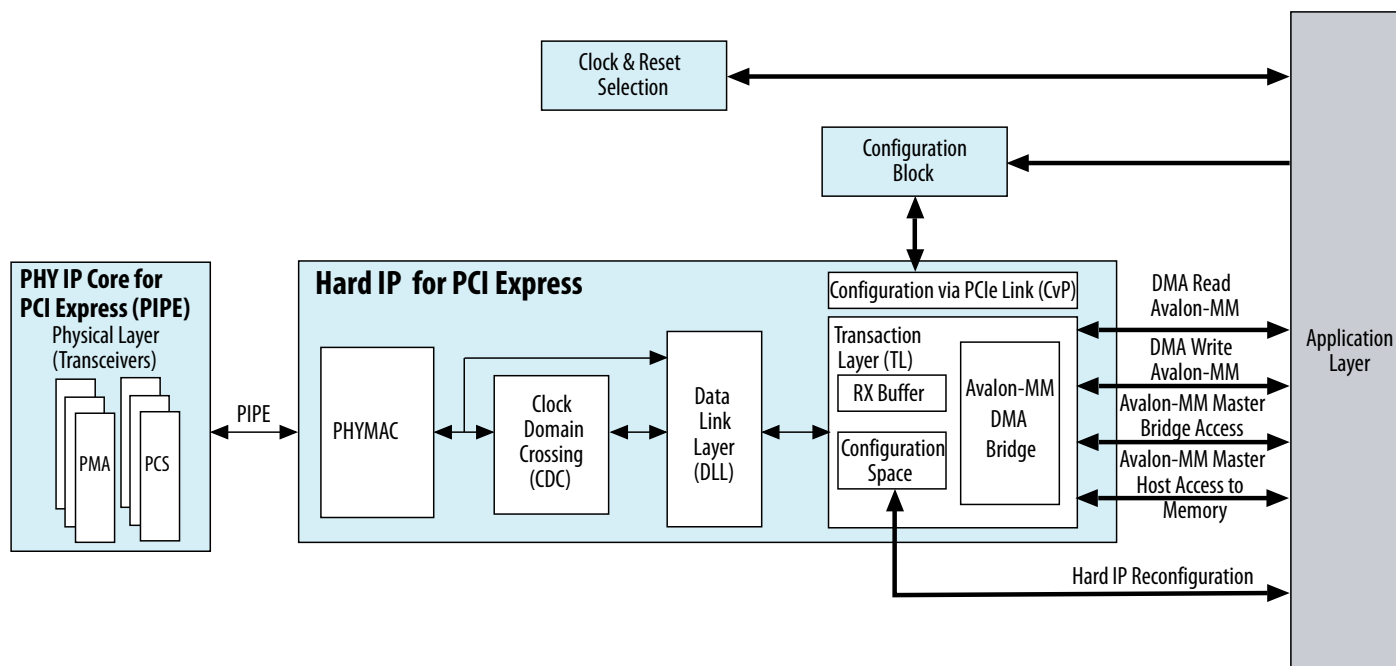


Table 8 9-1: Application Layer Clock Frequencies

Lanes	Gen1	Gen2	Gen3
×2	N/A	125 MHz @ 128 bits	125 MHz @ 128 bits
×4	N/A	125 MHz @ 128 bits	250 MHz @ 128 bits or 125 MHz @ 256 bits
×8	125 MHz @ 128 bits	250 MHz @ 128 bits or 125 MHz @ 256 bits	250 MHz @ 256 bits

**Related Information**

[PCI Express Base Specification 2.1 or 3.0](#)

## Top-Level Interfaces

### Avalon-MM DMA Interface

An Avalon-MM interface with DMA connects the Application Layer and the Transaction Layer. This interface includes high-performance, burst capable Read DMA and Write DMA modules. This variant is available for the following Endpoint configurations:

- Gen1 x8
- Gen2 x4, x8
- Gen3 x2, x4, x8

#### Related Information

[V-Series DMA Avalon-MM DMA Interface to the Application Layer](#) on page 4-1

### Clocks and Reset

The *PCI Express Base Specification* requires an input reference clock, which is called `refclk` in this design. The *PCI Express Base Specification* stipulates that the frequency of this clock be 100 MHz.

The *PCI Express Base Specification* also requires a system configuration time of 100 ms. To meet this specification, IP core includes an embedded hard reset controller. This reset controller exits the reset state after the periphery of the device is initialized.

### Transceiver Reconfiguration

The transceiver reconfiguration interface allows you to dynamically reconfigure the values of analog settings in the PMA block of the transceiver. Dynamic reconfiguration is necessary to compensate for process variations.

#### Related Information

[Transceiver PHY IP Reconfiguration](#) on page 10-1

### Interrupts

The Hard IP for PCI Express offers the following interrupt mechanisms:

- Message Signaled Interrupts (MSI)— MSI uses the TLP single dword memory writes to implement interrupts. This interrupt mechanism conserves pins because it does not use separate wires for interrupts. In addition, the single dword provides flexibility in data presented in the interrupt message. The MSI Capability structure is stored in the Configuration Space and is programmed using Configuration Space accesses.
- MSI-X—The Transaction Layer generates MSI-X messages which are single dword memory writes. The MSI-X Capability structure points to an MSI-X table structure and MSI-X PBA structure which are stored in memory. This scheme is in contrast to the MSI capability structure, which contains all of the control and status information for the interrupt vectors.

#### Related Information

[MSI Interrupts for Endpoints](#) on page 4-25

## PIPE

The PIPE interface implements the Intel-designed PIPE interface specification. You can use this parallel interface to speed simulation; however, you cannot use the PIPE interface in actual hardware.

- The simulation models support PIPE and serial simulation.

### Related Information

[PIPE Interface Signals](#) on page 4-38

## Data Link Layer

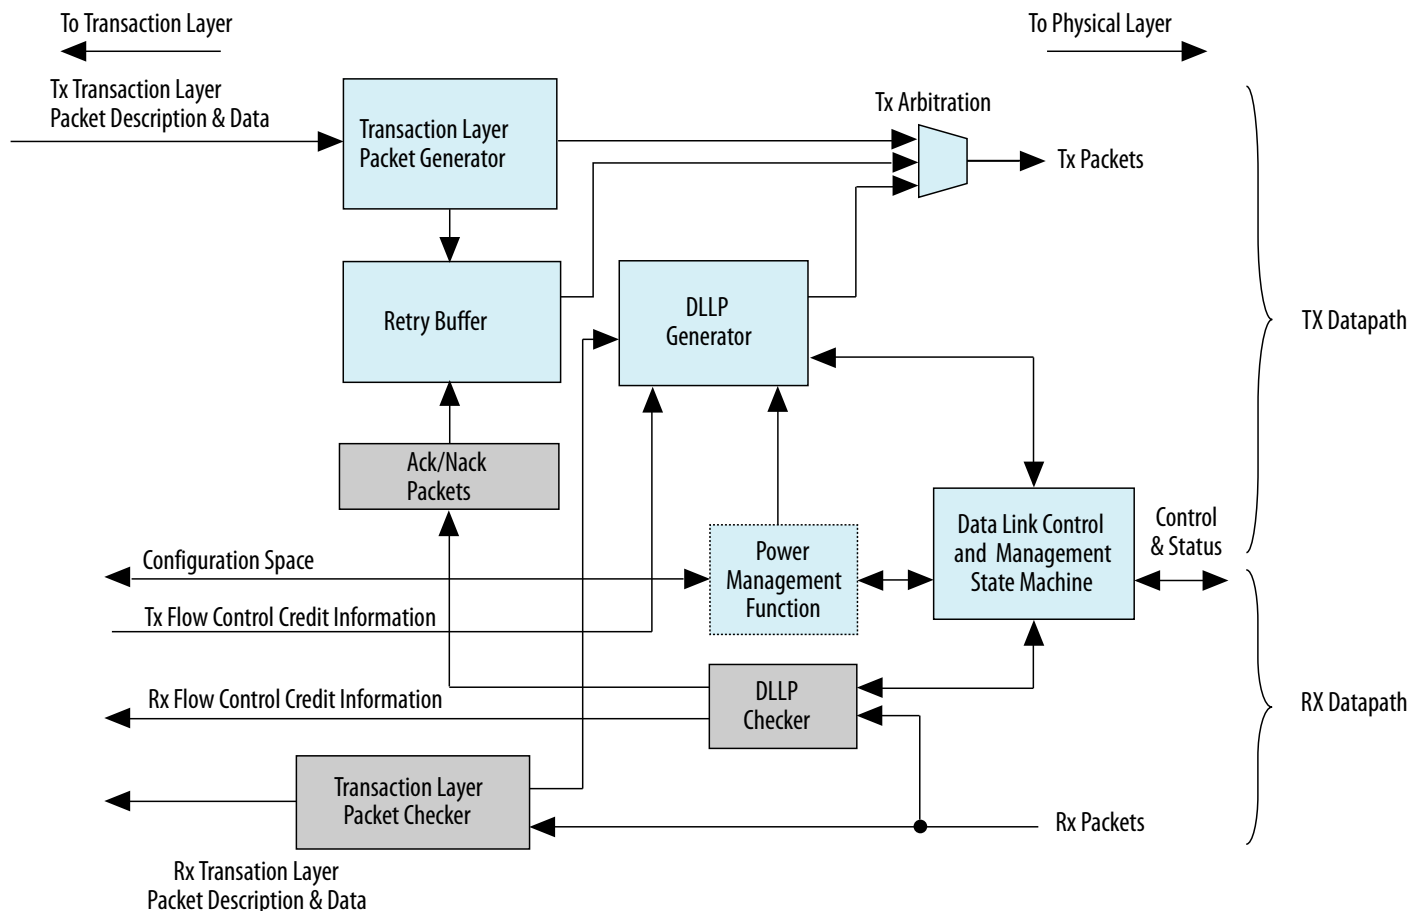
The Data Link Layer is located between the Transaction Layer and the Physical Layer. It maintains packet integrity and communicates (by DLL packet transmission) at the PCI Express link level.

The DLL implements the following functions:

- Link management through the reception and transmission of DLL Packets (DLLP), which are used for the following functions:
  - Power management of DLLP reception and transmission
  - To transmit and receive `ACK/NAK` packets
  - Data integrity through generation and checking of CRCs for TLPs and DLLPs
  - TLP retransmission in case of `NAK` DLLP reception or replay timeout, using the retry (replay) buffer
  - Management of the retry buffer
  - Link retraining requests in case of error through the Link Training and Status State Machine (LTSSM) of the Physical Layer



Figure 8 9-2: Data Link Layer



The DLL has the following sub-blocks:

- **Data Link Control and Management State Machine**—This state machine connects to both the Physical Layer's LTSSM state machine and the Transaction Layer. It initializes the link and flow control credits and reports status to the Transaction Layer.
- **Power Management**—This function handles the handshake to enter low power mode. Such a transition is based on register values in the Configuration Space and received Power Management (PM) DLLPs. All of the V-Series Hard IP for PCIe IP core variants do not support low power modes.
- **Data Link Layer Packet Generator and Checker**—This block is associated with the DLLP's 16-bit CRC and maintains the integrity of transmitted packets.
- **Transaction Layer Packet Generator**—This block generates transmit packets, including a sequence number and a 32-bit Link CRC (LCRC). The packets are also sent to the retry buffer for internal storage. In retry mode, the TLP generator receives the packets from the retry buffer and generates the CRC for the transmit packet.
- **Retry Buffer**—The retry buffer stores TLPs and retransmits all unacknowledged packets in the case of NAK DLLP reception. In case of ACK DLLP reception, the retry buffer discards all acknowledged packets.

- ACK/NAK Packets—The ACK/NAK block handles ACK/NAK DLLPs and generates the sequence number of transmitted packets.
- Transaction Layer Packet Checker—This block checks the integrity of the received TLP and generates a request for transmission of an ACK/NAK DLLP.
- TX Arbitration—This block arbitrates transactions, prioritizing in the following order:
  - Initialize FC Data Link Layer packet
  - ACK/NAK DLLP (high priority)
  - Update FC DLLP (high priority)
  - PM DLLP
  - Retry buffer TLP
  - TLP
  - Update FC DLLP (low priority)
  - ACK/NAK FC DLLP (low priority)

## Physical Layer

The Physical Layer is the lowest level of the PCI Express protocol stack. It is the layer closest to the serial link. It encodes and transmits packets across a link and accepts and decodes received packets. The Physical Layer connects to the link through a high-speed SERDES interface running at 2.5 Gbps for Gen1 implementations, at 2.5 or 5.0 Gbps for Gen2 implementations, and at 2.5, 5.0 or 8.0 Gbps for Gen3 implementations.

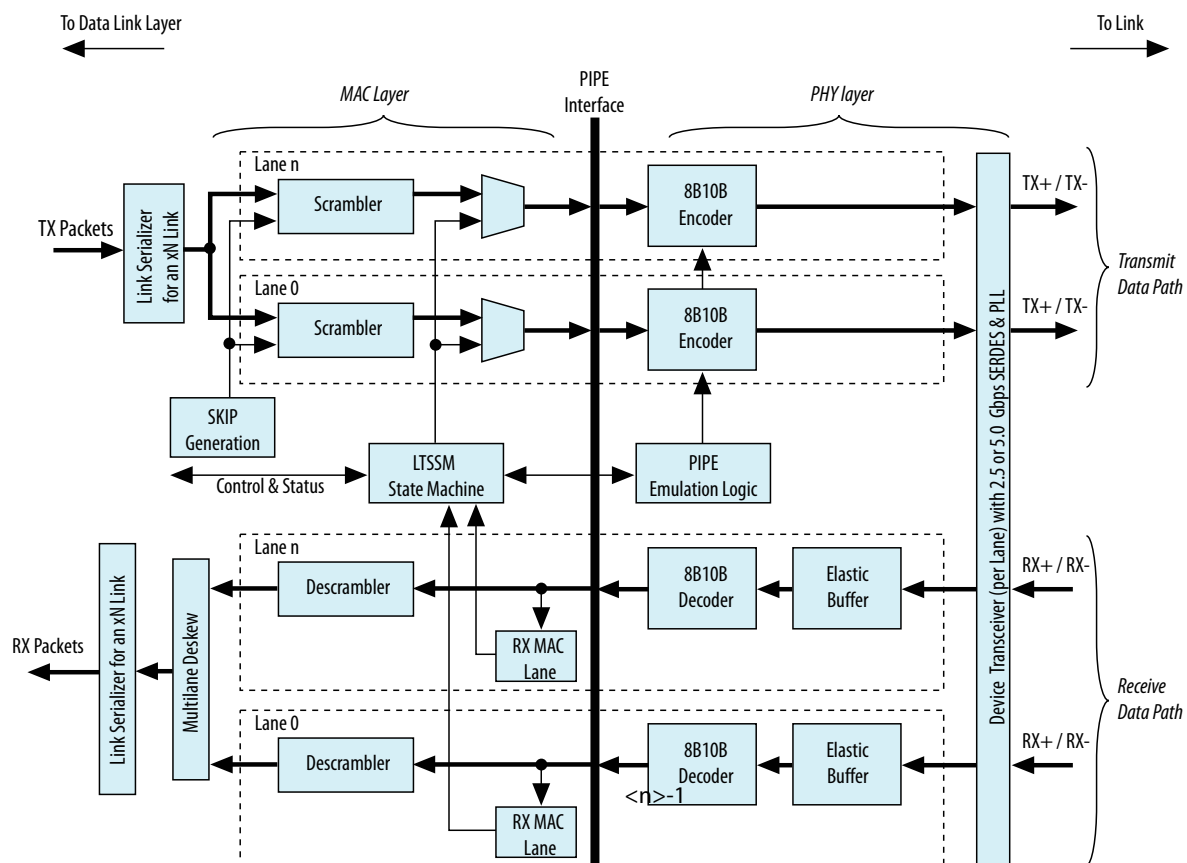
The Physical Layer is the lowest level of the PCI Express protocol stack. It is the layer closest to the serial link. It encodes and transmits packets across a link and accepts and decodes received packets. The Physical Layer connects to the link through a high-speed SERDES interface running at 2.5 Gbps for Gen1 implementations and at 2.5 or 5.0 Gbps for Gen2 implementations.

The Physical Layer is responsible for the following actions:

- Training the link
- Scrambling/descrambling and 8B/10B encoding/decoding for 2.5 Gbps (Gen1), 5.0 Gbps (Gen2), or 128b/130b encoding/decoding of 8.0 Gbps (Gen3) per lane
- Scrambling/descrambling and 8B/10B encoding/decoding for 2.5 Gbps (Gen1) and 5.0 Gbps (Gen2) per lane
- Serializing and deserializing data
- Equalization (Gen3)
- Operating the PIPE 3.0 Interface
- Implementing auto speed negotiation (Gen2 and Gen3)
- Implementing auto speed negotiation (Gen2)
- Transmitting and decoding the training sequence
- Providing hardware autonomous speed control
- Implementing auto lane reversal



Figure 8 9-3: Physical Layer Architecture



PHY Layer—The PHY layer includes the 8B/10B encode and decode functions for Gen1 and Gen2. The PHY also includes elastic buffering and serialization/deserialization functions.

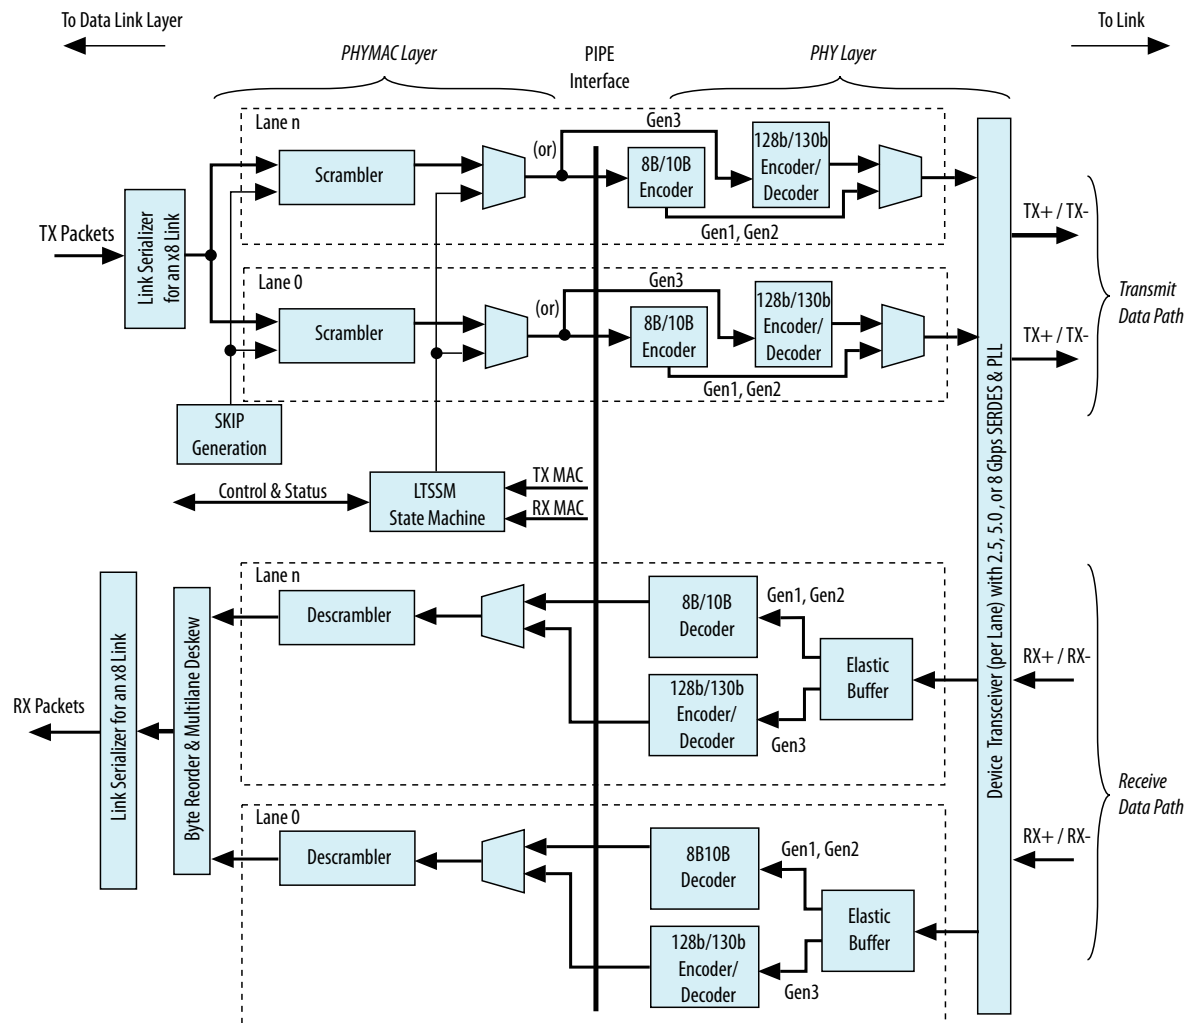
The Physical Layer is subdivided by the PIPE Interface Specification into two layers (bracketed horizontally in above figure):

- PHYMAC—The MAC layer includes the LTSSM and the scrambling/descrambling, byte reordering, and multilane deskew functions.
- Media Access Controller (MAC) Layer—The MAC layer includes the LTSSM and the scrambling and descrambling and multilane deskew functions.
- PHY Layer—The PHY layer includes the 8B/10B encode and decode functions for Gen1 and Gen2. It includes 128b/130b encode and decode functions for Gen3. The PHY also includes elastic buffering and serialization/deserialization functions.
- PHY Layer—The PHY layer includes the 8B/10B encode and decode functions for Gen1 and Gen2. The PHY also includes elastic buffering and serialization/deserialization functions.

The Physical Layer integrates both digital and analog elements. Intel designed the PIPE interface to separate the PHYMAC from the PHY. The V-Series Hard IP for PCI Express complies with the PIPE interface specification.

**Note:** The internal PIPE interface is visible for simulation. It is not available for debugging in hardware using a logic analyzer such as Signal Tap. If you try to connect Signal Tap to this interface the design fails compilation.

**Figure 8 9-4: Physical Layer Architecture**



The PHYMAC block comprises four main sub-blocks:

- MAC Lane—Both the RX and the TX path use this block.
  - On the RX side, the block decodes the Physical Layer packet and reports to the LTSSM the type and number of TS1/TS2 ordered sets received.
  - On the TX side, the block multiplexes data from the DLL and the Ordered Set and SKP sub-block (LTSTX). It also adds lane specific information, including the lane number and the force PAD value when the LTSSM disables the lane during initialization.
- LTSSM—This block implements the LTSSM and logic that tracks TX and RX training sequences on each lane.
- For transmission, it interacts with each MAC lane sub-block and with the LTSTX sub-block by asserting both global and per-lane control bits to generate specific Physical Layer packets.
  - On the receive path, it receives the Physical Layer packets reported by each MAC lane sub-block. It also enables the multilane deskew block. This block reports the Physical Layer status to higher layers.
  - LTSTX (Ordered Set and SKP Generation)—This sub-block generates the Physical Layer packet. It receives control signals from the LTSSM block and generates Physical Layer packet for each lane. It generates the same Physical Layer Packet for all lanes and PAD symbols for the link or lane number in the corresponding TS1/TS2 fields. The block also handles the receiver detection operation to the PCS sub-layer by asserting predefined PIPE signals and waiting for the result. It also generates a SKP Ordered Set at every predefined timeslot and interacts with the TX alignment block to prevent the insertion of a SKP Ordered Set in the middle of packet.
  - Deskew—This sub-block performs the multilane deskew function and the RX alignment between the initialized lanes and the datapath. The multilane deskew implements an eight-word FIFO buffer for each lane to store symbols. Each symbol includes eight data bits, one disparity bit, and one control bit. The FIFO discards the FTS, COM, and SKP symbols and replaces PAD and IDL with D0.0 data. When all eight FIFOs contain data, a read can occur. When the multilane lane deskew block is first enabled, each FIFO begins writing after the first COM is detected. If all lanes have not detected a COM symbol after seven clock cycles, they are reset and the resynchronization process restarts, or else the RX alignment function recreates a 64-bit data word which is sent to the DLL.



# V-Series Avalon-MM DMA for PCI Express

2018.07.31

UG-01154



Subscribe



Send Feedback

The V-Series Avalon-MM DMA for PCI Express IP Core includes highly efficient Read DMA, Write DMA, and DMA Descriptor Controller modules. The typical throughput for hardware systems using a 256-bit interface to the Application Layer is 6 GB/sec or higher. For hardware systems using a 128-bit interface to the Application Layer the throughput scales proportionately to 3 GB/sec.

Using a 64-byte payload, the maximum theoretical throughput is far less due to the increased proportion of the bandwidth taken by the TLP headers. The throughput for back-to-back TX memory write completions, RX read completions, and simultaneous reads and writes is 2 GB/sec.

**Note:** A 64-byte packet is the minimum packet size for Ethernet.

## Related Information

- [Getting Started with the Avalon-MM DMA](#) on page 2-1
- [DMA Descriptor Controller Registers](#) on page 5-15

## Understanding the Internal DMA Descriptor Controller

When you select **Instantiate internal descriptor controller** in the parameter editor, the Avalon-MM with DMA includes an internal DMA Descriptor Controller to manage read and write DMA operations. The DMA Descriptor Controller includes read and write data movers to perform local memory reads and writes. It supports up to 128 descriptors for read and write DMAs. Host software programs the DMA Descriptor Controller internal registers with the location and size of the descriptor table residing in the PCI Express main memory. The descriptor control logic directs the DMA read logic to copy the entire table to its local FIFOs.

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

**ALTERA**  
now part of Intel

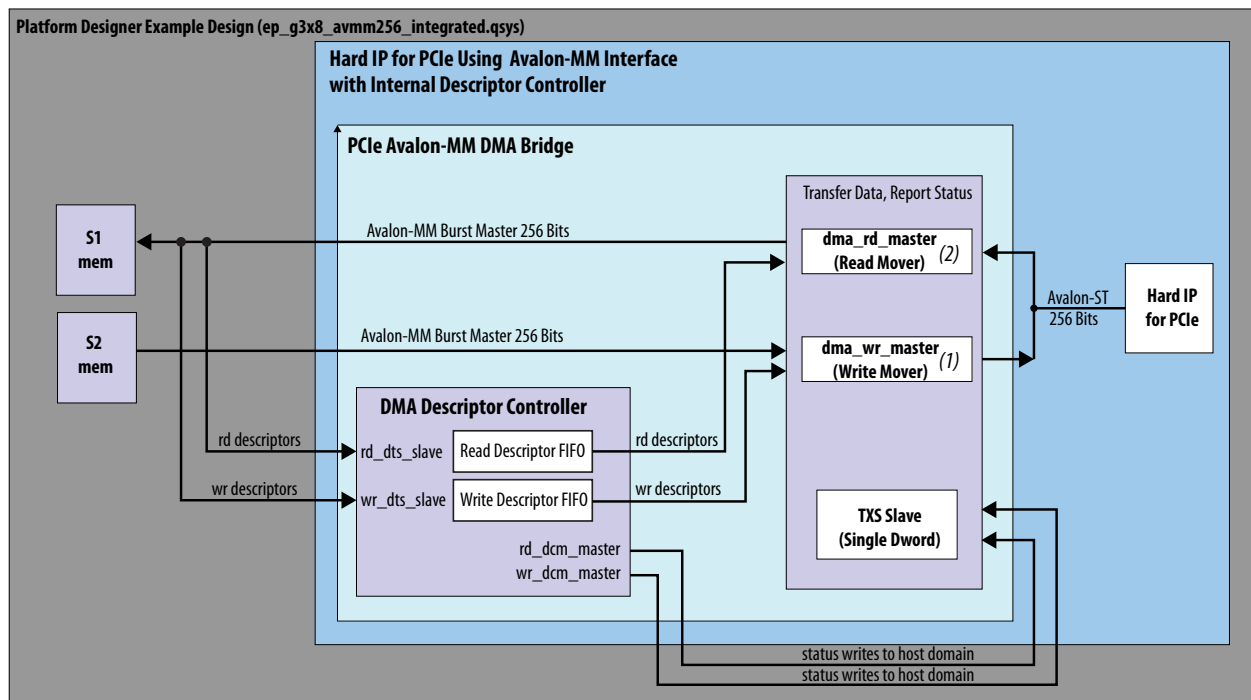
**Figure -1: Platform Designer Design Example with the Internal DMA Descriptor Controller**

This Platform Designer design example, **ep\_g3x8\_avmm256\_integrated.qsys**, is available in the `<install_dir>/ip/altera/altera_pcie/altera_pcie_a10_ed/example_design/a10` directory. Refer to *Getting Started with the V-Series Avalon-MM DMA* for instructions on simulating and compiling this example design. This screen capture filters out some interface types for clarity.

System Contents					
Address Map					
Interconnect Requirements					
System: ep_g3x8_avmm256_integrated Path: altpcie_devkit_0					
Use	Connections	Name	Description	Base	
<input checked="" type="checkbox"/>		<b>DUT</b>	Arria 10 Hard IP for PCI Express		
		txs	Avalon Memory Mapped Slave	0x0000_0000_0000_0000	
		rxm_bar4	Avalon Memory Mapped Master		
		dma_rd_master	Avalon Memory Mapped Master		
		dma_wr_master	Avalon Memory Mapped Master		
		rd_dts_slave	Avalon Memory Mapped Slave	0x0000_0000_0100_0000	
		wr_dts_slave	Avalon Memory Mapped Slave	0x0000_0000_0100_2000	
		rd_dcm_master	Avalon Memory Mapped Master		
		wr_dcm_master	Avalon Memory Mapped Master		
<input checked="" type="checkbox"/>		<b>onchip_memory2_0</b>	On-Chip Memory (RAM or ROM)		
		s1	Avalon Memory Mapped Slave	0x0000_0000_0000_0000	
		s2	Avalon Memory Mapped Slave	0x0000_0000_0000_0000	

**Figure -2: Avalon-MM DMA Block Diagram with the Internal DMA Descriptor Controller**

This block diagram corresponds to the Platform Designer system shown in the previous figure.

**Notes:**

- (1) Write Mover transfers data from local domain to the host domain.  
 (2) Read Mover transfers data from the host domain to local domain.

This design uses BAR0 and BAR1 to create a 64-bit address to access the DMA Descriptor Controller. These BARs cannot connect to any other interface. If BAR0 must access a different interface, you must use an external DMA descriptor controller. Intel recommends that you select the internal DMA Descriptor Controller if you do not plan to modify this component.

The high-performance BAR2 or BAR2 and BAR3 for 64-bit addresses is available to receive data for other high performance functions.

## Understanding the External DMA Descriptor Controller

Using the External DMA Descriptor Controller provides more flexibility. You can either modify or replace it to meet your system requirements. You may need to modify the DMA Descriptor Controller for the following reasons:

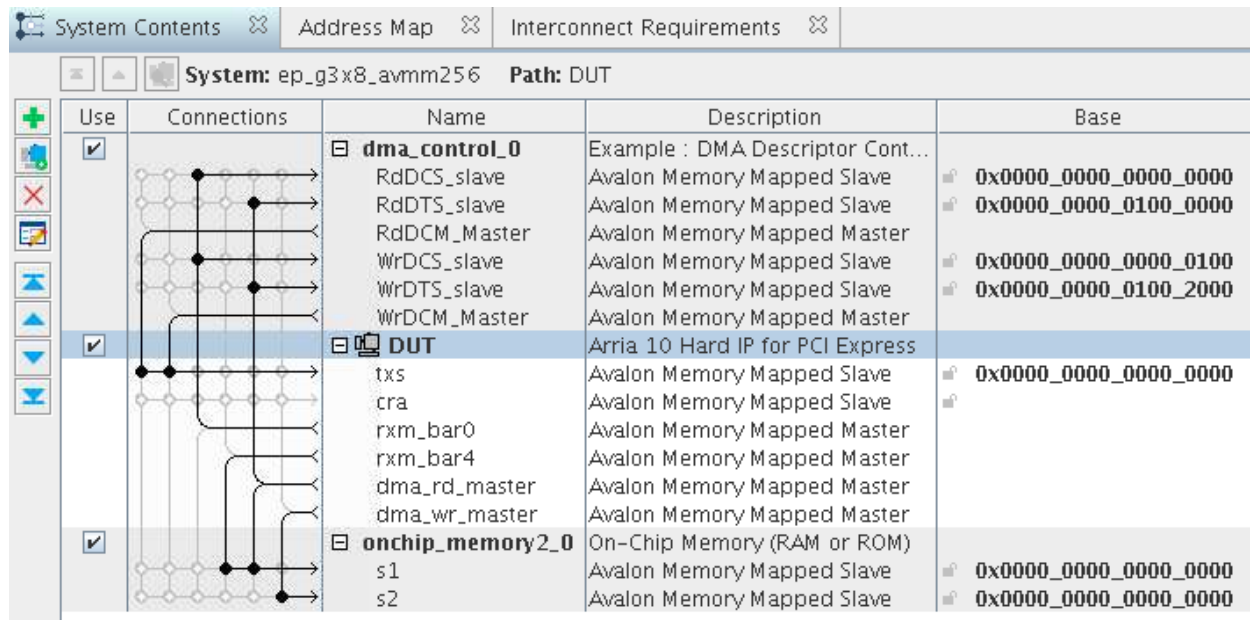
- To implement multi-channel operation
- To implement the descriptors as a linked list or to implement a custom DMA programming model
- To store descriptors in a local memory, instead of system (host-side) memory

To interface to the DMA logic included in this variant, the custom DMA descriptor controller must implement the following functions:

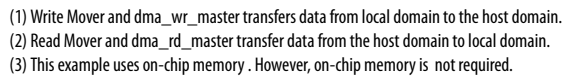
- It must communicate with the Write Mover and Read Mover to copy the descriptor table to local memory.
- The Write Mover and Read Mover must execute the descriptors stored in local memory.
- The DMA Avalon-MM write (WrDCM\_Master) and read (RdDCM\_Master) masters must be able to update status to the TX slave (TXS).

**Figure -3: Avalon-MM DMA Block Diagram with External DMA Descriptor Controller**

This Platform Designer design example, `ep_g3x8_avmm256.qsys`, is available in the `<install_dir>/ip/altera/altera_pcie/altera_pcie_a10_ed/example_design/a10` directory. This screen capture filters out some interface types for clarity.



This block diagram corresponds to the Platform Designer system shown in the previous figure. When the DMA Descriptor Controller is instantiated as a external component, it drives table entries on the `RdDmaRxData_i[159:0]` and `WrDmaRxData_i[159:0]` buses.



The DMA modules shown in the block diagram implements the following functionality:

- Read DMA (Read Mover and dma\_rd\_master) –Transfers data from the host domain to the local domain. It sends Memory Read TLPs upstream and writes the completion data to external Avalon-MM components using its own high performance master port. It follows the *PCI Express Base Specification* rules concerning tags, flow control credits, read completion boundary, maximum read size, and 4 KB boundaries.
- Write DMA (Write Mover and dma\_wr\_master) – Transfers data from the local domain to the host domain. It reads data from an Avalon-MM slave component using its own high performance master port. It sends data upstream using Memory Write TLPs. It follows the *PCI Express Base Specification* rules concerning tags, flow control credits, RX buffer completion rules, maximum payload size, and 4 KB boundaries.
- DMA Descriptor Controller–Manages read and write DMA operations. Host software programs its internal registers with the location of the descriptor table residing in the PCI Express system memory. The descriptor control logic directs the DMA read logic to copy the entire table to the local FIFO. It then fetches the table entries from the FIFO one at a time. It directs the appropriate DMA to transfer the data between the local and host domains. It also sends DMA status upstream via the TX slave single dword port (TXS). For more information about the DMA Descriptor Controller registers, refer to [DMA Descriptor Controller Registers](#) on page 5-15.
- RX Master (PCIe BAR0-1) –Allows the host to program internal registers of the DMA Descriptor Controller.
- TX Slave (TXS) – The DMA Descriptor Controller reports status on each read and write descriptor to this Avalon-MM slave. It also uses this port to send MSI requests.

# Transceiver PHY IP Reconfiguration 10

2018.07.31

UG-01154



Subscribe



Send Feedback

As silicon progresses towards smaller process nodes, circuit performance is affected by variations due to process, voltage, and temperature (PVT). Consequently, Gen3 designs require offset cancellation and adaptive equalization (AEQ) to ensure correct operation. Altera's Qsys example designs all include Transceiver Reconfiguration Controller and Altera PCIe Reconfig Driver IP cores that automatically perform these functions during the LTSSM equalization states.

As silicon progresses towards smaller process nodes, circuit performance is affected by variations due to process, voltage, and temperature (PVT). Designs typically require offset cancellation to ensure correct operation. At Gen2 data rates, designs also require DCD calibration. Altera's Qsys example designs all include Transceiver Reconfiguration Controller and Altera PCIe Reconfig Driver IP cores to perform these functions.

## Connecting the Transceiver Reconfiguration Controller IP Core

The Transceiver Reconfiguration Controller IP Core is available for V-series devices and can be found in the **Interface Protocols/Transceiver PHY** category in the IP Catalog. When you instantiate the Transceiver Reconfiguration Controller the **Enable offset cancellation block** and **Enable PLL calibration** options are enabled by default. For Gen3 variants, you should also turn on **Enable adaptive equalization (AEQ) block**.

A software driver for the Transceiver Reconfiguration Controller IP Core, Altera PCIe Reconfig Driver IP core, is also available in the IP Catalog under **Interface Protocols/PCIe**. The PCIe Reconfig Driver is implemented in clear text that you can modify if your design requires different reconfiguration functions.

**Note:** You must include a software driver in your design to program the Transceiver Reconfiguration Controller IP Core.

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

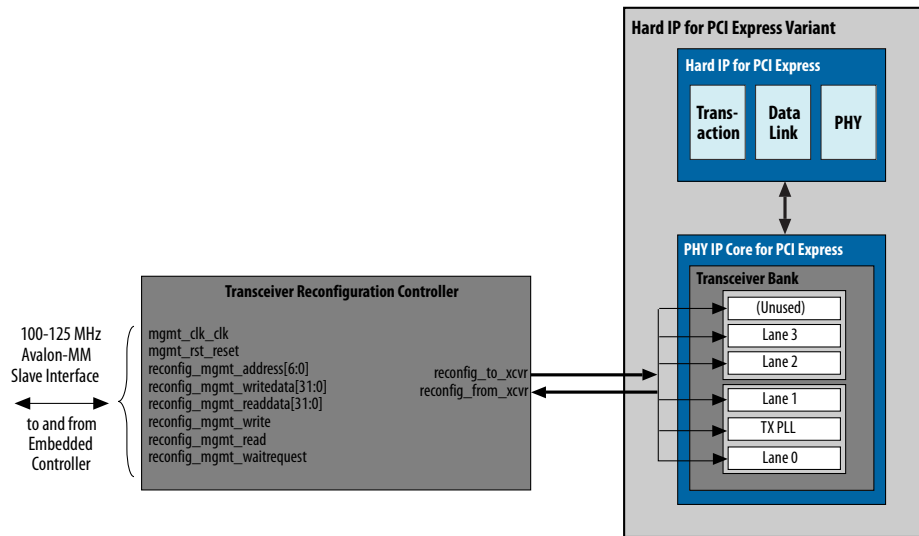
\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

**ALTERA**  
now part of Intel

**Figure 10-1: Altera Transceiver Reconfiguration Controller Connectivity**

The following figure shows the connections between the Transceiver Reconfiguration Controller instance and the PHY IP Core for PCI Express instance for a ×4 variant.



As this figure illustrates, the `reconfig_to_xcvr[ <n> 70-1:0]` and `reconfig_from_xcvr[ <n> 46-1:0]` buses connect the two components. You must provide a 100–125 MHz free-running clock to the `mgmt_clk_clk` clock input of the Transceiver Reconfiguration Controller IP Core.

Initially, each lane and TX PLL require a separate reconfiguration interface. The parameter editor reports this number in the message pane. You must take note of this number so that you can enter it as a parameter value in the Transceiver Reconfiguration Controller parameter editor. The following figure illustrates the messages reported for a Gen2 ×4 variant. The variant requires five interfaces: one for each lane and one for the TX PLL.

**Figure 10-2: Number of External Reconfiguration Controller Interfaces**

<code>ep_g2x4.DUT</code>	5 reconfiguration interfaces are required for connection to the external reconfiguration controller and the reconfig driver
<code>ep_g2x4.DUT</code>	Credit allocation in the 16 KBytes receive buffer:
<code>ep_g2x4.DUT</code>	Posted : header=16 data=16
<code>ep_g2x4.DUT</code>	Non posted: header=16 data=0
<code>ep_g2x4.DUT</code>	Completion: header=195 data=781

When you instantiate the Transceiver Reconfiguration Controller, you must specify the required **Number of reconfiguration interfaces** as the following figure illustrates.



**Figure 10-3: Specifying the Number of Transceiver Interfaces for Arria V GZ and Stratix V Devices**

## Transceiver Reconfiguration Controller

alt\_xcvr\_reconfig

Parameters

Interface Bundles

Number of reconfiguration interfaces:

Optional interface grouping:

(e.g. '2,2' or leave blank for a single bundle)

Transceiver Calibration functions

NOTE – please refer to the device handbook for reset sequence requirements between the reconfiguration co

☒ Enable offset cancellation

☒ Enable PLL calibration

☐ Create optional calibration status ports

Analog Features

☒ Enable Analog controls

☐ Enable EyeQ block

☐ Enable Bit Error Rate Block

☐ Enable decision feedback equalizer (DFE) block

☐ Enable adaptive equalization (AEQ) block

Reconfiguration Features

☐ Enable channel/PLL reconfiguration

☐ Enable PLL reconfiguration support block

**Figure 10-4: Specifying the Number of Transceiver Interfaces for Arria V and Cyclone V Devices**

**Transceiver Reconfiguration Controller**  
alt\_xcvr\_reconfig

**Parameters**  
Device family: Arria V

**Interface Bundles**  
Number of reconfiguration interfaces:   
Optional interface grouping:   
(e.g. '2,2' or leave blank for a single bundle)

**Transceiver Calibration functions**  
NOTE – please refer to the device handbook for reset sequence requirements between the reconfiguration controller and transceiver PHY.  
☒ Enable offset cancellation  
☐ Enable duty cycle calibration  
☐ Create optional calibration status ports

**Analog Features**  
☒ Enable Analog controls

**Reconfiguration Features**  
☐ Enable channel/PLL reconfiguration  
☐ Enable PLL reconfiguration support block

The Transceiver Reconfiguration Controller includes an **Optional interface grouping** parameter. Transceiver banks include six channels. For a ×4 variant, no special interface grouping is required because all 4 lanes and the TX PLL fit in one bank.

**Note:** Although you must initially create a separate logical reconfiguration interface for each lane and TX PLL in your design, when the Quartus Prime software compiles your design, it reduces the original number of logical interfaces by merging them. Allowing the Quartus Prime software to merge reconfiguration interfaces gives the Fitter more flexibility in placing transceiver channels.

**Note:** You cannot use SignalTap to observe the reconfiguration interfaces.

## Transceiver Reconfiguration Controller Connectivity for Designs Using CvP

If your design meets the following criteria:

- It enables CvP
- It includes an additional transceiver PHY that connect to the same Transceiver Reconfiguration Controller

then you must connect the PCIe `refclk` signal to the `mgmt_clk_clk` signal of the Transceiver Reconfiguration Controller and the additional transceiver PHY. In addition, if your design includes more than one Transceiver Reconfiguration Controller on the same side of the FPGA, they all must share the `mgmt_clk_clk` signal.

For more information about using the Transceiver Reconfiguration Controller, refer to the *Transceiver Reconfiguration Controller* chapter in the *Altera Transceiver PHY IP Core User Guide*.

#### Related Information

- [Altera Transceiver PHY IP Core User Guide](#)
- [Application Note 645: Dynamic Reconfiguration of PMA Controls in Stratix V Devices](#)

# Frequently Asked Questions for V-Series Avalon-MM DMA Interface for PCIe



2018.07.31

UG-01154



Subscribe



Send Feedback

The following miscellaneous facts might be of assistance in troubleshooting:

- Only the Root Ports can be loopback masters.
- Refer to Altera Solutions by searching in the Knowledge Base under Support on the Altera website.

## Related Information

[Known issues for V-Series PCIe solutions](#)

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

# V-Series Interface for PCIe Solutions User Guide Archive

# B

2018.07.31

UG-01154



Subscribe



Send Feedback

If an IP core version is not listed, the user guide for the previous IP core version applies.

IP Core Version	User Guide
15.1	<a href="#">V-Series Avalon-MM DMA Interface for PCIe Solutions User Guide</a>
14.1	<a href="#">V-Series Avalon-MM DMA Interface for PCIe Solutions User Guide</a>

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

# Document Revision History



2018.07.31

UG-01154



Subscribe



Send Feedback

## Document Revision History for the V-Series Avalon-MM DMA Interface for PCIe Solutions User Guide

Date	Version	Changes Made
2018.06.15	18.0.1	<p>Added note that Flush reads are not supported when burst mode for BAR2 is enabled.</p> <p>Updated the list of configurations supported by the Avalon-MM and Avalon-MM with DMA variants.</p>
2018.05.07	18.0	Changed all references to Intel Cyclone 10 to Intel Cyclone 10 GX.
2017.10.06	17.1	<p>Made the following change to the user guide:</p> <ul style="list-style-type: none"><li>• Added support for Intel Cyclone 10 GX devices.</li><li>• Added optional parameter to invert the RX polarity.</li><li>• Corrected <i>Feature Comparison for all Hard IP for PCI Express IP Core</i> table: The Avalon-MM DMA interface does not automatically handle out-of-order completions.</li><li>• Added missing sequence of programming steps in <i>DMA Descriptor Controller Registers</i>.</li><li>• Rebranded as Intel.</li><li>• Corrected minor errors and typos.</li></ul>
2017.05.26	17.0	<p>Made the following changes to the user guide:</p> <ul style="list-style-type: none"><li>• Added note that starting with the Intel Quartus Prime Pro Edition Software, version 17.0, the QSF assignments in the following answer <i>What assignments do I need for a PCIe Gen1, Gen2 or Gen3 design that targets an Intel Arria 10 ES2, ES3 or production device?</i> are already included in the design.</li></ul>

Intel Corporation. All rights reserved. Intel, the Intel logo, Altera, Arria, Cyclone, Enpirion, MAX, Nios, Quartus and Stratix words and logos are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Intel warrants performance of its FPGA and semiconductor products to current specifications in accordance with Intel's standard warranty, but reserves the right to make changes to any products and services at any time without notice. Intel assumes no responsibility or liability arising out of the application or use of any information, product, or service described herein except as expressly agreed to in writing by Intel. Intel customers are advised to obtain the latest version of device specifications before relying on any published information and before placing orders for products or services.

\*Other names and brands may be claimed as the property of others.

ISO  
9001:2015  
Registered

Date	Version	Changes Made
2017.05.08	17.0	<p>Made the following changes to the user guide:</p> <ul style="list-style-type: none"> <li>Corrected <i>Comparison of Avalon-ST, Avalon-MM and Avalon-MM with DMA Interfaces</i> table. Out-of-order Completions are not supported transparently for the Avalon-MM with DMA interface.</li> </ul>
2017.03.15	16.1.1	<p>Made the following changes:</p> <ul style="list-style-type: none"> <li>Rebranded as Intel.</li> </ul>
2016.10.28	16.1	<p>Made the following to the IP core changes:</p> <ul style="list-style-type: none"> <li>Increased the max DMA transfer to 1 MB for both the 128- and 256-bit interfaces.</li> </ul> <p>Made the following changes to the user guide:</p> <ul style="list-style-type: none"> <li>Corrected the number of tags supported in the <i>Feature Comparison for all Hard IP for PCI Express IP Cores</i> table.</li> <li>Added PCIe bifurcation to the <i>Feature Comparison for all Hard IP for PCI Express IP Cores</i> table. PCI bifurcation is not supported.</li> <li>Removed reference to a Linux software driver for the DMA modules which is not available.</li> <li>Added section covering design example limitations.</li> </ul>
2016.05.02	16.0	<p>Made the following changes:</p> <ul style="list-style-type: none"> <li>Redesigned the 128-bit interface to the Application Layer resulting in consistently high throughput for both on-chip and external memory.</li> <li>Added simulation support for Gen3 PIPE mode using the ModelSim, VCS, and NCSim simulators.</li> <li>Revised discussion of the DMA Descriptor Controller in the <i>Avalon-MM with DMA IP Core Architecture</i>.</li> <li>Revised <i>Read DMA Example</i> to reflect current maximum transfer size of 64 KB for 256-bit interface. The example now corresponds to an example design provided in the <i>&lt;install_dir&gt;</i>.</li> <li>Corrected description of Write Descriptor Table Avalon-MM Slave Port.</li> </ul>



Date	Version	Changes Made
2015.11.30	15.1	<ul style="list-style-type: none"><li>Added description of the Altera PCIe Reconfig Driver in the <i>Connecting the Transceiver Reconfiguration Controller IP Core</i> topic.</li><li>Added note explaining that the <i>Getting Started</i> design examples do not generate all the files necessary to download to an Altera FPGA Development Kit. Provided link to <i>AN 690: PCI Express DMA Reference Design for Stratix V Devices</i> and <i>AN 708: PCI Express DMA Reference Design Using External DDR3 Memory for Stratix V and Arria V GZ Devices</i> that includes all necessary files.</li><li>Added a FAQ chapter.</li><li>Removed <code>dlup</code> signal. This signal is no longer part of the <code>Hard IP Status</code> interface.</li><li>Fixed minor errors and typos.</li></ul>
2014.12.18	14.1	Made the following changes: <ul style="list-style-type: none"><li>Updated the instructions to create a Quartus II project and compile the design in the <i>Getting Started with the Avalon-MM DMA</i> chapter.</li><li>Added links to the <i>PCI Express Multi-Channel DMA Interface Example Design User Guide</i> and the <i>Avalon-MM DMA FIFO Example Design User Guide</i>.</li></ul>
2014.12.15	14.1	Made the following changes to the Intel Arria 10 user guide: <ul style="list-style-type: none"><li>In the <i>Getting Started</i> chapter, corrected directory path for the simulation.</li><li>Added the fact that the RX Burst Master only support dword granularity.</li><li>Added definitions for <code>test_in[2]</code>, <code>test_in[6]</code> and <code>test_in[7]</code>.</li><li>Added instructions for Quartus II compilation.</li></ul>